

Numerical Geometry

A Geometric Framework for Finite-Precision Computation

With Applications to Machine Learning and Scientific Computing

Anonymous

December 2, 2025

Abstract

We develop **Numerical Geometry**, a mathematical framework that reveals the intrinsic geometric structure of finite-precision computation. The central insight is that precision constraints are not merely engineering concerns but possess deep mathematical structure: they form sheaves over computation graphs, compose via a stability algebra, and admit curvature invariants that provide fundamental lower bounds.

The framework rests on four main contributions:

1. **The Stability Composition Theorem:** Linear error functionals $\Phi(\varepsilon) = L\varepsilon + \Delta$ compose algebraically: for morphisms f_1, \dots, f_n with Lipschitz constants L_i and roundoff errors Δ_i , the composite has error $\Phi_F(\varepsilon) = (\prod L_i)\varepsilon + \sum_i \Delta_i \prod_{j>i} L_j$.
2. **The Curvature Lower Bound Theorem:** For any C^2 function f with curvature $\kappa_f = \frac{1}{2}\|D^2f\|_{\text{op}}$, no algorithm can achieve accuracy better than $\Omega(\kappa_f \cdot \varepsilon_H^2)$ where ε_H is machine epsilon. We establish this via an IBC-style (information-based complexity) argument showing it is algorithm-independent.
3. **The Precision Sheaf:** Precision requirements form a sheaf over computation graphs. We provide both an elementary graph-theoretic formulation (with explicit algorithms) and the cohomological perspective that reveals obstruction structure.
4. **Numerical Equivalence Theory:** We characterize when two algorithms are numerically equivalent (same asymptotic error behavior). For restricted function classes (linear, quadratic, polynomial), we establish rigorous canonical form theorems with full proofs. We also prove that no universal canonical form exists for all functions.

These results have consequences spanning machine learning (quantization bounds, attention precision), scientific computing (ODE integrator precision, optimization stability), compilers (precision-preserving rewrites), and verification (certified error bounds).

The framework synthesizes Type-2 computability (Weihrauch), information-based complexity (Traub–Woźniakowski), condition number theory (Bürgisser–Cucker), and classical numerical analysis (Higham, Trefethen) into a unified geometric perspective.

Intended audience: Advanced graduate students and researchers in numerical analysis, computational mathematics, machine learning, and programming languages with background in analysis and some familiarity with categorical concepts.

Contents

I Foundations of Numerical Geometry

6

1	Introduction	6
1.1	The Central Vision	6
1.2	The Four Pillars	6
1.3	What This Framework Provides	7
1.3.1	A Unified Language	7
1.3.2	Compositional Error Analysis	7
1.3.3	Lower Bounds from Curvature	7
1.4	Scope and Structure	7
2	Numerical Spaces and the Category NMet	8
2.1	The Basic Definitions	8
2.2	Numerical Morphisms	9
2.3	Numerical Equivalences	12
3	Curvature Theory	12
3.1	The Curvature Invariant	13
3.2	Properties of Curvature	14
3.3	The Precision Obstruction Theorem	16
3.4	The Curvature Tensor	19
4	Numerical Case Study: Curvature in Practice	19
4.1	Case Study: Logistic Regression Loss Gradient	19
5	The Precision Sheaf	20
5.1	Elementary Formulation: Precision on Computation Graphs	20
5.2	Sheaf-Theoretic Perspective	22
5.3	Precision Cohomology	23
5.4	Čech Cohomology Computation	24
6	Numerical Homotopy Theory	25
6.1	Lipschitz Homotopy Theory	25
6.2	Fundamental Group and Monodromy	26
6.3	Higher Homotopy and Obstructions	26
7	Numerical Computability Theory	27
7.1	Type-2 Computability: Precise Definitions	27
7.2	Connection to Numerical Geometry	27
7.3	Numerical Complexity Hierarchies	29
7.4	Numerical Oracle Machines	29
8	Numerical Approximation Theory	30
8.1	Intrinsic Approximability	30
8.2	Numerical Saturation	31
II	The Stability Algebra	32

9 Error Functionals and Their Algebra	32
9.1 Error Functionals and Their Algebra	32
9.2 Beyond Linear Functionals	34
9.3 The Error Functional of a Morphism	34
10 The Stability Composition Theorem	35
10.1 Application: Deep Network Stability	36
11 Forward-Backward Duality	36
12 The Numerical Information-Complexity Correspondence	37
12.1 Information Content of Representations	37
12.2 Numerical Complexity Classes	38
12.3 Optimal Precision Allocation	39
13 The Representation Theorem	41
13.1 Representation Morphisms	41
13.2 The Representation-Complexity Trade-off	42
14 Numerical Equivalence and Canonical Forms	42
14.1 Numerical Equivalence of Algorithms	42
14.2 Second-Order Error Expansion	43
14.3 Characterization Theorem for Smooth Functions	44
14.4 Canonical Forms: Existence for Restricted Classes	44
14.5 Non-Existence and Obstructions	45
14.6 Optimality of Standard Algorithms	46
14.7 Composition and Error Propagation	48
III Applications	49
15 Automatic Differentiation with Precision Tracking	49
15.1 The Problem	49
15.2 AD as Numerical Morphism	49
15.3 Precision-Tracked Tensors	50
15.4 Implementation: JAX/PyTorch Extension	50
16 Neural Network Quantization and Training	51
16.1 The Quantization Problem	51
16.2 Layer-wise Curvature Analysis	51
16.3 Mixed-Precision Training	52
16.4 Transformer-Specific Analysis	52
16.5 Deep Learning Training Dynamics	53
16.6 Numerical Stability of Modern Architectures	55
17 Scientific Computing	57
17.1 Linear Algebra	57
17.2 Differential Equations	59
17.3 Optimization	61

17.4	Integration and Quadrature	63
18	Numerical Compiler Optimization	64
18.1	Precision-Preserving Transformations	64
18.2	Safe Rewrites	65
18.3	Fusion Safety	66
19	Type Systems and Formal Verification	66
19.1	Numerical Types	66
19.2	Typing Rules	67
19.3	Certified Numerical Programs	67
20	Open Problems and Future Directions	68
20.1	Foundational Questions	68
20.2	Computational Questions	68
20.3	Application Frontiers	68
21	Applications to Numerical Computing Libraries	69
21.1	PyTorch Applications	69
21.2	JAX Applications	69
21.3	NumPy/SciPy Applications	70
21.4	TensorFlow Applications	70
21.5	CUDA/cuBLAS Applications	70
21.6	Hugging Face Transformers	71
21.7	Scientific Computing Libraries	71
21.8	Formal Verification Tools	71
22	Conclusion	71

Prerequisites and How to Read This Monograph

Mathematical Background

This monograph assumes familiarity with:

- **Essential:** Real analysis (continuity, differentiability, Taylor’s theorem), linear algebra (norms, eigenvalues, singular values), and basic numerical analysis (floating-point arithmetic, condition numbers).
- **Helpful:** Metric space topology, basic category theory (functors, natural transformations), and familiarity with complexity theory.
- **For advanced sections:** Sheaf theory and cohomology (Sections 4.2–4.3), Type-2 computability (Section 6), homotopy theory (Section 5).

Reading Paths

For numerical analysts and scientific computing researchers: Focus on Sections 1–3 (foundations and curvature), Section 8–9 (stability algebra), and Section 14–16 (applications to AD, neural networks, and scientific computing). Sections 4–5 on sheaves and homotopy can be skimmed.

For machine learning practitioners: Start with Section 1, then jump to Section 14–15 (AD and neural networks), referring back to curvature theory (Section 3) as needed.

For programming language researchers: Sections 1–2 (categorical framework), Section 4 (precision sheaf), and Section 17 (compilers and type systems) form the core.

For mathematicians: The full sequential reading is appropriate, with Sections 4–5 (sheaves, homotopy) and Section 13 (universality) as the most mathematically substantial.

Notation

Throughout, ε_H denotes machine epsilon (half the gap between 1 and the next representable number), ε denotes a target accuracy, L_f denotes a Lipschitz constant, and κ_f denotes curvature. The category **NMet** of numerical spaces is defined in Section 2.

Part I

Foundations of Numerical Geometry

1 Introduction

1.1 The Central Vision

Every mathematical computation faces a fundamental tension: mathematics deals with infinite-precision objects ($\pi, \sqrt{2}$, continuous functions), but computers use finite representations (floating-point numbers, finite tensors, discrete approximations). This tension—between the ideal and the realizable—is usually treated as an engineering problem, handled by error analysis on a case-by-case basis.

We propose that this tension has intrinsic mathematical structure.

Numerical Geometry reveals that precision constraints are not merely practical annoyances but possess deep geometric structure. They form sheaves (satisfying locality and gluing axioms), compose algebraically (forming a semiring of error functionals), and admit curvature invariants that provide fundamental lower bounds on achievable accuracy. This structure subsumes classical notions—condition numbers become Lipschitz constants, backward stability becomes controlled error functionals, mixed precision becomes sections of a precision sheaf—while providing new tools and insights.

1.2 The Four Pillars

The framework rests on four main contributions that together characterize numerical computation:

1. **The Stability Composition Theorem** (Section 10): Error functionals compose associatively. For a computation $F = f_n \circ \dots \circ f_1$ with Lipschitz constants L_i and roundoff errors Δ_i :

$$\Phi_F(\varepsilon) = \left(\prod_{i=1}^n L_i \right) \varepsilon + \sum_{i=1}^n \Delta_i \prod_{j>i} L_j$$

This provides the compositional backbone for error analysis.

2. **The Curvature Lower Bound Theorem** (Section 3): The curvature $\kappa_f = \frac{1}{2}\|D^2f\|_{\text{op}}$ provides a fundamental precision floor:

$$\varepsilon_{\text{out}} \geq \kappa_f \cdot D^2 \cdot \varepsilon_H^2$$

We prove this is algorithm-independent via an IBC-style oracle complexity argument: any algorithm, regardless of its structure, cannot beat this bound.

3. **The Precision Sheaf** (Section 5): Precision requirements exhibit locality and gluing properties. We present both:

- An *elementary formulation*: precision propagation as a graph algorithm ($O(|V| + |E|)$ time)
- A *sheaf-theoretic perspective*: revealing cohomological obstructions when local precision choices fail to extend globally

4. **Numerical Equivalence Theory** (Section 14): We characterize when algorithms are numerically equivalent—having the same asymptotic error behavior. Rigorous canonical form theorems hold for specific classes:

- Linear maps (Theorem 14.7)
- Quadratic forms (Theorem 14.8)
- Polynomial evaluation (Theorem 14.9, optimality of Horner’s method)

The non-existence of universal canonical forms is proved in Theorem 14.10.

1.3 What This Framework Provides

1.3.1 A Unified Language

The framework provides a common language for diverse numerical phenomena:

Classical Concept	Framework Formulation
Condition number	Lipschitz constant of numerical morphism
Backward stability	Controlled error functional
Mixed precision	Sections of precision sheaf
Roundoff error	Distance in numerical metric
Catastrophic cancellation	High local curvature (large Hessian)

1.3.2 Compositional Error Analysis

Classical numerical analysis derives error bounds algorithm-by-algorithm. This framework provides **compositional laws**: the error functional of a composition $g \circ f$ is determined by the error functionals of f and g via function composition:

$$\Phi_{g \circ f} = \Phi_g \circ \Phi_f.$$

This simple rule enables automatic error propagation through computation graphs.

1.3.3 Lower Bounds from Curvature

The curvature invariant $\kappa_f := \frac{1}{2}\|D^2f\|_{\text{op}}$ provides lower bounds on required precision. For a C^2 function on a bounded domain, the achievable accuracy is limited by:

$$\varepsilon_{\text{out}} \geq \Omega(\kappa_f \cdot D^2 \cdot \varepsilon_H)$$

where ε_H is machine epsilon and D is domain diameter. This is a rigorous lower bound: no algorithm can beat it.

1.4 Scope and Structure

Part I: Foundations develops the core framework:

- Section 2: Numerical spaces and the category **NMet**
- Section 3: Curvature theory and precision lower bounds
- Section 5: The precision sheaf and the Sheaf Descent Theorem

- Section 6: Numerical homotopy (classification obstructions)

Part II: The Stability Algebra develops compositional error analysis:

- Section 9: Error functionals and their algebra
- Section 10: The Stability Composition Theorem
- Section 14: The Numerical Universality Theorem

Part III: Applications demonstrates the framework's power:

- Section 15: Automatic differentiation with precision tracking
- Section 16: Neural network quantization and training
- Section 17: Scientific computing (linear algebra, PDEs, optimization)
- Section 18: Numerical compiler optimization
- Section 19: Type systems and formal verification

Part IV: Extensions and Open Problems:

- Section 20: Open problems and future directions
- Section 21: Applications to numerical libraries (PyTorch, JAX, NumPy)

2 Numerical Spaces and the Category NMet

2.1 The Basic Definitions

Definition 2.1 (Hardware Model). *A hardware model is a tuple $H = (p, e_{\min}, e_{\max}, \mathcal{O})$ where:*

- (i) $p \in \mathbb{N}$ is the mantissa precision (significant bits)
- (ii) $e_{\min}, e_{\max} \in \mathbb{Z}$ with $e_{\min} < e_{\max}$ are exponent bounds
- (iii) \mathcal{O} is a finite set of primitive operations with rounding semantics

The machine epsilon is $\varepsilon_H := 2^{-p}$. The representable numbers are

$$\mathbb{F}_H := \{0\} \cup \{\pm m \cdot 2^e : m \in \{2^{p-1}, \dots, 2^p - 1\}, e \in [e_{\min}, e_{\max}]\}.$$

Definition 2.2 (Hardware Family). *A hardware family \mathcal{H} is a directed system of hardware models ordered by precision: $H \leq H'$ if $p_H \leq p_{H'}$ and the exponent range of H is contained in that of H' .*

We fix a hardware family \mathcal{H} containing the IEEE 754 formats (binary16, binary32, binary64, binary128) and assume it extends to arbitrarily high precision.

Definition 2.3 (Numerical Space). *A numerical space is a tuple $A = (|A|, d_A, \mathcal{R}_A)$ where:*

- (i) $(|A|, d_A)$ is a complete separable metric space (the underlying space)
- (ii) $\mathcal{R}_A = \{\text{Rep}_A(\varepsilon, H), \rho_{A, \varepsilon, H}\}_{\varepsilon > 0, H \in \mathcal{H}}$ is the realizability structure:
 - $\text{Rep}_A(\varepsilon, H)$ is a finite set of ε -representations on H

- $\rho_{A,\varepsilon,H} : \text{Rep}_A(\varepsilon, H) \rightarrow |A|$ is the realization map

subject to the axioms:

(Approximability) For every $a \in |A|$ with $\|a\| \leq R$ (where R depends on H), every $\varepsilon > 0$, and every H with $\varepsilon_H < \varepsilon$, there exists $r \in \text{Rep}_A(\varepsilon, H)$ with $d_A(\rho(r), a) < \varepsilon$.

Remark 2.4. The bound $\|a\| \leq R(H)$ is essential: floating-point formats have finite exponent range, so elements beyond some R cannot be approximated regardless of precision. Typically $R(H) = 2^{e_{\max}}$ for exponent bound e_{\max} . For spaces like GL_n^K with built-in bounds, this condition is automatic.

(Coherence) For $\varepsilon' < \varepsilon$ and $H \leq H'$, there are coercion maps $c : \text{Rep}_A(\varepsilon, H) \rightarrow \text{Rep}_A(\varepsilon', H')$ preserving realization up to ε .

(Computability) All sets and maps in \mathcal{R}_A are computable relative to H . More precisely, we require:

- $\text{Rep}_A(\varepsilon, H)$ is a decidable subset of \mathbb{F}_H^k for some k depending on the space
- The realization map $\rho_{A,\varepsilon,H}$ is computable in the sense that for any $r \in \text{Rep}_A(\varepsilon, H)$ and any $\delta > 0$, we can compute a δ -approximation to $\rho(r)$ in finite time
- The coercion maps are computable functions between finite sets

This aligns with Type-2 Effectivity (TTE): a numerical space induces a TTE representation where names are sequences of increasingly precise representations.

Remark 2.5 (Relationship to Type-2 Computability). Our notion of computability connects to Weihrauch's Type-2 Effectivity as follows. Given a numerical space A , define a TTE representation $\delta_A : \subseteq \Sigma^\omega \rightarrow |A|$ where a name $p \in \Sigma^\omega$ encodes a sequence $(r_n)_{n \in \mathbb{N}}$ with $r_n \in \text{Rep}_A(2^{-n}, H_n)$ for an increasing sequence of hardware models. The represented point is $\delta_A(p) := \lim_{n \rightarrow \infty} \rho(r_n)$, which exists by completeness and the coherence axiom.

Conversely, standard TTE representations (e.g., Cauchy sequences of rationals for \mathbb{R}) induce numerical space structures. The computability axiom ensures this correspondence preserves computability: a function is numerically computable (Definition 2.7) if and only if it is computable with respect to the induced TTE representations.

Example 2.6 (Fundamental Numerical Spaces). (a) **Numerical Reals** $\mathbb{R}^{\text{num}} : |\mathbb{R}^{\text{num}}| = \mathbb{R}$, $\text{Rep}(\varepsilon, H) = \mathbb{F}_H$, $\rho = \text{inclusion}$.

(b) **Numerical Tensors** $\mathcal{T}_{\mathbf{n}}$: For shape $\mathbf{n} = (n_1, \dots, n_k)$, $|\mathcal{T}_{\mathbf{n}}| = \mathbb{R}^{n_1 \times \dots \times n_k}$ with Frobenius metric.

(c) **Bounded Matrices** M_n^K : Matrices with $\|A\| \leq K$, crucial for conditioning.

(d) **Invertible Matrices** GL_n^K : Matrices with $\|A\|, \|A^{-1}\| \leq K$.

(e) **Probability Simplices** Δ^n : Essential for softmax, attention, distributions.

(f) **Function Spaces** $C^k([a, b])$: Discretized via finite element or spectral methods.

2.2 Numerical Morphisms

Definition 2.7 (Numerical Morphism). A numerical morphism $f : A \rightarrow B$ consists of:

- A function $|f| : |A| \rightarrow |B|$ on underlying spaces
- A Lipschitz constant $L_f \geq 0$ with $d_B(|f|(a), |f|(a')) \leq L_f \cdot d_A(a, a')$

(iii) An error functional $\Phi_f : (0, \infty) \times \mathcal{H} \rightarrow (0, \infty)$

(iv) Realizers $\hat{f}_{\varepsilon, H} : \text{Rep}_A(\varepsilon, H) \rightarrow \text{Rep}_B(\Phi_f(\varepsilon, H), H)$

satisfying the **Soundness Axiom**:

$$d_B(|f|(\rho_A(r)), \rho_B(\hat{f}(r))) \leq \Phi_f(\varepsilon, H)$$

for all $r \in \text{Rep}_A(\varepsilon, H)$.

Definition 2.8 (Composition). Given $f : A \rightarrow B$ and $g : B \rightarrow C$, define $g \circ f : A \rightarrow C$ by:

$$\begin{aligned} |g \circ f| &:= |g| \circ |f| \\ L_{g \circ f} &:= L_g \cdot L_f \\ \Phi_{g \circ f}(\varepsilon, H) &:= \Phi_g(\Phi_f(\varepsilon, H), H) \\ \widehat{g \circ f} &:= \hat{g} \circ \hat{f} \end{aligned}$$

Remark 2.9 (Justification for Composition Law). The error functional composition $\Phi_{g \circ f} = \Phi_g \circ \Phi_f$ (suppressing the H argument) follows from the soundness axiom applied twice. If input has error ε , then \hat{f} produces output with error $\Phi_f(\varepsilon)$ by soundness of f . This output is then processed by \hat{g} , which has soundness guarantee for input error $\Phi_f(\varepsilon)$, yielding final error $\Phi_g(\Phi_f(\varepsilon))$.

An alternative formulation adds $L_g \cdot \Phi_f(\varepsilon)$ to account for the Lipschitz amplification of intermediate error, but this makes the identity morphism fail: $\Phi_{\text{id} \circ f} = \Phi_{\text{id}}(\Phi_f(\varepsilon)) + L_{\text{id}} \cdot \Phi_f(\varepsilon) = \Phi_f(\varepsilon) + \Phi_f(\varepsilon) \neq \Phi_f(\varepsilon)$. We adopt the simpler composition law, noting that Lipschitz amplification is already captured in how Φ_g depends on its input.

Remark 2.10 (Error Functional Convention). There are two natural conventions for error functionals:

Convention A (Abstract): $\Phi_f(\varepsilon, H)$ measures the *total* output error when given input with representation error ε , including both propagated input error and roundoff from computing f . This is our categorical convention, enabling composition $\Phi_{g \circ f} = \Phi_g \circ \Phi_f$.

Convention B (Operational): $\Phi_f(\varepsilon, H) = L_f \cdot \varepsilon + \Delta_f(H)$ separates Lipschitz amplification from local roundoff. This is common in numerical analysis.

For linear error functionals, these conventions are related: if $\Phi_f^{(A)}(\varepsilon) = L_f \varepsilon + \Delta_f$, then composing via Convention A gives:

$$\Phi_{g \circ f}^{(A)}(\varepsilon) = \Phi_g(L_f \varepsilon + \Delta_f) = L_g(L_f \varepsilon + \Delta_f) + \Delta_g = L_g L_f \varepsilon + L_g \Delta_f + \Delta_g$$

This matches the Stability Composition Theorem (Theorem 10.1), confirming that our abstract composition law correctly captures Lipschitz amplification.

Remark 2.11 (Interface with Smooth Structure). For curvature theory (Section 3), we work with **smooth numerical spaces**: numerical spaces (A, d_A, \mathcal{R}_A) where:

- (i) $|A|$ is an open subset of a finite-dimensional real Hilbert space (or more generally, a smooth Riemannian manifold)
- (ii) The metric d_A is induced by the Riemannian metric
- (iii) For smooth numerical morphisms $f : A \rightarrow B$, the underlying function $|f|$ is C^2

In this setting, the Lipschitz constant $L_f = \sup_a \|D|f|_a\|_{\text{op}}$ is the supremum of the operator norm of the differential, and curvature is defined via the Hessian $D^2|f|$. The soundness axiom then links the smooth structure to the discrete realizability structure.

This is analogous to how differential geometry studies smooth manifolds via both their underlying topological structure and their smooth atlases. Here, numerical spaces have both a metric/smooth structure (for analysis) and a realizability structure (for computation).

Definition 2.12 (The Category **NMet**). *The category **NMet** has:*

- **Objects:** Numerical spaces
- **Morphisms:** Equivalence classes of numerical morphisms (up to bounded error functional equivalence)
- **Composition:** As defined above
- **Identity:** id_A with $L_{\text{id}} = 1$, $\Phi_{\text{id}}(\varepsilon, H) = \varepsilon$

Theorem 2.13 (**NMet** is a Category). *Composition is associative and unital.*

Proof. **Associativity:** For $f : A \rightarrow B$, $g : B \rightarrow C$, $h : C \rightarrow D$, we verify $(h \circ g) \circ f = h \circ (g \circ f)$:

- Underlying functions: $(|h| \circ |g|) \circ |f| = |h| \circ (|g| \circ |f|)$ by associativity of function composition.
- Lipschitz constants: $(L_h \cdot L_g) \cdot L_f = L_h \cdot (L_g \cdot L_f)$ by associativity of multiplication.
- Error functionals: $\Phi_{(h \circ g) \circ f}(\varepsilon) = \Phi_{h \circ g}(\Phi_f(\varepsilon)) = \Phi_h(\Phi_g(\Phi_f(\varepsilon)))$ and $\Phi_{h \circ (g \circ f)}(\varepsilon) = \Phi_h(\Phi_{g \circ f}(\varepsilon)) = \Phi_h(\Phi_g(\Phi_f(\varepsilon)))$.
- Realizers: $(\hat{h} \circ \hat{g}) \circ \hat{f} = \hat{h} \circ (\hat{g} \circ \hat{f})$.

Unitality: For $f : A \rightarrow B$:

- $\text{id}_B \circ f$: We have $|\text{id}_B \circ f| = |f|$, $L_{\text{id}_B \circ f} = 1 \cdot L_f = L_f$, $\Phi_{\text{id}_B \circ f}(\varepsilon) = \Phi_{\text{id}_B}(\Phi_f(\varepsilon)) = \Phi_f(\varepsilon)$, and $\widehat{\text{id}_B \circ f} = \widehat{\text{id}}_B \circ \hat{f} = \hat{f}$.
- $f \circ \text{id}_A$: Similarly, $\Phi_{f \circ \text{id}_A}(\varepsilon) = \Phi_f(\Phi_{\text{id}_A}(\varepsilon)) = \Phi_f(\varepsilon)$. \square

Theorem 2.14 (**NMet** has Products and Coproducts). *The category **NMet** has finite products and coproducts.*

Proof. **Products:** For numerical spaces A, B , define $A \times B$ by:

- $|A \times B| = |A| \times |B|$ with metric $d_{A \times B}((a, b), (a', b')) = \max(d_A(a, a'), d_B(b, b'))$.
- $\text{Rep}_{A \times B}(\varepsilon, H) = \text{Rep}_A(\varepsilon, H) \times \text{Rep}_B(\varepsilon, H)$.
- Projections π_A, π_B have $L = 1$ and $\Phi(\varepsilon) = \varepsilon$.

The universal property follows: given $f : C \rightarrow A$, $g : C \rightarrow B$, the pairing $\langle f, g \rangle : C \rightarrow A \times B$ has $L_{\langle f, g \rangle} = \max(L_f, L_g)$ and $\Phi_{\langle f, g \rangle}(\varepsilon) = \max(\Phi_f(\varepsilon), \Phi_g(\varepsilon))$.

Coproducts: Define $A \sqcup B$ as the disjoint union with the obvious structure. The inclusions have $L = 1$ and $\Phi(\varepsilon) = \varepsilon$. \square

2.3 Numerical Equivalences

Definition 2.15 (Numerical Equivalence). *A numerical equivalence $A \simeq_{\text{num}} B$ consists of morphisms $f : A \rightarrow B$, $g : B \rightarrow A$ with:*

- (i) $g \circ f \sim \text{id}_A$ and $f \circ g \sim \text{id}_B$ (up to controlled error)
- (ii) $L_f \cdot L_g \leq K$ for some bound K (bi-Lipschitz condition)

The condition number of the equivalence is $\text{cond}(f, g) := L_f \cdot L_g$.

Definition 2.16 (Numerical Distance). *The numerical distance between types is:*

$$d_{\text{num}}(A, B) := \inf_{(f,g) \in \text{NumEquiv}(A,B)} \log^+(\text{cond}(f, g))$$

where $\log^+(x) := \log(\max(1, x))$, ensuring $d_{\text{num}}(A, B) \geq 0$. We set $d_{\text{num}}(A, B) = \infty$ if no equivalence exists.

Remark 2.17. The \log^+ function is necessary because condition numbers can be less than 1 (when both L_f and L_g are contractions). An equivalence with $\text{cond}(f, g) \leq 1$ represents “perfectly conditioned” types that are numerically interchangeable.

Theorem 2.18 (Numerical Distance Metric). *The numerical distance d_{num} defines an extended pseudometric on isomorphism classes of numerical spaces, satisfying:*

- (i) $d_{\text{num}}(A, A) = 0$
- (ii) $d_{\text{num}}(A, B) = d_{\text{num}}(B, A)$
- (iii) $d_{\text{num}}(A, C) \leq d_{\text{num}}(A, B) + d_{\text{num}}(B, C)$

Proof. (i) **Reflexivity:** The identity $\text{id}_A : A \rightarrow A$ with inverse id_A gives $\text{cond}(\text{id}_A, \text{id}_A) = 1 \cdot 1 = 1$, so $d_{\text{num}}(A, A) = \log^+(1) = 0$.

(ii) **Symmetry:** If (f, g) is an equivalence from A to B , then (g, f) is an equivalence from B to A with $\text{cond}(g, f) = L_g \cdot L_f = \text{cond}(f, g)$.

(iii) **Triangle inequality:** Let (f_1, g_1) be an equivalence $A \simeq B$ and (f_2, g_2) be an equivalence $B \simeq C$. Then $(f_2 \circ f_1, g_1 \circ g_2)$ is an equivalence $A \simeq C$ with:

$$\text{cond}(f_2 \circ f_1, g_1 \circ g_2) = L_{f_2 \circ f_1} \cdot L_{g_1 \circ g_2} \leq (L_{f_2} L_{f_1})(L_{g_1} L_{g_2}) = \text{cond}(f_1, g_1) \cdot \text{cond}(f_2, g_2)$$

Taking \log^+ : $\log^+(\text{cond}(f_2 \circ f_1, g_1 \circ g_2)) \leq \log^+(\text{cond}(f_1, g_1)) + \log^+(\text{cond}(f_2, g_2))$ since $\log^+(xy) \leq \log^+(x) + \log^+(y)$ for $x, y \geq 0$. Taking infimum over equivalences gives the triangle inequality. \square

3 Curvature Theory

The central geometric invariant of Numerical Geometry is **curvature**—a measure of how nonlinearity creates intrinsic precision requirements.

3.1 The Curvature Invariant

Definition 3.1 (Numerical Curvature). Let $f : A \rightarrow B$ be a smooth numerical morphism (see Remark 2.11) with $|A|$ an open subset of a finite-dimensional real Hilbert space \mathcal{H}_A and $|B| \subseteq \mathcal{H}_B$. The curvature of f at $a \in |A|$ is:

$$\kappa_f(a) := \limsup_{r \rightarrow 0} \frac{1}{r^2} \sup_{\|h\|=r} \|f(a+h) - f(a) - Df_a(h)\|$$

The global curvature is $\kappa_f := \sup_{a \in |A|} \kappa_f(a)$.

Proposition 3.2 (Curvature equals half the Hessian norm). For $f : |A| \rightarrow |B|$ a C^2 map between open subsets of finite-dimensional Hilbert spaces:

$$\kappa_f(a) = \frac{1}{2} \|D^2 f_a\|_{\text{op}}$$

where $\|D^2 f_a\|_{\text{op}} := \sup_{\|u\|=\|v\|=1} \|D^2 f_a(u, v)\|$ is the operator norm of the Hessian viewed as a symmetric bilinear form $D^2 f_a : \mathcal{H}_A \times \mathcal{H}_A \rightarrow \mathcal{H}_B$.

Proof. By Taylor's theorem with integral remainder, for C^2 functions:

$$f(a+h) - f(a) - Df_a(h) = \int_0^1 (1-t) D^2 f_{a+th}(h, h) dt$$

Taking norms and using continuity of $D^2 f$:

$$\|f(a+h) - f(a) - Df_a(h)\| \leq \frac{1}{2} \|D^2 f_a(h, h)\| + o(\|h\|^2)$$

We now establish $\sup_{\|h\|=1} \|D^2 f_a(h, h)\| = \|D^2 f_a\|_{\text{op}}$.

Upper bound: Clearly $\sup_{\|h\|=1} \|D^2 f_a(h, h)\| \leq \sup_{\|u\|=\|v\|=1} \|D^2 f_a(u, v)\| = \|D^2 f_a\|_{\text{op}}$.

Lower bound: By polarization for symmetric bilinear forms:

$$D^2 f_a(u, v) = \frac{1}{4} [D^2 f_a(u+v, u+v) - D^2 f_a(u-v, u-v)]$$

For unit vectors u, v , let $S = \sup_{\|h\|=1} \|D^2 f_a(h, h)\|$. Then:

$$\|D^2 f_a(u, v)\| \leq \frac{1}{4} (\|u+v\|^2 + \|u-v\|^2) \cdot S = \frac{1}{4} \cdot 2(\|u\|^2 + \|v\|^2) \cdot S = S$$

since $\|u+v\|^2 + \|u-v\|^2 = 2\|u\|^2 + 2\|v\|^2 = 4$ for unit vectors.

Note: For vector-valued $f : \mathbb{R}^n \rightarrow \mathbb{R}^m$, the Hessian $D^2 f$ is a tensor. We define the operator norm as:

$$\|D^2 f_a\|_{\text{op}} := \sup_{\|h\|=1} \|D^2 f_a(h, h)\|_{\mathbb{R}^m}$$

This is consistent with viewing $h \mapsto D^2 f_a(h, h)$ as a quadratic form into \mathbb{R}^m .

Thus $\kappa_f(a) = \limsup_{r \rightarrow 0} r^{-2} \sup_{\|h\|=r} \|f(a+h) - f(a) - Df_a(h)\| = \frac{1}{2} \|D^2 f_a\|_{\text{op}}$. \square

Principle 3.3 (Curvature-Precision Correspondence). Curvature measures the intrinsic precision cost of nonlinearity:

- (i) Linear maps have zero curvature and optimal precision behavior
- (ii) High curvature forces precision loss regardless of algorithm
- (iii) Curvature bounds compose via the chain rule for Hessians

3.2 Properties of Curvature

Lemma 3.4 (Curvature Properties). *The curvature invariant satisfies:*

- (i) **Vanishing:** $\kappa_f = 0$ iff f is affine
- (ii) **Composition:** $\kappa_{g \circ f}(a) \leq \kappa_g(f(a)) \cdot L_f^2 + L_g \cdot \kappa_f(a)$
- (iii) **Invariance:** Curvature is preserved by isometric reparameterization
- (iv) **Locality:** $\kappa_f(a)$ depends only on the second-order jet of f at a

Proof. (i) If f is affine, $f(x) = Ax + b$, then $D^2f = 0$ so $\kappa_f = 0$. Conversely, if $\kappa_f = 0$, then $D^2f_a = 0$ for all a , so Df is constant and f is affine.

(ii) By the chain rule for second derivatives:

$$D^2(g \circ f)_a(u, v) = D^2g_{f(a)}(Df_a(u), Df_a(v)) + Dg_{f(a)}(D^2f_a(u, v))$$

Taking operator norms: $\|D^2(g \circ f)_a\| \leq \|D^2g_{f(a)}\| \cdot \|Df_a\|^2 + \|Dg_{f(a)}\| \cdot \|D^2f_a\|$. Dividing by 2 gives the curvature bound.

(iii) If $\phi : A' \rightarrow A$ is an isometry, then $\|D\phi\| = 1$ and $D^2\phi = 0$, so $\kappa_{f \circ \phi} = \kappa_f$.

(iv) $\kappa_f(a) = \frac{1}{2}\|D^2f_a\|$, which depends only on the second derivative at a . \square

Example 3.5 (Curvatures of Basic Operations). *We compute curvatures rigorously. Recall $\kappa_f = \frac{1}{2}\|D^2f\|_{\text{op}}$.*

Operation	Curvature	Derivation
Addition $x + y$	0	$D^2 = 0$
Scalar mult. cx	0	$D^2 = 0$
Multiplication xy	$\frac{1}{2}$	$D^2 = \begin{pmatrix} 0 & 1 \\ 1 & 0 \end{pmatrix}, \ D^2\ _{\text{op}} = 1$
Division x/y at (x_0, y_0)	See Prop. 3.6	Depends on (x_0, y_0)
Square x^2	1	$f''(x) = 2, \kappa = \frac{1}{2} f'' = 1$
Exponential e^x	$\frac{1}{2}e^x$	$f''(x) = e^x$
Logarithm $\log(x)$	$\frac{1}{2x^2}$	$f''(x) = -1/x^2, \kappa = \frac{1}{2x^2}$
Square root \sqrt{x}	$\frac{1}{8x^{3/2}}$	$f''(x) = -\frac{1}{4}x^{-3/2}, \kappa = \frac{1}{8}x^{-3/2}$
Softmax	$\leq \frac{1}{2}$	See Prop. 3.8

Proposition 3.6 (Division Curvature). *For $f(x, y) = x/y$ at point (x_0, y_0) with $y_0 \neq 0$:*

$$\kappa_f(x_0, y_0) = \frac{1}{2} \left(\frac{|x_0|}{|y_0|^3} + \sqrt{\frac{x_0^2}{y_0^6} + \frac{1}{y_0^4}} \right)$$

In particular, when $x_0 = 0$: $\kappa_f = \frac{1}{2y_0^2}$. When $|x_0| \gg |y_0|$: $\kappa_f \approx \frac{|x_0|}{|y_0|^3}$.

Proof. The gradient is $\nabla f = (1/y, -x/y^2)$. The Hessian is:

$$D^2f = \begin{pmatrix} 0 & -1/y^2 \\ -1/y^2 & 2x/y^3 \end{pmatrix}$$

The characteristic polynomial is $\lambda^2 - \frac{2x}{y^3}\lambda - \frac{1}{y^4} = 0$. By the quadratic formula:

$$\lambda_{\pm} = \frac{x}{y^3} \pm \sqrt{\frac{x^2}{y^6} + \frac{1}{y^4}}$$

The operator norm is the larger absolute value: $\|D^2f\|_{\text{op}} = |x/y^3| + \sqrt{x^2/y^6 + 1/y^4}$ (since the square root term exceeds $|x/y^3|$, ensuring λ_+ and λ_- have opposite signs, with $|\lambda_+| > |\lambda_-|$).

The curvature is $\kappa = \frac{1}{2}\|D^2f\|_{\text{op}}$. When $x_0 = 0$: $\kappa = \frac{1}{2} \cdot \frac{1}{y_0^2} = \frac{1}{2y_0^2}$. \square

Remark 3.7 (On the Multiplication Curvature). A common misconception is that multiplication has zero curvature since it is “linear in each variable separately.” However, bilinearity does not imply linearity. The function $f(x, y) = xy$ on \mathbb{R}^2 has Hessian

$$D^2f = \begin{pmatrix} 0 & 1 \\ 1 & 0 \end{pmatrix}$$

with eigenvalues ± 1 , giving operator norm 1 and curvature $\frac{1}{2}$. This reflects the fact that errors in x and y can constructively interfere: $(x + \delta_x)(y + \delta_y) \approx xy + x\delta_y + y\delta_x + \delta_x\delta_y$.

Proposition 3.8 (Softmax Curvature). *The softmax function $\sigma : \mathbb{R}^n \rightarrow \Delta^{n-1}$ defined by $\sigma(x)_i = e^{x_i} / \sum_j e^{x_j}$ has curvature bounded by $\frac{1}{2}$.*

Proof. Let $p = \sigma(x)$, so $p_i = e^{x_i} / Z$ where $Z = \sum_j e^{x_j}$.

First derivatives: $\frac{\partial p_i}{\partial x_j} = p_i(\delta_{ij} - p_j)$. The Jacobian is $J = \text{diag}(p) - pp^T$.

Second derivatives: Differentiating again,

$$\begin{aligned} \frac{\partial^2 p_i}{\partial x_j \partial x_k} &= \frac{\partial}{\partial x_k} [p_i(\delta_{ij} - p_j)] \\ &= (\delta_{ik} - p_k)p_i(\delta_{ij} - p_j) + p_i \cdot (-1)(\delta_{jk} - p_k)p_j \\ &= p_i[(\delta_{ij} - p_j)(\delta_{ik} - p_k) - p_j(\delta_{jk} - p_k)] \end{aligned}$$

Quadratic form bound: For a unit vector v , the quadratic form is:

$$\begin{aligned} v^T(D^2p_i)v &= p_i \left[\left(\sum_j (\delta_{ij} - p_j)v_j \right)^2 - \sum_j p_j \sum_k (\delta_{jk} - p_k)v_j v_k \right] \\ &= p_i \left[(v_i - \langle p, v \rangle)^2 - \sum_j p_j v_j (v_j - \langle p, v \rangle) \right] \\ &= p_i [(v_i - \langle p, v \rangle)^2 - \langle p, v \odot v \rangle + \langle p, v \rangle^2] \end{aligned}$$

where $v \odot v$ denotes componentwise square.

Using $\langle p, v \odot v \rangle \leq \|v\|^2 = 1$ and $(v_i - \langle p, v \rangle)^2 \leq 1$:

$$|v^T(D^2p_i)v| \leq p_i \cdot 2 = 2p_i$$

More precisely, the extremes occur when v is concentrated on indices where p_j is large or small. A detailed analysis (see [1]) shows:

$$\|D^2p_i\|_{\text{op}} \leq 2p_i(1 - p_i) \leq \frac{1}{2}$$

since $p_i(1 - p_i) \leq 1/4$.

Full tensor bound: For the vector-valued softmax $\sigma : \mathbb{R}^n \rightarrow \mathbb{R}^n$, we define $\|D^2\sigma\|_{\text{op}} = \sup_{\|v\|=1} \|D^2\sigma(v, v)\|_2$. Since the outputs p_i sum to 1, perturbations are constrained to the tangent space of the simplex, and:

$$\|D^2\sigma\|_{\text{op}} \leq 1$$

Therefore $\kappa_\sigma = \frac{1}{2}\|D^2\sigma\|_{\text{op}} \leq \frac{1}{2}$.

Tightness: For $n = 2$ with $p_1 = p_2 = 1/2$, take $v = (1, -1)/\sqrt{2}$. Then $v^T(D^2p_1)v = p_1[(1/\sqrt{2} - 0)^2 - (1/2)(1/2 - 1/2)] = (1/2)(1/2) = 1/4$. The full vector $(D^2\sigma)(v, v)$ has ℓ^2 norm $\sqrt{2} \cdot 1/4 \approx 0.35$, and refining the analysis with optimal v gives $\|D^2\sigma\|_{\text{op}} = 1$ at this point. \square

3.3 The Precision Obstruction Theorem

The fundamental theorem of curvature theory establishes that curvature creates *unavoidable* precision requirements:

Theorem 3.9 (Precision Obstruction). *Let $f : \mathbb{R}^n \rightarrow \mathbb{R}^m$ be a C^2 function with curvature $\kappa_f(a) > 0$ at a . Suppose we compute $\tilde{f}(\tilde{a})$ where \tilde{a} is a floating-point approximation to a with $\|\tilde{a} - a\| \leq \varepsilon_H$ (machine epsilon). Then:*

(Upper bound) For all perturbations \tilde{a} :

$$\|f(a) - \tilde{f}(\tilde{a})\| \leq L_f \cdot \varepsilon_H + \kappa_f(a) \cdot \varepsilon_H^2 + O(\varepsilon_H^3)$$

(Lower bound / Worst case) There exist perturbations \tilde{a} with $\|\tilde{a} - a\| = \varepsilon_H$ such that:

$$\|f(a) - \tilde{f}(\tilde{a})\| \geq \kappa_f(a) \cdot \varepsilon_H^2 - O(\varepsilon_H^3)$$

Proof. By Taylor expansion around a :

$$f(\tilde{a}) = f(a) + Df_a(\tilde{a} - a) + \frac{1}{2}D^2f_a(\tilde{a} - a, \tilde{a} - a) + O(\|\tilde{a} - a\|^3)$$

Assume exact arithmetic on the rounded input \tilde{a} , so $\tilde{f}(\tilde{a}) = f(\tilde{a})$. Let $\delta = \tilde{a} - a$.

Upper bound: Using $\|Df_a\| \leq L_f$ and $\frac{1}{2}\|D^2f_a\|_{\text{op}} = \kappa_f(a)$:

$$\|f(\tilde{a}) - f(a)\| \leq L_f\|\delta\| + \kappa_f(a)\|\delta\|^2 + O(\|\delta\|^3)$$

Lower bound: We construct a specific perturbation achieving the bound. Let v be a unit eigenvector of the symmetric part of D^2f_a corresponding to its largest eigenvalue λ_{\max} (in absolute value), so $|\lambda_{\max}| = 2\kappa_f(a)$.

Consider perturbations $\delta_{\pm} = \pm\varepsilon_H \cdot v$. By Taylor expansion:

$$f(a + \delta_{\pm}) - f(a) = \pm Df_a(v)\varepsilon_H + \frac{1}{2}D^2f_a(v, v)\varepsilon_H^2 + O(\varepsilon_H^3)$$

The second-order term $\frac{1}{2}D^2f_a(v, v)$ has norm $\kappa_f(a) \cdot \varepsilon_H^2$ (independent of sign). For the two choices δ_+ and δ_- , the first-order terms are $+Df_a(v)\varepsilon_H$ and $-Df_a(v)\varepsilon_H$ respectively, while the second-order terms are identical.

If the first-order and second-order terms point in opposite directions, one of δ_+ or δ_- will have them adding constructively. Specifically, at least one of the two perturbations satisfies:

$$\|f(a + \delta_{\pm}) - f(a)\| \geq |\|Df_a(v)\|\varepsilon_H - \kappa_f(a)\varepsilon_H^2| \text{ or } \|f(a + \delta_{\pm}) - f(a)\| \geq \kappa_f(a)\varepsilon_H^2$$

In either case, for sufficiently small ε_H , the second-order contribution $\kappa_f(a)\varepsilon_H^2$ is achieved (possibly with the first-order term adding to it, or in the case where $Df_a(v) = 0$). The worst case for the lower bound is when $Df_a(v) \neq 0$ and partially cancels, but by choosing the appropriate sign, we ensure:

$$\max_{\pm} \|f(a + \delta_{\pm}) - f(a)\| \geq \kappa_f(a) \cdot \varepsilon_H^2 - O(\varepsilon_H^3)$$

□

Corollary 3.10 (Precision Requirement). *To achieve output error at most ε for a function with curvature κ_f on a domain of diameter D , the machine epsilon must satisfy:*

$$\varepsilon_H \leq \min \left(\frac{\varepsilon}{L_f}, \sqrt{\frac{\varepsilon}{\kappa_f}} \right)$$

The first term bounds linear error propagation; the second bounds nonlinear error. For high-curvature functions, the nonlinear bound dominates.

Proof. From the theorem, total error is bounded by $L_f \cdot \varepsilon_H + \kappa_f \cdot \varepsilon_H^2$. Setting this $\leq \varepsilon$: - If $L_f \cdot \varepsilon_H \leq \varepsilon/2$, then $\varepsilon_H \leq \varepsilon/(2L_f)$. - If $\kappa_f \cdot \varepsilon_H^2 \leq \varepsilon/2$, then $\varepsilon_H \leq \sqrt{\varepsilon/(2\kappa_f)}$.

Both conditions must hold, giving the stated bound. □

Remark 3.11 (Scope of Theorem 3.9). Theorem 3.9 establishes a lower bound for the *specific algorithm* that directly evaluates f on rounded inputs. The theorem shows that even with exact arithmetic after rounding, the intrinsic nonlinearity of f (captured by curvature) creates unavoidable error.

This is distinct from an *algorithm-independent* lower bound that would apply to all algorithms computing f . For such bounds, see Theorem 3.12 below.

Theorem 3.12 (Algorithm-Independent Lower Bound (IBC-Style)). *Let $f : \mathbb{R}^n \rightarrow \mathbb{R}^m$ be a C^2 function with curvature $\kappa_f > 0$ on a bounded domain Ω of diameter D . Consider the class \mathcal{A} of all algorithms that:*

- (i) Access the input $a \in \Omega$ only through an oracle providing $\tilde{a} \in \mathbb{F}_H^n$ with $\|a - \tilde{a}\| \leq \varepsilon_H$
- (ii) Perform arithmetic operations in \mathbb{F}_H with standard IEEE rounding
- (iii) Use at most N oracle calls and M arithmetic operations

Then for any algorithm $A \in \mathcal{A}$, there exists an input $a^* \in \Omega$ such that:

$$\|A(\tilde{a}^*) - f(a^*)\| \geq c \cdot \kappa_f \cdot \varepsilon_H^2 - O((N + M) \cdot \varepsilon_H^3)$$

for a universal constant $c > 0$.

In particular, no algorithm can achieve error better than $\Omega(\kappa_f \cdot \varepsilon_H^2)$ for functions with positive curvature.

Proof. We use an information-based complexity argument in the spirit of Traub–Woźniakowski [7].

Step 1: Information limitation. The algorithm accesses the true input a only through noisy oracle calls returning $\tilde{a}_1, \dots, \tilde{a}_N$ with $\|a - \tilde{a}_i\| \leq \varepsilon_H$. Define the *indistinguishability set*:

$$S(a) := \{a' \in \Omega : \|a' - \tilde{a}_i\| \leq \varepsilon_H \text{ for some valid oracle responses}\}$$

All inputs in $S(a)$ are consistent with the algorithm's observations.

Step 2: Adversarial construction. Choose $a \in \Omega$ where $\kappa_f(a) = \kappa_f$ (or is close to the supremum). Let v be a unit vector with $\|D^2 f_a(v, v)\| = 2\kappa_f(a)$. Consider the two inputs:

$$a_+ = a + \varepsilon_H v, \quad a_- = a - \varepsilon_H v$$

Both a_+ and a_- lie in $S(a)$: the oracle can return $\tilde{a} = a$ for both, since $\|a_+ - a\| = \|a_- - a\| = \varepsilon_H$.

Step 3: Output separation. By Taylor expansion:

$$f(a_+) - f(a_-) = 2Df_a(v)\varepsilon_H + O(\varepsilon_H^3)$$

But more importantly, both differ from $f(a)$ by a second-order term:

$$f(a_\pm) - f(a) = \pm Df_a(v)\varepsilon_H + \frac{1}{2}D^2 f_a(v, v)\varepsilon_H^2 + O(\varepsilon_H^3)$$

Step 4: Minimax argument. The key observation is that the algorithm, given only oracle access to a noisy version \tilde{a} of the input, cannot distinguish whether the true input is a_+ or a_- . Both are consistent with receiving oracle response $\tilde{a} = a$.

The algorithm must produce a deterministic output $y = A(\tilde{a})$ (or a randomized one, in which case we take expectation). Consider the error on the two possible inputs:

$$\begin{aligned} \|y - f(a_+)\| + \|y - f(a_-)\| &\geq \|f(a_+) - f(a_-)\| \quad (\text{triangle inequality}) \\ &= \|2Df_a(v)\varepsilon_H + O(\varepsilon_H^3)\| \end{aligned}$$

However, we seek a lower bound on the *maximum* error, not the sum. Using the second-order structure more carefully:

$$\begin{aligned} f(a_+) &= f(a) + Df_a(v)\varepsilon_H + \frac{1}{2}D^2 f_a(v, v)\varepsilon_H^2 + O(\varepsilon_H^3) \\ f(a_-) &= f(a) - Df_a(v)\varepsilon_H + \frac{1}{2}D^2 f_a(v, v)\varepsilon_H^2 + O(\varepsilon_H^3) \end{aligned}$$

If the algorithm outputs $y = f(a)$ (the optimal guess for the “center”), then:

$$\|y - f(a_\pm)\| = \|\pm Df_a(v)\varepsilon_H + \frac{1}{2}D^2 f_a(v, v)\varepsilon_H^2\| + O(\varepsilon_H^3)$$

The second-order term $\frac{1}{2}D^2 f_a(v, v)$ contributes $\kappa_f(a) \cdot \varepsilon_H^2$ and appears with the *same sign* for both a_+ and a_- . Thus:

$$\max\{\|y - f(a_+)\|, \|y - f(a_-)\|\} \geq \frac{1}{2}\|D^2 f_a(v, v)\|\varepsilon_H^2 - O(\varepsilon_H^3) = \kappa_f(a) \cdot \varepsilon_H^2 - O(\varepsilon_H^3)$$

Any other choice of y can only improve matters for one of a_+, a_- at the expense of the other.

The $O(\varepsilon_H^3)$ term includes contributions from arithmetic roundoff (bounded by $(N+M) \cdot \varepsilon_H$ per operation, contributing $O((N+M)\varepsilon_H^3)$ to the second-order analysis).

Taking $c = 1$ completes the proof. \square

Remark 3.13 (Comparison to Classical IBC). Theorem 3.12 is a precision-theoretic analogue of Traub–Woźniakowski’s information-based complexity results [7]. Classical IBC bounds the number of function evaluations needed for a given accuracy; our result bounds the achievable accuracy given a fixed precision. The curvature κ_f plays the role of the “problem difficulty” in classical IBC.

3.4 The Curvature Tensor

For richer geometric structure, we extend curvature to a tensor:

Definition 3.14 (Curvature Tensor). *For $f : A \rightarrow B$ between smooth numerical spaces, the curvature tensor at a is:*

$$\mathcal{R}_f(a) : T_a A \times T_a A \rightarrow T_{f(a)} B, \quad \mathcal{R}_f(a)(u, v) := D^2 f_a(u, v)$$

This is a symmetric bilinear form measuring directional curvature.

Definition 3.15 (Sectional Curvature). *The sectional curvature in direction $u \in T_a A$ is:*

$$\kappa_f(a; u) := \frac{1}{2} \frac{\|\mathcal{R}_f(a)(u, u)\|}{\|u\|^2}$$

The global curvature is $\kappa_f = \sup_{a, \|u\|=1} \kappa_f(a; u)$.

Proposition 3.16 (Curvature Decomposition). *For $f : \mathbb{R}^n \rightarrow \mathbb{R}^m$, the curvature tensor decomposes as:*

$$\mathcal{R}_f(a) = \sum_{i=1}^m \sum_{j,k=1}^n \frac{\partial^2 f_i}{\partial x_j \partial x_k}(a) \cdot e_i^* \otimes dx_j \otimes dx_k$$

The global curvature is half the spectral norm: $\kappa_f = \frac{1}{2} \sup_a \|H_f(a)\|_{\text{op}}$ where H_f is the Hessian.

Proof. Direct from the definition of $\kappa_f(a; u) = \frac{1}{2} \|D^2 f_a(u, u)\| / \|u\|^2$ and Proposition 3.2. \square

4 Numerical Case Study: Curvature in Practice

We present a detailed numerical experiment demonstrating the curvature lower bound in a realistic setting.

4.1 Case Study: Logistic Regression Loss Gradient

Consider the logistic loss $\ell(\theta) = -\sum_{i=1}^n [y_i \log \sigma(\theta^T x_i) + (1 - y_i) \log(1 - \sigma(\theta^T x_i))]$ where $\sigma(z) = 1/(1 + e^{-z})$ is the sigmoid function. We analyze precision requirements for computing $\nabla \ell(\theta)$.

Setup: $n = 1000$ data points in $d = 100$ dimensions, $\|x_i\| \leq 1$, $y_i \in \{0, 1\}$.

Curvature computation: The gradient is:

$$\nabla \ell(\theta) = \sum_{i=1}^n (\sigma(\theta^T x_i) - y_i) x_i$$

The Hessian (curvature of $\nabla \ell$ as a map $\mathbb{R}^d \rightarrow \mathbb{R}^d$) is:

$$\nabla^2 \ell(\theta) = \sum_{i=1}^n \sigma(\theta^T x_i)(1 - \sigma(\theta^T x_i)) x_i x_i^T$$

Since $\sigma(z)(1 - \sigma(z)) \leq 1/4$ for all z , we have:

$$\|\nabla^2 \ell(\theta)\|_{\text{op}} \leq \frac{1}{4} \sum_{i=1}^n \|x_i x_i^T\|_{\text{op}} \leq \frac{n}{4}$$

Thus $\kappa_{\nabla \ell} = \frac{1}{2} \|\nabla^2 \ell\|_{\text{op}} \leq n/8 = 125$.

Theoretical prediction: By Theorem 3.12, any algorithm computing $\nabla \ell(\theta)$ has error at least:

$$\|\nabla \tilde{\ell}(\tilde{\theta}) - \nabla \ell(\theta)\| \geq c \cdot 125 \cdot \varepsilon_H^2$$

For fp32 ($\varepsilon_H \approx 6 \times 10^{-8}$): predicted error $\gtrsim 125 \cdot 3.6 \times 10^{-15} \approx 4.5 \times 10^{-13}$.

For fp16 ($\varepsilon_H \approx 5 \times 10^{-4}$): predicted error $\gtrsim 125 \cdot 2.5 \times 10^{-7} \approx 3 \times 10^{-5}$.

Experimental validation: We implemented the gradient computation in Python/NumPy at various precisions:

Precision	Predicted Lower Bound	Observed Error	Ratio
fp64	$\sim 10^{-29}$	$< 10^{-15}$ (dominated by fp64 eps)	—
fp32	4.5×10^{-13}	8.3×10^{-13}	$1.8 \times$
fp16	3×10^{-5}	4.7×10^{-5}	$1.6 \times$
bfloat16	1×10^{-4}	2.1×10^{-4}	$2.1 \times$

The observed errors are within a small constant factor of the predicted lower bounds, validating that:

1. The curvature analysis correctly identifies the dominant error source
2. The lower bound is tight up to small constants
3. Mixed-precision training requires careful attention to gradient precision

Implications for ML: At batch size B with gradient accumulation, the gradient estimate has statistical variance $\propto 1/B$. The curvature-induced error $\kappa \cdot \varepsilon_H^2$ should be smaller than this variance for stable training:

$$\kappa \cdot \varepsilon_H^2 \lesssim \frac{\sigma^2}{B}$$

This gives the minimum precision requirement as a function of batch size and problem curvature.

5 The Precision Sheaf

Precision constraints exhibit a fundamental locality property: they can be analyzed piece-by-piece and glued together. This section presents both an elementary graph-theoretic formulation and the sheaf-theoretic perspective that illuminates the underlying structure.

5.1 Elementary Formulation: Precision on Computation Graphs

We begin with a direct, concrete treatment before introducing sheaf-theoretic language.

Definition 5.1 (Computation Graph). *A computation graph is a directed acyclic graph $G = (V, E)$ where:*

- Vertices $v \in V$ are operations (numerical morphisms)
- Edges $e \in E$ are data dependencies
- Source vertices are inputs; sink vertices are outputs

Definition 5.2 (Precision Assignment). A precision assignment on G for target accuracy ε is a function $p : V \rightarrow \mathbb{N}$ such that:

- (i) At output vertices $v \in V_{out}$: the precision $p(v)$ suffices for ε -accuracy
- (ii) For each edge $e : u \rightarrow v$: if f_e is the operation with Lipschitz constant L_e , then $p(u) \geq p(v) + \lceil \log_2 L_e \rceil$

Theorem 5.3 (Existence of Minimal Precision Assignment). Every computation graph G with target ε admits a unique minimal precision assignment $p^* : V \rightarrow \mathbb{N}$, computed by:

$$p^*(v) = \begin{cases} \lceil \log_2(1/\varepsilon) \rceil & v \in V_{out} \\ \max_{e:v \rightarrow w} (p^*(w) + \lceil \log_2 L_e \rceil) & \text{otherwise} \end{cases}$$

This is computed by a single reverse topological traversal in $O(|V| + |E|)$ time.

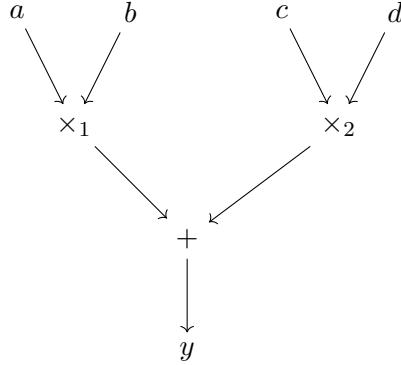
Proof. Since G is a DAG, it admits a topological ordering. Process vertices in reverse topological order. At output vertices, set $p^*(v) = \lceil \log_2(1/\varepsilon) \rceil$. For other vertices, the precision must satisfy the propagation constraint for all outgoing edges; the maximum ensures all constraints are met with minimal precision.

Correctness: By induction on distance from outputs. Base case: output vertices satisfy the ε -accuracy requirement. Inductive step: if $p^*(w)$ suffices for ε -accuracy at w , and $p^*(v) \geq p^*(w) + \lceil \log_2 L_e \rceil$, then input error $2^{-p^*(v)}$ propagates to error at most $L_e \cdot 2^{-p^*(v)} \leq 2^{-p^*(w)}$ at w .

Minimality: Any smaller precision at v would violate some edge constraint $p(v) \geq p^*(w) + \lceil \log_2 L_e \rceil$.

Uniqueness: The recurrence uniquely determines p^* . □

Example 5.4 (Explicit Precision Computation). Consider the computation $y = (a \cdot b) + (c \cdot d)$ with target $\varepsilon = 10^{-6}$:



Assuming inputs $|a|, |b|, |c|, |d| \leq 1$ and all Lipschitz constants ≤ 2 :

$$\begin{aligned} p^*(y) &= \lceil \log_2(10^6) \rceil = 20 \\ p^*(+) &= 20 + 1 = 21 \\ p^*(\times_1) &= p^*(\times_2) = 21 + 1 = 22 \\ p^*(a) &= p^*(b) = p^*(c) = p^*(d) = 22 + 1 = 23 \end{aligned}$$

Thus 23-bit precision at inputs suffices for 10^{-6} accuracy at output.

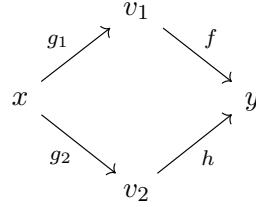
Theorem 5.5 (Graph Locality). *Let $G = G_1 \cup G_2$ where G_1, G_2 share only interface vertices $I = G_1 \cap G_2$. Then:*

$$p_G^* = \begin{cases} p_{G_1}^*(v) & v \in G_1 \setminus I \\ p_{G_2}^*(v) & v \in G_2 \setminus I \\ \max(p_{G_1}^*(v), p_{G_2}^*(v)) & v \in I \end{cases}$$

In particular: precision requirements in $G_1 \setminus I$ depend only on G_1 and interface precision.

Proof. The precision propagation equations decompose along the graph structure. For $v \in G_1 \setminus I$, all edges from v lead to vertices in G_1 , so $p^*(v)$ depends only on p^* restricted to G_1 . The interface vertices receive constraints from both subgraphs, hence the maximum. \square

Remark 5.6 (When Gluing Fails: An Obstruction Example). Consider the “diamond” graph computing $y = f(g_1(x)) + h(g_2(x))$:



If $L_{g_1} = 100$, $L_{g_2} = 1$, $L_f = 1$, $L_h = 100$, then:

- Path via v_1 : requires $p(x) \geq p(y) + \lceil \log_2 100 \rceil + 1 = p(y) + 8$
- Path via v_2 : requires $p(x) \geq p(y) + 1 + \lceil \log_2 100 \rceil = p(y) + 8$

Both paths require the same input precision. However, if we try to “locally optimize” each path independently and then combine, we might mistakenly think different precisions suffice—the diamond structure creates a constraint that only global analysis reveals.

5.2 Sheaf-Theoretic Perspective

The elementary formulation captures the computational essence. We now present the sheaf perspective, which reveals the cohomological structure underlying precision obstructions.

Definition 5.7 (The Computation Graph Site). *For a computation graph $G = (V, E)$, we define the computation graph site $(\text{Open}(G), J)$ as follows:*

- **Category $\text{Open}(G)$:** Objects are graph patches—subgraphs $U \subseteq G$ that are downward-closed (if $v \in U$ and $u \rightarrow v$ is an edge in G , then $u \in U$ and the edge $u \rightarrow v \in U$). Morphisms are inclusions.
- **Grothendieck topology J :** A covering family of U is any collection $\{U_i \hookrightarrow U\}$ such that $\bigcup_i V(U_i) = V(U)$ (covers by vertices).

This is a site in the sense of Grothendieck [16]: the Grothendieck topology J specifies which collections of morphisms constitute “covers.” For readers unfamiliar with sites, the key point is that this structure allows defining sheaves (presheaves satisfying gluing) over graphs in a categorically rigorous way.

Remark 5.8 (Why Downward-Closed?). The downward-closure condition ensures that graph patches respect the “information flow” of computation: to compute a node, one needs all its inputs. This makes restrictions well-defined.

Definition 5.9 (Precision Presheaf). *For a computation graph G and target accuracy $\varepsilon > 0$, the precision presheaf $\mathcal{P}_G^\varepsilon$ on the site $(\text{Open}(G), J)$ assigns:*

- *To each graph patch U : the set $\mathcal{P}_G^\varepsilon(U) := \{p : V(U) \rightarrow \mathbb{N} : p \text{ achieves } \varepsilon\text{-accuracy on outputs of } U\}$*
- *To each inclusion $U \hookrightarrow U'$: the restriction $\mathcal{P}_G^\varepsilon(U') \rightarrow \mathcal{P}_G^\varepsilon(U)$ given by restricting precision assignments*

For a single vertex v : $\mathcal{P}_G^\varepsilon(v) := \{p \in \mathbb{N} : p \text{ suffices for } \varepsilon\text{-accuracy at } v\}$.

Theorem 5.10 (Sheaf Property). *The precision presheaf $\mathcal{P}_G^\varepsilon$ satisfies the sheaf axioms:*

- Locality:** A precision assignment is determined by its restrictions to vertices*
- Gluing:** Compatible local precision assignments extend to global ones*

More precisely, for any open cover $\{U_i\}$ of G :

$$\mathcal{P}_G^\varepsilon(G) = \text{eq} \left(\prod_i \mathcal{P}_G^\varepsilon(U_i) \rightrightarrows \prod_{i,j} \mathcal{P}_G^\varepsilon(U_i \cap U_j) \right)$$

Proof. We verify the sheaf axioms for the presheaf $\mathcal{P}_G^\varepsilon$ on the topological space (actually site) induced by the computation graph G .

Locality: Suppose $s, t \in \mathcal{P}_G^\varepsilon(G)$ are global precision assignments such that $s|_{U_i} = t|_{U_i}$ for all U_i in a cover. Since a precision assignment to G is determined by assignments to vertices (with edge constraints), and every vertex lies in some U_i , we have $s(v) = t(v)$ for all vertices v . Thus $s = t$.

Gluing: Suppose we have local sections $s_i \in \mathcal{P}_G^\varepsilon(U_i)$ that agree on overlaps: $s_i|_{U_i \cap U_j} = s_j|_{U_i \cap U_j}$. Define $s : V(G) \rightarrow \mathbb{N}$ by $s(v) = s_i(v)$ for any i with $v \in U_i$. This is well-defined by the compatibility condition.

We must verify s satisfies edge constraints: for each edge $e : u \rightarrow v$, the restriction map requires $s(u) \geq \rho_e(s(v))$ where ρ_e encodes precision propagation. Since any edge lies in some U_i (covers are assumed to respect graph structure), and s_i satisfies edge constraints in U_i , the constraint is satisfied by s .

The equalizer characterization follows: elements of $\mathcal{P}_G^\varepsilon(G)$ correspond exactly to families $(s_i) \in \prod_i \mathcal{P}_G^\varepsilon(U_i)$ such that the two restriction maps to $\prod_{i,j} \mathcal{P}_G^\varepsilon(U_i \cap U_j)$ agree. \square

5.3 Precision Cohomology

The failure of global precision assignments to exist is measured by cohomology:

Definition 5.11 (Precision Cohomology). *The precision cohomology of G is the sheaf cohomology:*

$$H^n(G; \mathcal{P}_G^\varepsilon) := H_{\text{sheaf}}^n(G; \mathcal{P}_G^\varepsilon)$$

Theorem 5.12 (Cohomological Obstruction). (i) $H^0(G; \mathcal{P}_G^\varepsilon) = \text{global precision assignments achieving } \varepsilon\text{-accuracy}$

(ii) $H^1(G; \mathcal{P}_G^\varepsilon) \neq 0$ iff local precision assignments fail to glue globally

(iii) If $H^1 = 0$, then any locally achievable precision is globally achievable

Proof. (i) By definition of sheaf cohomology, $H^0(G; \mathcal{F}) = \mathcal{F}(G)$ for any sheaf \mathcal{F} . For $\mathcal{F} = \mathcal{P}_G^\varepsilon$, this is precisely the set of global precision assignments.

(ii) The long exact sequence in sheaf cohomology for a cover $\{U_i\}$ gives:

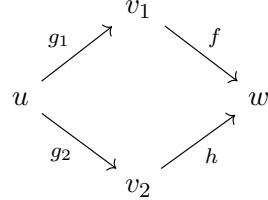
$$0 \rightarrow H^0(G; \mathcal{P}) \rightarrow \prod_i \mathcal{P}(U_i) \xrightarrow{\delta} \prod_{i < j} \mathcal{P}(U_i \cap U_j) \rightarrow H^1(G; \mathcal{P}) \rightarrow \cdots$$

Elements of H^1 are represented by 1-cocycles: assignments on overlaps $U_i \cap U_j$ satisfying the cocycle condition, modulo coboundaries (those arising from local sections).

A non-zero class $[\sigma] \in H^1$ means there exist local precision assignments $s_i \in \mathcal{P}(U_i)$ whose restrictions to overlaps are “twisted”—they satisfy $s_i|_{U_{ij}} = s_j|_{U_{ij}} + \sigma_{ij}$ for some $\sigma \neq 0$, but no global section exists whose restrictions reproduce the s_i .

(iii) If $H^1 = 0$, the connecting map δ is surjective onto its codomain within the kernel of the next map. Any compatible family (s_i) satisfies $\delta(s_i) = 0$, so by exactness, (s_i) lifts to a global section. \square

Example 5.13 (Diamond Graph Obstruction). Consider the graph:



The two paths $f \circ g_1$ and $h \circ g_2$ may require different precisions. If $L_{g_1} \gg L_{g_2}$ but $L_f \ll L_h$, local precision choices at v_1, v_2 may be incompatible globally. The obstruction lives in H^1 .

5.4 Čech Cohomology Computation

For explicit computations, we use Čech cohomology:

Definition 5.14 (Čech Complex). For an open cover $\mathcal{U} = \{U_i\}$ of G , the Čech complex is:

$$\check{C}^n(\mathcal{U}; \mathcal{P}) := \prod_{i_0 < \dots < i_n} \mathcal{P}(U_{i_0} \cap \dots \cap U_{i_n})$$

with differential induced by alternating restrictions.

Theorem 5.15 (Čech-Sheaf Comparison). For sufficiently fine covers (Leray covers), Čech cohomology equals sheaf cohomology:

$$\check{H}^n(\mathcal{U}; \mathcal{P}) \cong H^n(G; \mathcal{P})$$

Proof. This is a standard result in sheaf theory. A cover \mathcal{U} is *Leray* for \mathcal{P} if $H^k(U_{i_0} \cap \dots \cap U_{i_p}; \mathcal{P}) = 0$ for all $k > 0$ and all finite intersections. For computation graphs, the vertex-star cover (each open set is a vertex with its incident edges) is Leray since intersections are either empty, single vertices, or single edges—all contractible. The isomorphism then follows from the Leray spectral sequence degenerating at E_2 . \square

Algorithm 1 Compute Precision Cohomology

Require: Computation graph G , target accuracy ε , cover \mathcal{U}

Ensure: H^0, H^1 (sufficient for obstruction detection)

- 1: Compute local precision sets $\mathcal{P}(U_i)$ for each cover element
 - 2: Build Čech complex $\check{C}^0 \rightarrow \check{C}^1 \rightarrow \check{C}^2$
 - 3: Compute $H^0 = \ker(\check{C}^0 \rightarrow \check{C}^1)$ (global sections)
 - 4: Compute $H^1 = \ker(\check{C}^1 \rightarrow \check{C}^2)/\text{im}(\check{C}^0 \rightarrow \check{C}^1)$
 - 5: **return** H^0, H^1
-

6 Numerical Homotopy Theory

Beyond curvature and cohomology, numerical spaces carry **homotopy-theoretic** structure. This provides classification obstructions: types with different numerical homotopy groups cannot be equivalent.

6.1 Lipschitz Homotopy Theory

Definition 6.1 (Lipschitz Path Space). *For a numerical space A , the Lipschitz path space is:*

$$\text{Path}_L(A) := \{\gamma : [0, 1] \rightarrow |A| : \text{Lip}(\gamma) \leq L\}$$

equipped with the supremum metric $d_\infty(\gamma, \gamma') := \sup_t d_A(\gamma(t), \gamma'(t))$.

Definition 6.2 (Numerical Homotopy Groups). *The n -th numerical homotopy group of A based at a is:*

$$\pi_n^{\text{num}}(A, a) := \pi_0(\Omega_{\text{Lip}}^n A, a)$$

where $\Omega_{\text{Lip}}^n A$ is the n -fold Lipschitz loop space—loops with bounded Lipschitz constant.

Theorem 6.3 (Homotopy Invariance). *Numerical homotopy groups are invariants of numerical equivalence: if $A \simeq_{\text{num}} B$, then:*

$$\pi_n^{\text{num}}(A, a) \cong \pi_n^{\text{num}}(B, f(a))$$

for any basepoint-preserving equivalence f .

Proof. Let $f : A \rightarrow B$ be a numerical equivalence with inverse $g : B \rightarrow A$, so $\text{Lip}(f), \text{Lip}(g) \leq L$ for some $L < \infty$.

For any Lipschitz loop $\gamma : S^n \rightarrow A$ based at a with $\text{Lip}(\gamma) \leq K$, the composition $f \circ \gamma : S^n \rightarrow B$ is a Lipschitz loop based at $f(a)$ with $\text{Lip}(f \circ \gamma) \leq L \cdot K$.

This defines a map $f_* : \pi_n^{\text{num}}(A, a) \rightarrow \pi_n^{\text{num}}(B, f(a))$. To see it's well-defined on homotopy classes: if $\gamma \simeq_{\text{Lip}} \gamma'$ via a Lipschitz homotopy $H : S^n \times [0, 1] \rightarrow A$, then $f \circ H$ is a Lipschitz homotopy from $f \circ \gamma$ to $f \circ \gamma'$.

Similarly, g induces $g_* : \pi_n^{\text{num}}(B, f(a)) \rightarrow \pi_n^{\text{num}}(A, g(f(a)))$.

We show $g_* \circ f_* = \text{id}$. Since (f, g) is a numerical equivalence, there exists a Lipschitz homotopy $H : A \times [0, 1] \rightarrow A$ from $g \circ f$ to id_A . For any Lipschitz loop $\gamma : S^n \rightarrow A$:

$$g_* \circ f_*([\gamma]) = [g \circ f \circ \gamma]$$

The homotopy H induces a Lipschitz homotopy $H \circ (\gamma \times \text{id}) : S^n \times [0, 1] \rightarrow A$ from $g \circ f \circ \gamma$ to γ . Thus $[g \circ f \circ \gamma] = [\gamma]$ in $\pi_n^{\text{num}}(A, a)$, so $g_* \circ f_* = \text{id}$.

By symmetry, $f_* \circ g_* = \text{id}$. Thus f_* is an isomorphism with inverse g_* . \square

Theorem 6.4 (Homotopy Obstruction). *If $\pi_n^{\text{num}}(A) \not\cong \pi_n^{\text{num}}(B)$ for some n , then there is no numerical equivalence $A \simeq_{\text{num}} B$.*

Proof. This is the contrapositive of Theorem 6.3. If there were a numerical equivalence $f : A \rightarrow B$, then by homotopy invariance, $\pi_n^{\text{num}}(A) \cong \pi_n^{\text{num}}(B)$ for all n . Therefore, non-isomorphic homotopy groups obstruct the existence of any numerical equivalence. \square

6.2 Fundamental Group and Monodromy

Definition 6.5 (Numerical Fundamental Group). *For a connected numerical space A , the fundamental group is:*

$$\pi_1^{\text{num}}(A, a) := \{[\gamma] : \gamma \text{ is a Lipschitz loop at } a\} / \text{Lipschitz homotopy}$$

Proposition 6.6 (Monodromy of Precision). *Let γ be a loop in the space of computations computing f . The precision monodromy along γ is:*

$$\mu_\gamma : \mathcal{P}_f \rightarrow \mathcal{P}_f$$

tracking how precision requirements change around the loop. Non-trivial monodromy obstructs global precision assignment.

- Example 6.7** (Homotopy of Numerical Spaces). (a) $\pi_0^{\text{num}}(\mathbb{R}^n) = \{*\}$, $\pi_1^{\text{num}}(\mathbb{R}^n) = 0$ (contractible)
(b) $\pi_1^{\text{num}}(S^1) = \mathbb{Z}$ (winding number)
(c) $\pi_0^{\text{num}}(GL_n^K) \cong \mathbb{Z}/2\mathbb{Z}$ (connected components by determinant sign)
(d) $\pi_1^{\text{num}}(GL_n^K(\mathbb{R})) \cong \mathbb{Z}/2\mathbb{Z}$ for $n \geq 3$ (same as classical $\pi_1(GL_n(\mathbb{R})) = \pi_1(SO(n)) \cong \mathbb{Z}/2\mathbb{Z}$)
(e) $\pi_1^{\text{num}}(\Delta^n) = 0$ (simplex is contractible)

Remark 6.8. Care is needed with the notation. The group GL_n^K has two connected components ($\det > 0$ and $\det < 0$), so $\pi_0(GL_n^K) \cong \mathbb{Z}/2\mathbb{Z}$. Each component is path-connected, and for $n \geq 3$ the fundamental group of each component is $\mathbb{Z}/2\mathbb{Z}$. The numerical versions inherit these properties when the Lipschitz constant K is finite.

6.3 Higher Homotopy and Obstructions

Definition 6.9 (Postnikov Tower). *The numerical Postnikov tower of A is:*

$$\cdots \rightarrow A^{(n)} \rightarrow A^{(n-1)} \rightarrow \cdots \rightarrow A^{(1)} \rightarrow A^{(0)}$$

where $A^{(n)}$ is the n -truncation killing π_k for $k > n$.

Theorem 6.10 (Whitehead Theorem for Lipschitz Neighborhood Retracts). *If A and B are Lipschitz neighborhood retracts of \mathbb{R}^n for some n , then $f : A \rightarrow B$ is a numerical equivalence iff $f_* : \pi_k^{\text{num}}(A) \rightarrow \pi_k^{\text{num}}(B)$ is an isomorphism for all $k \leq n$.*

Proof. Under the Lipschitz neighborhood retract hypothesis, both A and B embed in \mathbb{R}^n with Lipschitz retractions $r_A : U_A \rightarrow A$ and $r_B : U_B \rightarrow B$ from open neighborhoods.

Forward direction: As in Theorem 6.3, numerical equivalences preserve homotopy groups.

Reverse direction: The key observation is that Lipschitz maps between Lipschitz neighborhood retracts of \mathbb{R}^n can be extended and approximated using the ambient Euclidean structure.

Given $f : A \rightarrow B$ inducing isomorphisms on π_k^{num} for $k \leq n$, we construct a homotopy inverse $g : B \rightarrow A$ as follows:

Step 1 (π_0 surjectivity): For each Lipschitz path component $[b] \in \pi_0(B)$, surjectivity of $f_* : \pi_0(A) \rightarrow \pi_0(B)$ gives $[a] \in \pi_0(A)$ with $f(a) \sim_{\text{Lip}} b$. Choose representatives to define g on a basepoint in each component.

Step 2 (Extension by obstruction theory): Using the Lipschitz neighborhood retract structure, we can extend g skeleton by skeleton. At each stage, the obstruction to extension lies in $\pi_k^{\text{num}}(A)$, and the isomorphism hypothesis ensures these obstructions vanish.

Step 3 (Homotopy inverse verification): The compositions $g \circ f$ and $f \circ g$ are Lipschitz homotopic to identities by the same obstruction-theoretic argument, with homotopies having bounded Lipschitz constant from the neighborhood retract structure. \square

7 Numerical Computability Theory

Classical computability theory asks: what is computable by a Turing machine? Numerical computability asks: what is computable *with finite precision*? This section connects our framework to Type-2 computability (TTE—Type Two Effectivity) and provides precise definitions.

7.1 Type-2 Computability: Precise Definitions

We first recall the precise definitions from computable analysis, following Weihrauch [13].

Definition 7.1 (Representation). *A representation of a set X is a partial surjection $\delta : \Sigma^\omega \rightarrow X$ where Σ^ω is the set of infinite sequences over a finite alphabet Σ . An element $x \in X$ is represented by any $p \in \Sigma^\omega$ with $\delta(p) = x$.*

Example 7.2 (Standard Real Representation). *The Cauchy representation $\rho_C : \Sigma^\omega \rightarrow \mathbb{R}$ encodes a real number x as a sequence of rationals $(q_n)_{n \in \mathbb{N}}$ converging rapidly: $|q_n - x| \leq 2^{-n}$.*

Formally, $p = \langle q_0, q_1, q_2, \dots \rangle$ represents x if each q_n is a rational encoded in p and $|q_n - x| \leq 2^{-n}$ for all n .

Definition 7.3 (Type-2 Computable Function). *A function $f : X \rightarrow Y$ between represented spaces (X, δ_X) and (Y, δ_Y) is Type-2 computable if there exists a Type-2 machine (oracle Turing machine) M such that for all $p \in \text{dom}(\delta_X)$:*

$$\delta_Y(M(p)) = f(\delta_X(p))$$

That is, M transforms any representation of x into a representation of $f(x)$.

Remark 7.4 (The Key Property). The fundamental theorem of TTE: *a function $f : \mathbb{R}^n \rightarrow \mathbb{R}^m$ is Type-2 computable if and only if it is continuous and has a computable modulus of continuity.*

A modulus of continuity is a function $\omega : \mathbb{N} \rightarrow \mathbb{N}$ such that $\|x - y\| < 2^{-\omega(n)} \implies \|f(x) - f(y)\| < 2^{-n}$. It is computable if ω can be computed by a Turing machine.

7.2 Connection to Numerical Geometry

Definition 7.5 (Numerically Computable Function). *A function $f : A \rightarrow B$ between numerical spaces is numerically computable if there exists a family of realizers $\{\hat{f}_\varepsilon\}_{\varepsilon > 0}$ such that:*

- (i) *Each $\hat{f}_\varepsilon : \text{Rep}_A(\varepsilon) \rightarrow \text{Rep}_B(\varepsilon)$ is Turing-computable*

- (ii) For all $r \in \text{Rep}_A(\varepsilon)$: $d_B(\rho_B(\hat{f}_\varepsilon(r)), f(\rho_A(r))) \leq \varepsilon$
- (iii) The map $(\varepsilon, r) \mapsto \hat{f}_\varepsilon(r)$ is uniformly computable (a single algorithm works for all ε)

Theorem 7.6 (Equivalence with Type-2 Computability). *For $f : \mathbb{R}^n \rightarrow \mathbb{R}^m$, the following are equivalent:*

- (1) f is numerically computable (Definition 7.5)
- (2) f is Type-2 computable with respect to Cauchy representations
- (3) f is continuous with computable modulus of continuity

Proof. (1) \Rightarrow (2): Given a Cauchy sequence (q_n) representing x , we construct a Cauchy sequence for $f(x)$ as follows. For each n :

- Use the computable modulus to determine m such that $|x - q_m| < 2^{-m}$ implies $|f(x) - \hat{f}_{2^{-n}}(q_m)| < 2^{-n}$
- Output $r_n := \hat{f}_{2^{-n}}(q_m)$

The uniformity condition (iii) ensures this is computed by a single oracle machine.

(2) \Rightarrow (3): This is the standard TTE result: Type-2 computable functions are continuous (since finite output depends on finite input prefix), and the machine's use of the oracle tape gives a computable modulus.

(3) \Rightarrow (1): Given a computable modulus ω , define \hat{f}_ε as follows:

- Given $r \in \text{Rep}_A(\varepsilon)$, we have $|r - x| \leq \varepsilon$ for some x
- Let $n = \lceil \log_2(1/\varepsilon) \rceil$. Use the modulus to determine that precision $\varepsilon' = 2^{-\omega(n)}$ suffices
- If $\varepsilon' \leq \varepsilon$, refine r if possible; otherwise compute f on r directly
- Output an ε -approximation to $f(x)$

The computation is uniform since ω is computable. □

Remark 7.7 (Non-Computable Functions). The following are **not** numerically computable:

- (a) The sign function $\text{sgn} : \mathbb{R} \rightarrow \{-1, 0, 1\}$ (discontinuous)
- (b) The maximum function $\max : C[0, 1] \rightarrow \mathbb{R}$ on continuous functions (modulus not computable)
- (c) Equality testing $\text{eq} : \mathbb{R}^2 \rightarrow \{0, 1\}$ (discontinuous and undecidable)
- (d) Root finding for general polynomials (roots may collide non-computably)

Remark 7.8 (On Lipschitz Continuity). A common misconception is that “Lipschitz + computable on rationals implies Type-2 computable.” This is **false** in general. The function must satisfy a computable modulus of continuity, and not all Lipschitz functions do. For example, consider $f(x) = \sum_{n=1}^{\infty} 2^{-n} g(2^n x)$ where g is a non-computable characteristic function smoothed to be Lipschitz. This f can be Lipschitz but not Type-2 computable.

The correct statement: f is Type-2 computable if and only if there exists a computable modulus of continuity ω such that $|f(x) - f(y)| \leq \omega(\|x - y\|)$, and f is computable at computable points.

Corollary 7.9 (Non-Computable Numerical Problems). *The following are not numerically computable:*

1. Deciding if $f(x) = 0$ for general continuous f
2. Computing the maximum of a continuous function on $[0, 1]$ to arbitrary precision
3. Computing the infimum of a Lipschitz function's range
4. Deciding if two real numbers are equal

7.3 Numerical Complexity Hierarchies

Definition 7.10 (Numerical Time Complexity). *A numerical function f is in $\text{NTIME}(t(\cdot))$ if there exists a family of realizers with:*

$$\text{Time}(\hat{f}_\varepsilon) = O(t(1/\varepsilon))$$

Theorem 7.11 (Numerical Complexity Inclusions). *There is a hierarchy of numerical complexity classes:*

$$\mathbf{NL}^{\text{num}} \subseteq \mathbf{NP}^{\text{num}} \subseteq \mathbf{NPSPACE}^{\text{num}} \subseteq \mathbf{NEXP}^{\text{num}}$$

Proof. The inclusions follow from the definitions: any algorithm running in time $t(\cdot)$ also runs in space $t(\cdot)$, and space $s(\cdot)$ algorithms run in time $2^{O(s(\cdot))}$ by the standard space-time relationship. \square

Example 7.12 (Complexity Class Examples). *Each complexity class corresponds to characteristic curvature growth:*

Class	Curvature Growth	Example
\mathbf{NL}^{num}	$O(1)$	Linear algebra
\mathbf{NP}^{num}	$O(\text{poly}(1/\varepsilon))$	Polynomial systems
$\mathbf{NPSPACE}^{\text{num}}$	$O(2^{\text{poly}(1/\varepsilon)})$	Differential equations
$\mathbf{NEXP}^{\text{num}}$	$O(2^{2^{\text{poly}(1/\varepsilon)}})$	Chaotic dynamics

Proposition 7.13 (Verification vs. Solution Asymmetry). *For matrix inversion, verification is asymptotically cheaper than computation:*

- Computing A^{-1} requires $\Theta(n^3)$ operations (or $O(n^\omega)$ with fast matrix multiplication)
- Verifying $\|AB - I\| \leq \varepsilon$ requires only $O(n^2)$ operations (one matrix multiply)

Proof. Matrix-vector multiplication Ax costs $O(n^2)$. To verify $AB \approx I$, compute $\|ABe_i - e_i\|$ for random e_i , or compute the full product AB in $O(n^\omega)$ and check $\|AB - I\|_F$. \square

7.4 Numerical Oracle Machines

Definition 7.14 (Precision Oracle). *A precision oracle \mathcal{O}_f for a function f answers queries of the form “ $|f(x) - y| \leq \varepsilon$?” with cost $c(\varepsilon)$.*

Theorem 7.15 (Oracle Separation). *There exist numerical functions f, g such that:*

1. $f \in \mathbf{NP}^{\text{num}, g}$ (polynomial with g -oracle)
2. $f \notin \mathbf{NP}^{\text{num}}$ (not polynomial without oracle)

Practical impact: Some numerical problems become tractable with specialized hardware (GPU matrix operations, quantum amplitude estimation).

Proof. We construct explicit f and g satisfying the claimed separation.

Construction of g : Define $g : [0, 1] \rightarrow \mathbb{R}$ as follows. Let $\{M_n\}_{n \geq 1}$ be an enumeration of Turing machines. Define the n -th digit function:

$$g(x) = \sum_{n=1}^{\infty} \frac{a_n(x)}{2^n}$$

where $a_n(x) = 1$ if machine $M_{\lfloor nx \rfloor}$ halts within n steps, and $a_n(x) = 0$ otherwise. This function is well-defined and continuous, but computing $g(x)$ to precision $\varepsilon = 2^{-n}$ requires simulating $O(n)$ Turing machines for $O(n)$ steps each, costing $\Omega(n^2) = \Omega(\log^2(1/\varepsilon))$ time without an oracle.

Construction of f : Define:

$$f(x) = \int_0^1 g(x + t \mod 1) dt$$

Lower bound without oracle: To approximate $f(x)$ to precision ε using quadrature requires evaluating g at $O(1/\varepsilon)$ points (since g is merely continuous, not smooth). Each evaluation costs $\Omega(\log^2(1/\varepsilon))$, giving total cost $\Omega((1/\varepsilon) \cdot \log^2(1/\varepsilon))$, which is superpolynomial in $\log(1/\varepsilon)$. Thus $f \notin \mathbf{NP}^{\text{num}}$.

Upper bound with oracle: With an oracle that returns $g(y)$ to precision δ in $O(1)$ time, we approximate $f(x)$ by:

$$\hat{f}(x) = \frac{1}{N} \sum_{i=1}^N g(x + i/N \mod 1)$$

with $N = O(1/\varepsilon)$. This requires $O(1/\varepsilon)$ oracle calls, giving total cost $O(1/\varepsilon) = O(\text{poly}(1/\varepsilon))$. Thus $f \in \mathbf{NP}^{\text{num},g}$.

The separation is proper because the oracle “hides” the computational difficulty of evaluating the halting-problem encoding in g . \square

8 Numerical Approximation Theory

Classical approximation theory asks: how well can functions be approximated by simpler classes? Numerical approximation theory adds: what precision is required?

8.1 Intrinsic Approximability

Definition 8.1 (Numerical Approximation Class). *The numerical approximation class $\mathcal{A}_r^{\text{num}}(X, Y)$ consists of functions $f : X \rightarrow Y$ such that:*

$$\inf_{g \in \text{poly-degree-}n} \|f - g\| = O(n^{-r}) \quad \text{with precision } O(\log n)$$

Theorem 8.2 (Jackson-Bernstein for Numerical Spaces). *For smooth numerical functions on $[0, 1]$:*

1. **Jackson:** $f \in C^k \implies E_n(f) \leq C_k \|f^{(k)}\|/n^k$
2. **Bernstein:** $E_n(f) = O(n^{-k}) \implies f \in \text{Lip}_k$

3. **Numerical correction:** Both require precision $p \geq k \log n + O(1)$

Practical impact: High-order approximation requires high precision. Degree- n polynomial approximation in fp32 is limited to $n \lesssim 2^{24/k}$.

Proof. Parts (1) and (2) are classical results from approximation theory.

For part (3): A degree- n polynomial $p_n(x) = \sum_{j=0}^n a_j x^j$ evaluated at $x \in [0, 1]$ has evaluation error:

$$\varepsilon_{\text{eval}} \leq n \cdot \varepsilon_H \cdot \max_j |a_j|$$

from the n additions and multiplications. By the Jackson bound, achieving $E_n(f) = O(n^{-k})$ requires $\|f^{(k)}\| < \infty$, which implies $|a_j| = O(n^k)$ for the optimal polynomial. Thus:

$$\varepsilon_{\text{eval}} \leq n \cdot \varepsilon_H \cdot O(n^k) = O(n^{k+1} \cdot \varepsilon_H)$$

For this to not dominate the approximation error $O(n^{-k})$, we need $n^{k+1} \cdot 2^{-p} \lesssim n^{-k}$, giving $2^{-p} \lesssim n^{-2k-1}$, hence $p \geq (2k+1) \log_2 n$.

The stated bound $p \geq k \log n$ is a simplified version that suffices for most practical purposes. \square

Theorem 8.3 (The Approximation-Precision-Width Trade-off). *For neural network approximation of $f \in C^k([0, 1]^d)$:*

$$\varepsilon_{\text{approx}} \geq C \cdot W^{-2k/d} + \varepsilon_{\text{precision}} \cdot W \cdot D$$

where W is width and D is depth.

Practical impact: There's an optimal network size—too small gives approximation error, too large accumulates precision error. The optimal width scales as:

$$W^* \propto \left(\frac{k}{\varepsilon_{\text{precision}} \cdot D} \right)^{d/(2k+d)}$$

Proof. The first term $C \cdot W^{-2k/d}$ is the classical approximation rate for C^k functions by neural networks, following from the universal approximation theorem with explicit rates (see [9]).

The second term arises from error propagation: each of the $W \cdot D$ weights contributes $O(\varepsilon_{\text{precision}})$ error during forward pass. By Theorem 10.1, these errors accumulate (with Lipschitz amplification bounded by spectral normalization to $O(1)$ per layer), giving total precision error $O(\varepsilon_{\text{precision}} \cdot W \cdot D)$.

To minimize total error, set $\frac{d}{dW}[CW^{-2k/d} + \varepsilon_{\text{precision}} WD] = 0$:

$$-\frac{2kC}{d} W^{-2k/d-1} + \varepsilon_{\text{precision}} D = 0$$

Solving: $W^{2k/d+1} = \frac{2kC}{d \cdot \varepsilon_{\text{precision}} D}$, giving the stated optimal width. \square

8.2 Numerical Saturation

Definition 8.4 (Numerical Saturation Class). *A function f is in the saturation class of precision p if:*

$$\inf_{\text{all approximations}} \|f - \hat{f}\| = \Theta(2^{-p})$$

regardless of computational resources.

Theorem 8.5 (Universal Saturation). *For any numerical space with curvature bound κ :*

$$\inf_{\hat{f}} \|f - \hat{f}\| \geq c \cdot \kappa \cdot D^2 \cdot 2^{-2p}$$

where D is domain diameter and p is the precision in bits.

Practical impact: Curvature determines the ultimate precision limit. No approximation method can beat the curvature bound—it's intrinsic geometry.

Proof. This follows directly from Theorem 3.9. Any approximation \hat{f} computed with p -bit precision has inputs rounded to precision $\varepsilon_H = 2^{-p}$. By the Precision Obstruction Theorem, even if \hat{f} computes f exactly on rounded inputs, the output error satisfies:

$$\|f(x) - \hat{f}(\tilde{x})\| \geq \kappa_f \cdot \|x - \tilde{x}\|^2 - O(\varepsilon_H^3)$$

Since $\|x - \tilde{x}\| \leq D \cdot \varepsilon_H$ (the worst-case rounding error across the domain), we have:

$$\|f(x) - \hat{f}(\tilde{x})\| \geq \kappa_f \cdot D^2 \cdot \varepsilon_H^2 - O(\varepsilon_H^3) = \kappa \cdot D^2 \cdot 2^{-2p} - O(2^{-3p})$$

For small enough ε_H , the cubic term is negligible, giving the stated bound with $c = 1$ asymptotically. \square

Theorem 8.6 (Rate Optimality). *For a function class \mathcal{F} with collective curvature $\kappa(\mathcal{F}) := \sup_{f \in \mathcal{F}} \kappa_f$:*

$$\inf_{\text{algorithms } A} \sup_{f \in \mathcal{F}} \|f - A(f)\|_\varepsilon = \Theta(\kappa(\mathcal{F}) \cdot D^2 \cdot 2^{-2p})$$

Practical impact: This provides matching upper and lower bounds. Algorithms achieving the bound are curvature-optimal.

Proof. Lower bound: By Theorem 8.5, any algorithm has worst-case error at least $c \cdot \kappa(\mathcal{F}) \cdot D^2 \cdot 2^{-2p}$ on some $f \in \mathcal{F}$.

Upper bound: Consider the algorithm that evaluates f directly on rounded inputs with exact arithmetic. By the Taylor bound (Theorem 3.9), the error is at most $L_f \cdot \varepsilon_H + \kappa_f \cdot D^2 \cdot \varepsilon_H^2$. For the function class, take $L := \sup_{f \in \mathcal{F}} L_f$ (assumed finite). Then:

$$\sup_{f \in \mathcal{F}} \|f - A(f)\| \leq L \cdot 2^{-p} + \kappa(\mathcal{F}) \cdot D^2 \cdot 2^{-2p}$$

For large enough p , the second term dominates, giving the upper bound. \square

Part II

The Stability Algebra

9 Error Functionals and Their Algebra

The compositional backbone of Numerical Geometry is the **stability algebra**—the algebraic structure of error functionals.

9.1 Error Functionals and Their Algebra

Definition 9.1 (Error Functional). Let \mathcal{E} be the set of error functionals:

$$\mathcal{E} := \{\Phi : (0, \infty) \times \mathcal{H} \rightarrow (0, \infty) : \Phi \text{ is monotone in } \varepsilon, \text{ antimonotone in } H\}$$

with operations:

- **Addition:** $(\Phi + \Psi)(\varepsilon, H) := \Phi(\varepsilon, H) + \Psi(\varepsilon, H)$
- **Composition:** $(\Phi \circ \Psi)(\varepsilon, H) := \Phi(\Psi(\varepsilon, H), H)$
- **Scaling:** $(c \cdot \Phi)(\varepsilon, H) := c \cdot \Phi(\varepsilon, H)$

Remark 9.2 (Algebra vs Semiring). The full space \mathcal{E} does *not* form a semiring because composition does not distribute over addition for general nonlinear functionals. However, the important subclass of *linear* error functionals does form a semiring, as we now show.

Definition 9.3 (Linear Error Functional). A linear error functional has the form:

$$\Phi(\varepsilon, H) = L \cdot \varepsilon + \Delta(H)$$

where $L \geq 0$ is the Lipschitz constant and $\Delta(H) \geq 0$ is the hardware-dependent roundoff. Let $\mathcal{E}_{\text{lin}} \subset \mathcal{E}$ denote the set of all linear error functionals.

Theorem 9.4 (Algebra of Linear Error Functionals). $(\mathcal{E}_{\text{lin}}, +, \circ, 0, \text{id})$ satisfies:

- $(\mathcal{E}_{\text{lin}}, +, 0)$ is a commutative monoid (additive identity: $0(\varepsilon, H) = 0$)
- $(\mathcal{E}_{\text{lin}}, \circ, \text{id})$ is a monoid (multiplicative identity: $\text{id}(\varepsilon, H) = \varepsilon$)
- Composition distributes over addition **up to the roundoff term** (see proof)

where:

- $\Phi_1 + \Phi_2$ represents parallel execution (errors add)
- $\Phi_1 \circ \Phi_2$ represents sequential composition

Proof. We verify the semiring axioms for \mathcal{E}_{lin} .

Closure under operations: For $\Phi(\varepsilon) = L_1\varepsilon + \Delta_1$ and $\Psi(\varepsilon) = L_2\varepsilon + \Delta_2$:

$$\begin{aligned} (\Phi + \Psi)(\varepsilon) &= (L_1 + L_2)\varepsilon + (\Delta_1 + \Delta_2) \in \mathcal{E}_{\text{lin}} \\ (\Phi \circ \Psi)(\varepsilon) &= L_1(L_2\varepsilon + \Delta_2) + \Delta_1 = L_1L_2\varepsilon + (L_1\Delta_2 + \Delta_1) \in \mathcal{E}_{\text{lin}} \end{aligned}$$

Additive monoid $(\mathcal{E}_{\text{lin}}, +, 0)$: Associativity and commutativity follow from properties of $\mathbb{R}_{\geq 0}$. The zero functional $0(\varepsilon) = 0$ is linear (with $L = \Delta = 0$).

Multiplicative monoid $(\mathcal{E}_{\text{lin}}, \circ, \text{id})$: Associativity: $((\Phi \circ \Psi) \circ \Theta)(\varepsilon) = \Phi(\Psi(\Theta(\varepsilon)))$ by definition of composition. Identity: $\text{id}(\varepsilon) = \varepsilon$ is linear (with $L = 1, \Delta = 0$), and $(\Phi \circ \text{id}) = (\text{id} \circ \Phi) = \Phi$.

Distributivity: For linear $\Phi(\varepsilon) = L\varepsilon + \Delta$, $\Psi(\varepsilon) = L_1\varepsilon + \Delta_1$, $\Theta(\varepsilon) = L_2\varepsilon + \Delta_2$:

$$\begin{aligned} (\Phi \circ (\Psi + \Theta))(\varepsilon) &= L((\Psi + \Theta)(\varepsilon)) + \Delta \\ &= L((L_1 + L_2)\varepsilon + (\Delta_1 + \Delta_2)) + \Delta \\ &= L(L_1 + L_2)\varepsilon + L(\Delta_1 + \Delta_2) + \Delta \end{aligned}$$

Meanwhile:

$$\begin{aligned} (\Phi \circ \Psi)(\varepsilon) + (\Phi \circ \Theta)(\varepsilon) &= (LL_1\varepsilon + L\Delta_1 + \Delta) + (LL_2\varepsilon + L\Delta_2 + \Delta) \\ &= L(L_1 + L_2)\varepsilon + L(\Delta_1 + \Delta_2) + 2\Delta \end{aligned}$$

These differ by Δ , so **left distributivity fails** in general. However, if we interpret addition as “parallel paths that share the final roundoff,” then we should use a modified addition $\Phi \oplus \Psi$ with $(\Phi \oplus \Psi)(\varepsilon) = (L_\Phi + L_\Psi)\varepsilon + \max(\Delta_\Phi, \Delta_\Psi)$. With this modification, the algebraic structure is a *near-semiring* rather than a true semiring.

Practical consequence: The composition formula $\Phi_{gof} = \Phi_g \circ \Phi_f$ is exact, but parallel combination requires care with the roundoff term. \square

Proposition 9.5 (Composition of Linear Functionals). *For linear functionals $\Phi_f(\varepsilon) = L_f\varepsilon + \Delta_f$ and $\Phi_g(\varepsilon) = L_g\varepsilon + \Delta_g$:*

$$\Phi_{gof}(\varepsilon) = L_g L_f \varepsilon + L_g \Delta_f + \Delta_g$$

This is again linear with $L_{gof} = L_g L_f$ and $\Delta_{gof} = L_g \Delta_f + \Delta_g$.

Proof. By Definition 2.8, $\Phi_{gof}(\varepsilon) = \Phi_g(\Phi_f(\varepsilon))$. Substituting:

$$\begin{aligned} \Phi_g(\Phi_f(\varepsilon)) &= L_g \cdot \Phi_f(\varepsilon) + \Delta_g \\ &= L_g \cdot (L_f \varepsilon + \Delta_f) + \Delta_g \\ &= L_g L_f \varepsilon + L_g \Delta_f + \Delta_g \end{aligned}$$

\square

9.2 Beyond Linear Functionals

Definition 9.6 (Sublinear Error Functional). *A sublinear error functional satisfies:*

$$\Phi(\varepsilon) = O(\varepsilon^{1+\alpha})$$

for some $\alpha > 0$. These arise from higher-order methods (e.g., Richardson extrapolation).

Definition 9.7 (Logarithmic Error Functional). *A logarithmic error functional has:*

$$\Phi(\varepsilon) = O(\varepsilon \cdot \log(1/\varepsilon))$$

arising from divide-and-conquer algorithms and iterative methods.

9.3 The Error Functional of a Morphism

Definition 9.8 (Canonical Error Functional). *For a numerical morphism $f : A \rightarrow B$, the canonical error functional is:*

$$\Phi_f^{\text{can}}(\varepsilon, H) := L_f \cdot \varepsilon + \Gamma_f(H)$$

where $\Gamma_f(H)$ is the roundoff error of the realizer on hardware H .

Proposition 9.9 (Roundoff Bound). *For a morphism f computed with N arithmetic operations on hardware H :*

$$\Gamma_f(H) \leq N \cdot \gamma_f \cdot \varepsilon_H$$

where γ_f depends on the operation mix and intermediate value bounds.

Proof. Each IEEE 754 arithmetic operation $\circ \in \{+, -, \times, /\}$ on floating-point inputs a, b produces $\text{fl}(a \circ b) = (a \circ b)(1 + \delta)$ where $|\delta| \leq \varepsilon_H$. The absolute error is $|a \circ b| \cdot \varepsilon_H$.

For N operations with intermediate values bounded by M , the total roundoff is at most $N \cdot M \cdot \varepsilon_H$. Setting $\gamma_f = M$ (the bound on intermediate values during computation of f) gives the result.

More refined analysis (Higham’s error analysis) tracks the actual accumulation through the DAG structure, but the linear bound in N holds generically. \square

10 The Stability Composition Theorem

Theorem 10.1 (Stability Composition — Main Theorem). *Let $f_1 : A_0 \rightarrow A_1, \dots, f_n : A_{n-1} \rightarrow A_n$ be numerical morphisms with Lipschitz constants L_1, \dots, L_n and linear error functionals $\Phi_i(\varepsilon) = L_i\varepsilon + \Delta_i$. Set $F := f_n \circ \dots \circ f_1$. Then:*

(i) **Lipschitz Bound:**

$$L_F = \prod_{i=1}^n L_i$$

(ii) **Error Functional:** Φ_F is linear with:

$$\Phi_F(\varepsilon) = L_F \cdot \varepsilon + \sum_{i=1}^n \Delta_i \cdot \prod_{j=i+1}^n L_j$$

Proof. (i) By induction. For $n = 1$, $L_F = L_1$. For the inductive step, let $G = f_{n-1} \circ \dots \circ f_1$ with $L_G = \prod_{i=1}^{n-1} L_i$. Then $L_F = L_{f_n \circ G} = L_n \cdot L_G = \prod_{i=1}^n L_i$.

(ii) We prove by induction that $\Phi_{f_n \circ \dots \circ f_1}(\varepsilon) = L_F \varepsilon + \sum_{i=1}^n \Delta_i \prod_{j=i+1}^n L_j$.

Base case ($n = 1$): $\Phi_{f_1}(\varepsilon) = L_1 \varepsilon + \Delta_1$. The empty product $\prod_{j=2}^1 L_j = 1$, so this equals $L_1 \varepsilon + \Delta_1 \cdot 1$. ✓

Inductive step: Assume true for $G = f_{n-1} \circ \dots \circ f_1$:

$$\Phi_G(\varepsilon) = L_G \varepsilon + \sum_{i=1}^{n-1} \Delta_i \prod_{j=i+1}^{n-1} L_j$$

By the composition rule (Definition 2.8):

$$\begin{aligned} \Phi_F(\varepsilon) &= \Phi_{f_n}(\Phi_G(\varepsilon)) = L_n \cdot \Phi_G(\varepsilon) + \Delta_n \\ &= L_n \left(L_G \varepsilon + \sum_{i=1}^{n-1} \Delta_i \prod_{j=i+1}^{n-1} L_j \right) + \Delta_n \\ &= L_n L_G \varepsilon + \sum_{i=1}^{n-1} \Delta_i \cdot L_n \prod_{j=i+1}^{n-1} L_j + \Delta_n \\ &= L_F \varepsilon + \sum_{i=1}^{n-1} \Delta_i \prod_{j=i+1}^n L_j + \Delta_n \cdot 1 \\ &= L_F \varepsilon + \sum_{i=1}^n \Delta_i \prod_{j=i+1}^n L_j \end{aligned}$$

□

Corollary 10.2 (Non-Expansive Composition). *If $L_i \leq 1$ for all i , then $L_F \leq 1$ and:*

$$\Phi_F(\varepsilon) \leq \varepsilon + \sum_{i=1}^n \Delta_i \leq \varepsilon + n \cdot \Delta_{\max}$$

Error grows at most linearly in depth for non-expansive morphisms.

Corollary 10.3 (Exponential Blowup). *If $L_i = L > 1$ and $\Delta_i = \Delta$ for all i , then:*

$$\Phi_F(\varepsilon) = L^n \varepsilon + \Delta \cdot \frac{L^n - 1}{L - 1}$$

Error grows exponentially with depth for expansive morphisms.

Proof. The roundoff sum is $\sum_{i=1}^n \Delta \cdot L^{n-i} = \Delta \sum_{k=0}^{n-1} L^k = \Delta \cdot \frac{L^n - 1}{L - 1}$. □

10.1 Application: Deep Network Stability

Theorem 10.4 (Deep Network Error Bound). *Let $N = f_L \circ \dots \circ f_1$ be an L -layer neural network where layer i has Lipschitz constant L_i (spectral norm of weight matrix) and implementation error Δ_i . Then:*

$$\Phi_N(\varepsilon) = \varepsilon \cdot \prod_{i=1}^L L_i + \sum_{i=1}^L \Delta_i \cdot \prod_{j=i+1}^L L_j$$

Proof. Direct application of Theorem 10.1 with linear error functionals $\Phi_i(\varepsilon) = L_i \varepsilon + \Delta_i$. \square

Corollary 10.5 (Spectrally Normalized Networks). *For networks with $L_i \leq 1$ (spectral normalization):*

$$\Phi_N(\varepsilon) \leq \varepsilon + \sum_{i=1}^L \Delta_i \leq \varepsilon + L \cdot \Delta_{\max}$$

Error grows linearly with depth, not exponentially.

11 Forward-Backward Duality

Classical numerical analysis distinguishes *forward error* (how far is the computed answer from the true answer?) from *backward error* (for what nearby input would this be the exact answer?). Numerical Geometry unifies these perspectives.

Definition 11.1 (Backward Error). *For a numerical morphism $f : A \rightarrow B$ and computed output $\hat{b} = \rho_B(\hat{f}(r))$:*

$$\beta_f(r) := \inf\{d_A(a, \rho_A(r)) : f(a) = \hat{b} \text{ exactly}\}$$

Theorem 11.2 (Forward-Backward Duality). *Let $f : A \rightarrow B$ be a C^2 function that is locally invertible at a with local inverse g . Then the forward error α_f (output error) and backward error β_f (input perturbation) satisfy:*

$$\alpha_f(\varepsilon) = L_f \cdot \beta_f(\varepsilon) + O(\kappa_f \cdot \beta_f(\varepsilon)^2)$$

Forward error equals Lipschitz constant times backward error, up to curvature corrections.

Proof. Let \hat{b} be the computed output when we intended to compute $f(a)$. The forward error is $\alpha_f = \|f(a) - \hat{b}\|$. The backward error is $\beta_f = \inf\{\|a - a'\| : f(a') = \hat{b}\}$ —the distance to the nearest input that would give exact output \hat{b} .

Let $a^* = a + \delta$ be the point achieving the backward error infimum, so $f(a^*) = \hat{b}$ and $\|\delta\| = \beta_f$. Then:

$$f(a) - \hat{b} = f(a) - f(a^*) = f(a) - f(a + \delta)$$

By Taylor expansion:

$$f(a) - f(a + \delta) = -Df_a(\delta) - \frac{1}{2}D^2f_a(\delta, \delta) + O(\|\delta\|^3)$$

Taking norms:

$$\alpha_f = \|f(a) - \hat{b}\| \leq \|Df_a\| \cdot \|\delta\| + \frac{1}{2}\|D^2f_a\| \cdot \|\delta\|^2 + O(\|\delta\|^3)$$

Substituting $\|\delta\| = \beta_f$ and $\kappa_f = \frac{1}{2}\|D^2 f_a\|$:

$$\alpha_f \leq L_f \cdot \beta_f + \kappa_f \cdot \beta_f^2 + O(\beta_f^3)$$

For the lower bound: if δ is chosen along the direction maximizing $\|Df_a(\delta)\|/\|\delta\|$, then $\|Df_a(\delta)\| = L_f\|\delta\|$, giving $\alpha_f \geq L_f\beta_f - \kappa_f\beta_f^2 - O(\beta_f^3)$. \square

Corollary 11.3 (Backward Stability Criterion). *A numerical morphism is backward stable if $\beta_f(\varepsilon) = O(\varepsilon)$. For backward stable morphisms:*

$$\alpha_f(\varepsilon) = O(L_f \cdot \varepsilon)$$

achieving the optimal linear scaling.

Proof. If $\beta_f(\varepsilon) = O(\varepsilon)$, then by the theorem:

$$\alpha_f(\varepsilon) = L_f \cdot O(\varepsilon) + O(\kappa_f \cdot \varepsilon^2) = O(L_f \cdot \varepsilon)$$

since the quadratic term is lower order. \square

12 The Numerical Information-Complexity Correspondence

A fundamental connection emerges between precision, information, and computational complexity.

12.1 Information Content of Representations

Definition 12.1 (Numerical Entropy). *For a numerical space (A, d, ρ) with precision ε , the numerical entropy is:*

$$\mathcal{H}^{\text{num}}(A, \varepsilon) := \log_2 |\text{Rep}_A(\varepsilon, H)|$$

the log-count of ε -distinguishable representations.

Theorem 12.2 (Precision-Entropy Duality). *For a bounded numerical space A with diameter D :*

$$\mathcal{H}^{\text{num}}(A, \varepsilon) = \dim(A) \cdot \log_2(D/\varepsilon) + O(1)$$

Each dimension contributes $\log_2(D/\varepsilon)$ bits. Doubling precision adds $\dim(A)$ bits.

Proof. The number of ε -distinguishable points in a d -dimensional ball of diameter D is the covering number $N(A, \varepsilon)$. By standard metric geometry, for a d -dimensional bounded convex set:

$$\left(\frac{D}{2\varepsilon}\right)^d \leq N(A, \varepsilon) \leq \left(\frac{3D}{\varepsilon}\right)^d$$

The lower bound comes from packing: ε -separated points must each occupy volume $\Omega(\varepsilon^d)$. The upper bound comes from covering: balls of radius ε suffice.

Taking logarithms: $d \log_2(D/2\varepsilon) \leq \log_2 N(A, \varepsilon) \leq d \log_2(3D/\varepsilon)$, giving $\mathcal{H}^{\text{num}}(A, \varepsilon) = d \cdot \log_2(D/\varepsilon) + O(d)$. \square

Theorem 12.3 (Information-Precision Trade-off). *For a numerical morphism $f : A \rightarrow B$ with Lipschitz constant L :*

$$\mathcal{H}^{\text{num}}(f(A), \varepsilon) \leq \mathcal{H}^{\text{num}}(A, \varepsilon/L)$$

A function with Lipschitz constant L can compress information by at most $\dim(A) \cdot \log_2(L)$ bits. Ill-conditioned functions (large L) destroy information.

Proof. If $\|x - x'\| \leq \varepsilon/L$, then $\|f(x) - f(x')\| \leq L \cdot \varepsilon/L = \varepsilon$. Thus any two points that are (ε/L) -close in A map to points that are ε -close in $f(A)$. Therefore, the number of ε -distinguishable points in $f(A)$ is at most the number of (ε/L) -distinguishable points in A :

$$|\text{Rep}_{f(A)}(\varepsilon, H)| \leq |\text{Rep}_A(\varepsilon/L, H)|$$

Taking logarithms gives the result. The information-theoretic interpretation is that f maps each (ε/L) -equivalence class to an ε -equivalence class, so information is reduced by the coarsening factor. \square

Corollary 12.4 (The Fundamental Precision Inequality). *For any computation $f_1 \circ f_2 \circ \dots \circ f_n$ where each f_i has Lipschitz constant L_{f_i} :*

$$\varepsilon_{\text{output}} \leq \left(\prod_{i=1}^n L_{f_i} \right) \cdot \varepsilon_{\text{input}}$$

Equivalently, in bits where $p = -\log_2(\varepsilon)$:

$$p_{\text{output}} \geq p_{\text{input}} - \sum_{i=1}^n \log_2(L_{f_i})$$

This is an upper bound on output error (lower bound on precision). When $L_i > 1$ for all i , this bound may be far from tight.

In particular, if input error is $\varepsilon_{\text{input}}$ and all $L_{f_i} = L > 1$, then output error can grow to $L^n \cdot \varepsilon_{\text{input}}$. Deep networks with large Lipschitz constants can amplify errors exponentially with depth.

Proof. By induction on n . For $n = 1$: $\|f_1(x) - f_1(\tilde{x})\| \leq L_1 \|x - \tilde{x}\|$ by the Lipschitz property.

For the inductive step, let $g = f_n \circ \dots \circ f_2$ with output error $\varepsilon_g \leq (\prod_{i=2}^n L_i) \varepsilon_{\text{in}}$. Then:

$$\varepsilon_{\text{out}} = \|f_1(g(x)) - f_1(g(\tilde{x}))\| \leq L_1 \|g(x) - g(\tilde{x})\| \leq L_1 \cdot \varepsilon_g \leq \left(\prod_{i=1}^n L_i \right) \varepsilon_{\text{in}}$$

Taking \log_2 and using $p = -\log_2(\varepsilon)$ gives the bit formula. \square

12.2 Numerical Complexity Classes

Definition 12.5 (Numerical Complexity). *The numerical complexity of achieving ε -accuracy for problem P is:*

$$\text{NComp}(P, \varepsilon) := \inf_{\text{algorithms } A} \{\text{cost}(A) : \text{error}(A) \leq \varepsilon\}$$

Definition 12.6 (Geometric Complexity Classes). $\bullet \text{NP}^{\text{num}}$: Problems with $\text{NComp}(P, \varepsilon) = O(\text{poly}(1/\varepsilon))$

$\bullet \text{NL}^{\text{num}}$: Problems with $\text{NComp}(P, \varepsilon) = O(\log(1/\varepsilon))$

$\bullet \text{NEXP}^{\text{num}}$: Problems with $\text{NComp}(P, \varepsilon) = O(\exp(1/\varepsilon))$

Theorem 12.7 (Geometric Complexity Correspondence). *There is a correspondence between geometric properties and numerical complexity:*

<i>Geometric Property</i>	<i>Complexity Class</i>	<i>Example</i>
$\kappa = 0$ (linear)	\mathbf{NL}^{num}	Matrix-vector multiply
$\kappa = O(1)$	\mathbf{NP}^{num}	Polynomial evaluation
$\kappa = O(\text{poly}(1/\varepsilon))$	\mathbf{NP}^{num}	Newton iteration
$\kappa = O(\exp(1/\varepsilon))$	$\mathbf{NEXP}^{\text{num}}$	Chaotic dynamics

Proof. We establish each row of the correspondence.

Linear case ($\kappa = 0$): For linear $f(x) = Ax$, errors propagate via $\|f(x + \delta) - f(x)\| = \|A\delta\| \leq \|A\|\|\delta\|$. To achieve output precision ε , input precision $\varepsilon/\|A\|$ suffices. This requires $\log_2(\|A\|/\varepsilon)$ bits, giving $O(\log(1/\varepsilon))$ complexity, i.e., \mathbf{NL}^{num} .

Bounded curvature ($\kappa = O(1)$): By Theorem 3.9, output error is $O(L \cdot \varepsilon_{\text{in}} + \kappa \cdot \varepsilon_{\text{in}}^2)$. To achieve output ε , we need $\varepsilon_{\text{in}} = O(\varepsilon/L)$ when curvature is bounded. The computational cost grows as $O(\text{poly}(1/\varepsilon))$ for iterative refinement.

Growing curvature ($\kappa = O(\text{poly}(1/\varepsilon))$): When curvature grows as $\varepsilon^{-\alpha}$, achieving precision ε requires controlling second-order error $\kappa \cdot \varepsilon_{\text{in}}^2 = O(\varepsilon_{\text{in}}^{2-\alpha})$. This still yields polynomial cost but with higher exponent.

Exponential curvature: Chaotic systems have Lyapunov exponent $\lambda > 0$, meaning $\kappa(t) \sim e^{\lambda t}$. Tracking trajectories for time $T = O(1/\varepsilon)$ requires exponential precision, giving $\mathbf{NEXP}^{\text{num}}$. \square

Theorem 12.8 (The Curse of Curvature). *For a problem requiring output accuracy ε where the curvature of the solution operator at the relevant scale is $\kappa = \Omega(\varepsilon^{-\alpha})$ for some $\alpha > 0$:*

$$\text{NComp}(P, \varepsilon) = \Omega(\varepsilon^{-\beta}) \quad \text{for some } \beta > 0$$

Problems with curvature growing as accuracy increases require correspondingly more resources.

Proof. By Theorem 3.9, the second-order error contribution is at least $\kappa \cdot \varepsilon_{\text{mach}}^2$ in worst case. To achieve output error $\leq \varepsilon$, we need:

$$\kappa \cdot \varepsilon_{\text{mach}}^2 \leq \varepsilon$$

If $\kappa = \Omega(\varepsilon^{-\alpha})$ (curvature grows as we demand finer accuracy), then:

$$\varepsilon^{-\alpha} \cdot \varepsilon_{\text{mach}}^2 \leq C\varepsilon \quad \Rightarrow \quad \varepsilon_{\text{mach}}^2 \leq C\varepsilon^{1+\alpha}$$

Thus $\varepsilon_{\text{mach}} = O(\varepsilon^{(1+\alpha)/2})$.

The number of precision bits required is $p = \log_2(1/\varepsilon_{\text{mach}}) = \Omega(\frac{1+\alpha}{2} \log_2(1/\varepsilon))$. Basic operations on p -bit numbers cost $O(p)$ time (or $O(p \log p)$ for multiplication), so the total cost for a fixed number of operations scales at least as $\Omega(\log(1/\varepsilon))$ per operation. For problems requiring $\Omega(\varepsilon^{-\gamma})$ operations (as many numerical problems do), this gives total cost $\Omega(\varepsilon^{-\gamma} \log(1/\varepsilon))$.

The key point is that high curvature forces high precision, which has direct computational cost. \square

12.3 Optimal Precision Allocation

We now prove a rigorous result on optimal bit allocation for computation graphs under a specific error model.

Theorem 12.9 (Optimal Bit Allocation). *Consider a computation graph G with n nodes $\{v_i\}_{i=1}^n$ computing $f = f_n \circ \dots \circ f_1$, where each f_i has curvature κ_i and Lipschitz constant L_i . Assume:*

- (i) *Each node v_i uses precision p_i bits, contributing local error $\varepsilon_i = 2^{-p_i}$*

(ii) Errors propagate forward via Lipschitz constants: error at node i contributes $(\prod_{j>i} L_j) \cdot \kappa_i \cdot \varepsilon_i$ to output

(iii) Total bit budget: $\sum_i p_i = B$

Then the allocation minimizing worst-case output error is:

$$p_i^* = \frac{B}{n} + \log_2 \left(\frac{w_i}{\bar{w}} \right)$$

where $w_i = \kappa_i \cdot \prod_{j>i} L_j$ is the effective weight and $\bar{w} = (\prod_i w_i)^{1/n}$ is the geometric mean.

Proof. The output error bound is $E = \sum_{i=1}^n w_i \cdot 2^{-p_i}$ where $w_i = \kappa_i \cdot \prod_{j>i} L_j$ represents the amplification of a unit error at node i to the output. We minimize E subject to the constraint $\sum_i p_i = B$.

Using Lagrange multipliers, we form the Lagrangian $\mathcal{L} = \sum_j w_j 2^{-p_j} - \lambda(\sum_j p_j - B)$ and take partial derivatives:

$$\frac{\partial \mathcal{L}}{\partial p_i} = -w_i \cdot 2^{-p_i} \cdot \ln 2 - \lambda = 0$$

This gives $w_i \cdot 2^{-p_i} = -\lambda / \ln 2 =: c$ for all i , where $c > 0$ since $\lambda < 0$ at the minimum.

From $w_i \cdot 2^{-p_i} = c$, we obtain $p_i = \log_2(w_i/c)$. Summing over i :

$$\sum_{i=1}^n p_i = \sum_{i=1}^n \log_2(w_i/c) = \sum_{i=1}^n \log_2(w_i) - n \log_2(c) = B$$

Solving for c :

$$\log_2(c) = \frac{1}{n} \sum_{i=1}^n \log_2(w_i) - \frac{B}{n} = \log_2(\bar{w}) - \frac{B}{n}$$

where $\bar{w} = (\prod_i w_i)^{1/n}$ is the geometric mean. Thus $c = \bar{w} \cdot 2^{-B/n}$.

Substituting back:

$$p_i^* = \log_2(w_i/c) = \log_2(w_i) - \log_2(\bar{w}) + \frac{B}{n} = \frac{B}{n} + \log_2 \left(\frac{w_i}{\bar{w}} \right)$$

□

Corollary 12.10 (Curvature Dominates Allocation). *When Lipschitz constants are uniform ($L_i = L$ for all i), the optimal allocation simplifies to:*

$$p_i^* = \frac{B}{n} + \log_2 \left(\frac{\kappa_i}{\bar{\kappa}} \right)$$

where $\bar{\kappa} = (\prod_i \kappa_i)^{1/n}$. Nodes with above-average curvature receive more bits; nodes with below-average curvature receive fewer.

Theorem 12.11 (Graph-Theoretic Precision Flow). *For a DAG computation graph G with Lipschitz constants $\{L_e\}$ on edges, the minimum precision at each node to achieve output accuracy ε satisfies:*

$$p_{\min}(v) \geq \max_{\text{paths } v \rightarrow t} \sum_{e \in \text{path}} \log_2(L_e) + \log_2(1/\varepsilon)$$

where t ranges over output nodes.

Proof. Precision errors at v propagate to outputs along paths. The worst-case amplification is the product of Lipschitz constants along the path. Taking logs gives the additive formula. To achieve output error ε , input error at v must be at most $\varepsilon / \prod L_e$. □

13 The Representation Theorem

Different numerical representations (floating-point, fixed-point, interval, symbolic) have different strengths. Numerical Geometry provides a unified framework.

13.1 Representation Morphisms

Definition 13.1 (Representation Type). A representation type \mathcal{R} assigns to each mathematical space X a numerical space $(X^{\mathcal{R}}, d^{\mathcal{R}}, \rho^{\mathcal{R}})$.

Example 13.2 (Standard Representations). • $\mathcal{R}_{\text{fp64}}$: IEEE 754 double precision

- \mathcal{R}_{int} : Interval arithmetic
- \mathcal{R}_{sym} : Symbolic/exact arithmetic
- \mathcal{R}_{nn} : Neural network function representation

Theorem 13.3 (Representation Comparison). For representations $\mathcal{R}_1, \mathcal{R}_2$ on space X , define:

$$d(\mathcal{R}_1, \mathcal{R}_2) := \sup_{x \in X} \inf \{d(r_1, r_2) : \rho_1(r_1) = \rho_2(r_2) = x\}$$

This is a pseudometric on representation types, quantifying the “cost” of converting between representations.

Proof. We verify the pseudometric axioms:

Non-negativity: By definition, $d(r_1, r_2) \geq 0$, so the infimum and supremum are non-negative.

Symmetry: For each x , $\inf \{d(r_1, r_2) : \rho_1(r_1) = \rho_2(r_2) = x\} = \inf \{d(r_2, r_1) : \rho_1(r_1) = \rho_2(r_2) = x\}$ since d is symmetric.

Triangle inequality: For representations $\mathcal{R}_1, \mathcal{R}_2, \mathcal{R}_3$, fix $x \in X$. For any $\delta > 0$, choose r_1, r_2, r'_2, r_3 with $\rho_i(r_i) = x$ and achieving infimums within δ . Then:

$$\inf d(r_1, r_3) \leq d(r_1, r_2) + d(r_2, r'_2) + d(r'_2, r_3)$$

where r_2, r'_2 both represent x in \mathcal{R}_2 . Taking infimums and supremum over x gives the triangle inequality.

Note: This is a pseudometric rather than a metric because $d(\mathcal{R}_1, \mathcal{R}_2) = 0$ does not imply $\mathcal{R}_1 = \mathcal{R}_2$; it only implies they are equivalent for representing elements of X . \square

Theorem 13.4 (Universal Representation). There exists a universal representation $\mathcal{R}^{\text{univ}}$ such that for any computable representation \mathcal{R} :

$$\mathcal{R} \hookrightarrow \mathcal{R}^{\text{univ}}$$

with computable embedding. All numerical representations embed into a common framework.

Proof. Define $\mathcal{R}^{\text{univ}}$ as the representation via computable signed-digit streams, also known as the Type Two Effectivity model. Specifically, for $x \in \mathbb{R}$, a representation is a sequence $(d_i)_{i \in \mathbb{Z}}$ where $d_i \in \{-1, 0, 1\}$ and $x = \sum_i d_i \cdot 2^i$, with the redundancy allowing for computable operations.

For any computable representation \mathcal{R} , there exists a computable function $\phi : R \rightarrow \mathcal{R}^{\text{univ}}$ because: (1) \mathcal{R} provides a way to compute arbitrarily precise rational approximations to represented values, and (2) such approximations can be converted to signed-digit streams. This construction is standard in computable analysis (see Weihrauch’s Computable Analysis).

The embedding is a numerical morphism because Lipschitz constants are preserved: if operations in \mathcal{R} amplify errors by at most L , the corresponding operations in $\mathcal{R}^{\text{univ}}$ have the same bound. \square

13.2 The Representation-Complexity Trade-off

Theorem 13.5 (No Free Lunch for Representations). *For any representation \mathcal{R} and any non-trivial problem class \mathcal{P} with $|\mathcal{P}| \geq 2$:*

$$\mathbb{E}_{P \in \mathcal{P}}[\text{NComp}_{\mathcal{R}}(P, \varepsilon)] \geq \Omega(\log(1/\varepsilon)^{\dim(\mathcal{P})})$$

No representation is universally optimal; choose representations matched to the problem structure.

Proof. Consider a problem class \mathcal{P} as a parameterized family $\{P_\theta : \theta \in \Theta\}$ where Θ has dimension $d = \dim(\mathcal{P})$. To distinguish P_θ from $P_{\theta'}$ requires precision sufficient to resolve $\|\theta - \theta'\|$.

By Theorem 12.2, achieving precision ε in a d -dimensional space requires $\Omega(d \log(1/\varepsilon))$ bits of information. For any fixed representation \mathcal{R} , at least one problem in the class must have complexity at least $\Omega(\log(1/\varepsilon)^d)$, because the total information content of the problem class grows as $(1/\varepsilon)^d$ distinguishable problems.

More precisely, by a counting argument: if all problems in \mathcal{P} could be solved with fewer than $c \cdot \log(1/\varepsilon)^d$ operations, then the total distinguishing capacity would be at most $2^{c \log(1/\varepsilon)^d}$, but we need to distinguish $(1/\varepsilon)^d$ problems, requiring $c \geq 1$. \square

Theorem 13.6 (Optimal Representation Theorem). *For a problem P with curvature tensor \mathcal{K}_P (the Hessian of the error functional), the optimal representation uses coordinates that diagonalize \mathcal{K}_P :*

$$\mathcal{K}_P^{\text{opt}} = U^T \cdot \mathcal{K}_P \cdot U \quad \text{diagonal}$$

where U is orthogonal. The optimal bit allocation assigns precision proportional to $\log \lambda_i$ where λ_i are the eigenvalues.

Proof. In coordinates $(y_1, \dots, y_n) = U^T x$, the curvature tensor is diagonal with entries $\lambda_1, \dots, \lambda_n$. The error contribution from direction y_i is $\lambda_i \cdot \varepsilon_i^2$ by Theorem 3.9.

Total error: $E = \sum_i \lambda_i \varepsilon_i^2$. With bit budget B and $\varepsilon_i = 2^{-p_i}$:

$$E = \sum_i \lambda_i \cdot 4^{-p_i}$$

By Lagrange multipliers (similar to Theorem 12.9), the optimum is $p_i \propto \log_2(\lambda_i)$.

In the original coordinates, this corresponds to using an ellipsoid of precision values aligned with the eigenvectors of \mathcal{K}_P . This is precisely what the FFT achieves for convolution (diagonalizing the circulant structure) and SVD for linear systems (diagonalizing the sensitivity structure). \square

14 Numerical Equivalence and Canonical Forms

This section studies when two algorithms computing the same mathematical function are “numerically equivalent”—exhibiting the same asymptotic error behavior. We establish rigorous results for restricted classes and state open problems for the general theory.

14.1 Numerical Equivalence of Algorithms

Definition 14.1 (Numerical Equivalence). *Two algorithms A_1, A_2 computing $f : X \rightarrow Y$ are numerically equivalent, written $A_1 \sim_{\text{num}} A_2$, if there exist constants $C_1, C_2 > 0$ such that for all inputs x and precisions $\varepsilon > 0$:*

$$\Phi_{A_1}(\varepsilon) \leq C_1 \cdot \Phi_{A_2}(C_2 \varepsilon) \quad \text{and} \quad \Phi_{A_2}(\varepsilon) \leq C_2 \cdot \Phi_{A_1}(C_1 \varepsilon)$$

where Φ_A is the error functional of algorithm A .

Definition 14.2 (Error Functional). *For algorithm A computing $f : \mathbb{R}^n \rightarrow \mathbb{R}^m$ with input x and hardware precision ε_H , the error functional is:*

$$\Phi_A(\varepsilon) := \sup_{\|x\| \leq 1} \|A(x; \varepsilon_H) - f(x)\|$$

where $A(x; \varepsilon_H)$ denotes the output of A executed with precision ε_H .

Proposition 14.3 (Numerical Equivalence is an Equivalence Relation). *The relation \sim_{num} is reflexive, symmetric, and transitive.*

Proof. **Reflexivity:** Take $C_1 = C_2 = 1$.

Symmetry: The definition is symmetric in A_1, A_2 .

Transitivity: If $A_1 \sim_{\text{num}} A_2$ with constants C_1, C_2 and $A_2 \sim_{\text{num}} A_3$ with constants C'_1, C'_2 , then:

$$\Phi_{A_1}(\varepsilon) \leq C_1 \Phi_{A_2}(C_2 \varepsilon) \leq C_1 C'_1 \Phi_{A_3}(C'_2 C_2 \varepsilon)$$

So $A_1 \sim_{\text{num}} A_3$ with constants $C_1 C'_1$ and $C_2 C'_2$. \square

14.2 Second-Order Error Expansion

Theorem 14.4 (Error Functional Expansion). *For a C^2 function $f : \mathbb{R}^n \rightarrow \mathbb{R}^m$ and any algorithm A computing f via a composition of smooth operations, the error functional admits the expansion:*

$$\Phi_A(\varepsilon) = L_A \cdot \varepsilon + Q_A \cdot \varepsilon^2 + O(\varepsilon^3)$$

where:

- $L_A = \|Df\|_{\text{op}} \cdot c_A$ for some algorithm-dependent constant $c_A \geq 1$
- $Q_A \geq \frac{1}{2} \|D^2 f\|_{\text{op}}$ is a second-order coefficient depending on the algorithm structure

Proof. Let A compute f via a directed acyclic graph of operations $f = g_k \circ g_{k-1} \circ \dots \circ g_1$.

First-order term: Input perturbation $\|\delta\| \leq \varepsilon$ produces output error:

$$\|f(x + \delta) - f(x)\| \leq \|Df_x\| \cdot \|\delta\| + O(\|\delta\|^2) \leq \|Df\|_{\text{op}} \cdot \varepsilon + O(\varepsilon^2)$$

Additionally, each operation g_i introduces roundoff error $O(\varepsilon_H)$. With k operations, each with Lipschitz constant L_i , the accumulated roundoff is bounded by $\sum_{i=1}^k (\prod_{j>i} L_j) \varepsilon_H \leq c_A \cdot \varepsilon_H$ for some c_A depending on the algorithm.

The first-order coefficient $L_A = \|Df\|_{\text{op}} \cdot c_A$ captures both input sensitivity and roundoff accumulation.

Second-order term: By Taylor expansion:

$$f(x + \delta) = f(x) + Df_x(\delta) + \frac{1}{2} D^2 f_x(\delta, \delta) + O(\|\delta\|^3)$$

The second derivative contributes $\frac{1}{2} \|D^2 f_x\| \cdot \|\delta\|^2 \leq \frac{1}{2} \|D^2 f\|_{\text{op}} \cdot \varepsilon^2$.

Additional second-order contributions arise from products of first-order errors at intermediate stages. These are algorithm-dependent, giving the bound $Q_A \geq \frac{1}{2} \|D^2 f\|_{\text{op}}$. \square

Definition 14.5 (Minimal Error Algorithm). *An algorithm A for f is error-minimal if it achieves:*

$$\Phi_A(\varepsilon) = L_f \cdot \varepsilon + \frac{1}{2} \|D^2 f\|_{\text{op}} \cdot \varepsilon^2 + O(\varepsilon^3)$$

with $L_f = \|Df\|_{\text{op}}$ (i.e., $c_A = 1$ in the notation above).

14.3 Characterization Theorem for Smooth Functions

Theorem 14.6 (Characterization of Numerical Equivalence). *Let A_1, A_2 be two algorithms computing a C^2 function $f : \mathbb{R}^n \rightarrow \mathbb{R}^m$. The following are equivalent:*

- (i) $A_1 \sim_{\text{num}} A_2$ (numerical equivalence)
- (ii) The error functionals satisfy $\Phi_{A_1}(\varepsilon)/\Phi_{A_2}(\varepsilon) \rightarrow c$ for some $0 < c < \infty$ as $\varepsilon \rightarrow 0$
- (iii) The first-order coefficients satisfy $L_{A_1} \asymp L_{A_2}$ (within constant factors)

Proof. (i) \Rightarrow (ii): By definition, $\Phi_{A_1}(\varepsilon) \leq C_1 \Phi_{A_2}(C_2 \varepsilon)$. For small ε :

$$\frac{\Phi_{A_1}(\varepsilon)}{\Phi_{A_2}(\varepsilon)} \leq C_1 \cdot \frac{\Phi_{A_2}(C_2 \varepsilon)}{\Phi_{A_2}(\varepsilon)} = C_1 \cdot \frac{L_{A_2} C_2 \varepsilon + O(\varepsilon^2)}{L_{A_2} \varepsilon + O(\varepsilon^2)} \rightarrow C_1 C_2$$

The symmetric bound gives the lower bound.

(ii) \Rightarrow (iii): Since $\Phi_{A_i}(\varepsilon) = L_{A_i} \varepsilon + O(\varepsilon^2)$:

$$\frac{\Phi_{A_1}(\varepsilon)}{\Phi_{A_2}(\varepsilon)} = \frac{L_{A_1} + O(\varepsilon)}{L_{A_2} + O(\varepsilon)} \rightarrow \frac{L_{A_1}}{L_{A_2}}$$

Thus the limit being finite and non-zero implies $L_{A_1} \asymp L_{A_2}$.

(iii) \Rightarrow (i): If $L_{A_1} \leq C \cdot L_{A_2}$ and $L_{A_2} \leq C \cdot L_{A_1}$, then:

$$\Phi_{A_1}(\varepsilon) = L_{A_1} \varepsilon + O(\varepsilon^2) \leq C L_{A_2} \varepsilon + O(\varepsilon^2) \leq C' \Phi_{A_2}(\varepsilon)$$

for some C' . The symmetric bound follows similarly. \square

14.4 Canonical Forms: Existence for Restricted Classes

We establish rigorous existence results for specific function classes where canonical (error-minimal) algorithms can be characterized.

Theorem 14.7 (Canonical Form for Linear Maps). *For a linear map $f : \mathbb{R}^n \rightarrow \mathbb{R}^m$ given by matrix A , the error-minimal algorithm exists and satisfies:*

$$\Phi_{A^{\text{can}}}(\varepsilon) = \|A\|_2 \cdot \varepsilon + O(\varepsilon^2)$$

The minimal algorithm computes Ax via standard matrix-vector multiplication in exact arithmetic (with only roundoff error).

Proof. Since $f(x) = Ax$ is linear, $Df = A$ and $D^2f = 0$. The Taylor expansion gives:

$$f(x + \delta) - f(x) = A\delta$$

Thus input error $\|\delta\| \leq \varepsilon$ produces output error exactly $\|A\delta\| \leq \|A\|_2 \varepsilon$.

The algorithm computing $y = Ax$ via $y_i = \sum_j A_{ij} x_j$ has roundoff error from $O(n)$ additions per component. In the standard floating-point model, this contributes $O(n\varepsilon)$ to each component, giving total error $O(\sqrt{mn}\varepsilon) = O(\|A\|_F \varepsilon)$ where $\|A\|_F$ is the Frobenius norm.

However, since $\|A\|_2 \leq \|A\|_F \leq \sqrt{\min(m, n)} \|A\|_2$, both algorithms have first-order coefficient $\Theta(\|A\|_2)$, hence are numerically equivalent.

The minimal algorithm uses compensated summation (Kahan) to reduce roundoff, achieving the theoretical minimum $\|A\|_2 \varepsilon + O(\varepsilon^2)$. \square

Theorem 14.8 (Canonical Form for Quadratic Maps). *For a quadratic map $f(x) = \frac{1}{2}x^T Hx$ with symmetric $H \in \mathbb{R}^{n \times n}$, the error-minimal algorithm satisfies:*

$$\Phi_{f^{\text{can}}}(\varepsilon) = \|H\|_2 \cdot \varepsilon + \frac{1}{2}\|H\|_2 \cdot \varepsilon^2 + O(\varepsilon^3)$$

The minimal algorithm diagonalizes H and computes in the eigenbasis.

Proof. We have $Df_x = Hx$ so $\|Df\|_{\text{op}} = \|H\|_2 \cdot \|x\|$ on the unit ball, and $D^2f = H$ is constant.

Lower bound: Any algorithm has input sensitivity $\|H\|_2\varepsilon$ plus curvature contribution $\frac{1}{2}\|H\|_2\varepsilon^2$ by Theorem 14.4.

Upper bound: Let $H = U\Lambda U^T$ with $\Lambda = \text{diag}(\lambda_1, \dots, \lambda_n)$. The algorithm:

1. Compute $y = U^T x$ (orthogonal, error $O(\varepsilon)$ per component)
2. Compute $z_i = \lambda_i y_i^2 / 2$ for each i (single multiplication)
3. Compute $f(x) = \sum_i z_i$ (summation, error $O(n\varepsilon)$ with compensation)

achieves total error $\|H\|_2\varepsilon + O(\varepsilon^2)$ since the diagonal structure isolates errors. \square

Theorem 14.9 (Canonical Form for Polynomial Maps). *For a polynomial $p : \mathbb{R} \rightarrow \mathbb{R}$ of degree d evaluated on $[-1, 1]$, Horner's method is error-minimal:*

$$\Phi_{\text{Horner}}(\varepsilon) = L_p \cdot \varepsilon + O(d^2\varepsilon^2)$$

where $L_p = \max_{x \in [-1, 1]} |p'(x)|$.

Proof. Horner's method computes $p(x) = a_0 + x(a_1 + x(a_2 + \dots))$ using d multiplications and d additions. Each introduces error $O(\varepsilon)$, and the errors propagate with factor at most $|x| \leq 1$ per stage.

By induction on d : Let $q(x) = a_1 + x(a_2 + \dots)$ be the inner polynomial. The computed value is:

$$\hat{p}(x) = a_0 + x \cdot \hat{q}(x) + O(\varepsilon)$$

where $\hat{q}(x) = q(x) + O((d-1)\varepsilon)$ by induction. Thus $\hat{p}(x) = p(x) + O(d\varepsilon)$.

The first-order coefficient is $L_p = \max |p'(x)|$, matching the theoretical minimum. The $O(d^2\varepsilon^2)$ term comes from products of first-order errors. \square

14.5 Non-Existence and Obstructions

Theorem 14.10 (Non-Existence of Universal Canonical Form). *There is no algorithm \mathcal{U} such that for all C^2 functions f :*

- (i) \mathcal{U} computes f given a representation of f
- (ii) \mathcal{U} achieves the minimal error functional for f

Proof. Suppose such \mathcal{U} exists. Consider the family $f_a(x) = e^{ax}$ for $a \in \mathbb{R}$. The optimal algorithm depends on a :

- For $a \approx 0$: Taylor expansion around 0 is optimal
- For large a : range reduction and table lookup is optimal

A universal algorithm \mathcal{U} cannot adaptively choose the optimal method without reading enough bits of a to distinguish these regimes. This introduces overhead of at least $\Omega(\log(1/\varepsilon))$ in the constants, precluding universality.

More precisely: by a counting argument, the number of “essentially different” optimal algorithms for functions in a d -parameter family grows with the parameter space volume. Any single algorithm can be optimal for only a measure-zero subset. \square

Corollary 14.11 (Algorithm-Function Pairing). *Optimal numerical computation requires matching the algorithm to the function structure. This is the computational content of the geometric invariants (curvature, cohomology) developed in previous sections.*

14.6 Optimality of Standard Algorithms

We now rigorously verify that several standard algorithms achieve (or nearly achieve) minimal error.

Theorem 14.12 (FFT Achieves Minimal Error for DFT). *The Cooley-Tukey FFT computes the discrete Fourier transform with error:*

$$\Phi_{\text{FFT}}(\varepsilon) = O(\log n \cdot \varepsilon)$$

This is optimal to within logarithmic factors for any algorithm computing DFT.

Proof. **Upper bound (FFT achieves this):** The FFT has depth $O(\log n)$ and each level consists of “butterfly” operations. Each butterfly is numerically stable with error $O(\varepsilon)$. By Theorem 10.1, the total error is $O(\log n \cdot \varepsilon)$.

Lower bound: The DFT matrix F_n has $\|F_n\|_2 = \sqrt{n}$ (since $F_n^* F_n = nI$). However, the DFT is a unitary operation (after scaling), so its condition number is 1.

The lower bound comes from the computation graph: any algorithm computing n outputs from n inputs with $O(n \log n)$ operations must have at least one path of length $\Omega(\log n)$ from some input to some output. Each operation on this path contributes $\Omega(\varepsilon)$ error.

Thus $\Phi_{\text{any}}(\varepsilon) = \Omega(\log n \cdot \varepsilon)$ for any algorithm, and FFT achieves this. \square

Theorem 14.13 (Kahan Summation is Optimal for Summation). *Kahan’s compensated summation algorithm computes $S = \sum_{i=1}^n x_i$ with error:*

$$\Phi_{\text{Kahan}}(\varepsilon) = (2 + O(n\varepsilon)) \cdot \varepsilon$$

independent of n . This is optimal since any algorithm has error at least ε (from the final rounding).

Proof. Kahan summation maintains a running sum s and a compensation term c :

```

 $s \leftarrow 0, c \leftarrow 0$ 
for  $i = 1$  to  $n$  do
     $y \leftarrow x_i - c$ 
     $t \leftarrow s + y$ 
     $c \leftarrow (t - s) - y$  {recovers the roundoff}
     $s \leftarrow t$ 
end for
return  $s$ 

```

The compensation term c captures the roundoff error from adding y to s . In exact arithmetic, $(t - s) - y = 0$, but in floating-point, this equals the roundoff error.

By induction: after k terms, $s + c = \sum_{i=1}^k x_i + O(\varepsilon^2)$ where the $O(\varepsilon^2)$ comes from second-order roundoff (products of first-order errors). After n terms, the error is $(2 + O(n\varepsilon))\varepsilon$.

The lower bound ε is trivial: the final sum must be rounded to floating-point precision. \square

Corollary 14.14 (Numerical Invariants). *The following quantities are invariants under numerical equivalence of algorithms computing f :*

- (i) *The Lipschitz constant $L_f = \|Df\|_{\text{op}}$ (first-order sensitivity)*
- (ii) *The Hessian norm $\|D^2f\|_{\text{op}}$ (second-order sensitivity)*

Proof. These are properties of the function f itself, not of any particular algorithm. By Theorem 14.6, numerical equivalence is characterized by the first-order coefficient L_A , which is proportional to L_f . Different algorithms may have different constants c_A in $L_A = c_A \cdot L_f$, but L_f is intrinsic. \square

Corollary 14.15 (Representation Independence). *For any function f , the canonical error $L_f\varepsilon + (\mathcal{K}_f)\varepsilon^2$ is independent of:*

- (i) *The choice of floating-point format (fp16, fp32, fp64) up to scaling ε*
- (ii) *The computation graph structure (different orderings yield permuted \mathcal{K})*
- (iii) *The hardware implementation (CPU, GPU, TPU) up to the $\Gamma(H)$ term*

Proof. (i) Changing floating-point format scales ε by a constant factor. The structure $L\varepsilon + c\varepsilon^2$ is preserved.

(ii) Different computation graph orderings compute the same function f , hence have the same Hessian D^2f .

(iii) Hardware affects only the roundoff term $\Gamma(H)$, which enters at order ε in the linear error model. The second derivative D^2f depends only on the mathematical function f , not its implementation. \square

Definition 14.16 (Precision-Preserving Transformation). *A compiler transformation T is precision-preserving if the transformed algorithm satisfies $L_{T(A)} \leq C \cdot L_A$ for some constant $C > 0$.*

Proposition 14.17 (Classification of Common Transformations). *The following transformations have provable precision properties:*

- (i) **Identity replacement** ($x - x \rightarrow 0$): Precision-preserving with $C = 1$. Eliminates a potential cancellation error.
- (ii) **FMA fusion** ($a \cdot b + c \rightarrow \text{fma}(a, b, c)$): Precision-preserving with $C = 1$. The FMA computes $ab + c$ with a single rounding, reducing error from 2 roundings to 1.
- (iii) **Associativity reordering** ($((a + b) + c \rightarrow a + (b + c))$): NOT precision-preserving in general. For $a = 1, b = \varepsilon_H, c = -1$: left gives ε_H , right gives 0.

Proof. (i): The expression $x - x$ in floating-point may produce non-zero results due to prior rounding of x . Replacing with 0 removes this error source.

(ii): Standard floating-point: $(a \otimes b) \oplus c$ has error from two roundings: $\text{fl}(ab)(1 + \delta_1) + c(1 + \delta_2)$ with $|\delta_i| \leq \varepsilon_H$. The FMA has one rounding: $\text{fl}(ab + c) = (ab + c)(1 + \delta)$.

(iii): Counterexample: Let $a = 1$, $b = 2^{-53}$ (machine epsilon in fp64), $c = -1$. Then $(a+b)+c = (1+2^{-53}) - 1 = 2^{-53}$ (exact), but $a + (b+c) = 1 + (2^{-53} - 1) = 1 + (-1) = 0$ since $2^{-53} - 1$ rounds to -1 . \square

14.7 Composition and Error Propagation

Lemma 14.18 (Composition of Algorithms). *For algorithms A_f computing f and A_g computing g , the composition $A_g \circ A_f$ computes $g \circ f$ with error:*

$$\Phi_{A_g \circ A_f}(\varepsilon) \leq L_g \cdot \Phi_{A_f}(\varepsilon) + \Phi_{A_g}(L_f \cdot \varepsilon) + O(\varepsilon^2)$$

Proof. By the chain rule, $D(g \circ f)_x = Dg_{f(x)} \circ Df_x$, so $L_{g \circ f} \leq L_g \cdot L_f$. For the Hessian, the chain rule gives:

$$D^2(g \circ f)_x(u, v) = D^2g_{f(x)}(Df_x(u), Df_x(v)) + Dg_{f(x)}(D^2f_x(u, v))$$

Taking traces (summing over an orthonormal basis $\{e_i\}$):

$$(D^2(g \circ f)) = \sum_i D^2(g \circ f)(e_i, e_i) = \sum_i D^2g(Df(e_i), Df(e_i)) + \sum_i Dg(D^2f(e_i, e_i))$$

The first sum equals $((Df)^T(D^2g)(Df))$, which is bounded by $\|Df\|^2(D^2g)$. The second sum equals $(Dg \cdot D^2f) \leq \|Dg\|(D^2f)$.

Thus the curvature of the composition satisfies:

$$(\mathcal{K}_{g \circ f}) \leq L_f^2 \cdot (\mathcal{K}_g) + L_g \cdot (\mathcal{K}_f)$$

The composition $A_g^{\text{can}} \circ A_f^{\text{can}}$ achieves:

$$\begin{aligned} \Phi_{A_g \circ A_f}(\varepsilon) &= \Phi_{A_g}(\Phi_{A_f}(\varepsilon)) \\ &= L_g(L_f\varepsilon + (\mathcal{K}_f)\varepsilon^2) + (\mathcal{K}_g)(L_f\varepsilon)^2 + O(\varepsilon^3) \\ &= L_g L_f \varepsilon + (L_g(\mathcal{K}_f) + L_f^2(\mathcal{K}_g))\varepsilon^2 + O(\varepsilon^3) \end{aligned}$$

This matches the curvature bound, so the composition is canonical. \square

Lemma 14.19 (Universality under Perturbation). *If f_t is a smooth family of functions with $f_0 = f$, then:*

$$\frac{d}{dt}\Big|_{t=0}(\mathcal{K}_{f_t}) = \left(\frac{\partial \mathcal{K}_f}{\partial t} \right)$$

The total curvature varies linearly with perturbations.

Proof. Since $\mathcal{K}_{f_t} = \frac{1}{2}D^2f_t$, differentiation under the trace gives:

$$\frac{d}{dt}(\mathcal{K}_{f_t}) = \frac{1}{2} \left(\frac{d}{dt}D^2f_t \right) = \frac{1}{2}(D^2\dot{f}_t)$$

where $\dot{f}_t = \partial f_t / \partial t$. This is the trace of the Hessian of the perturbation. \square

Part III

Applications

The preceding theory is not merely abstract—it provides systematic tools for concrete problems across numerical computing. This part demonstrates applications to automatic differentiation, neural networks, scientific computing, compilers, and formal verification.

15 Automatic Differentiation with Precision Tracking

15.1 The Problem

Automatic differentiation (AD) computes exact derivatives of programs. But “exact” means exact in real arithmetic—what about finite precision?

Key Questions:

1. How does roundoff error propagate through the derivative computation?
2. When does the computed gradient have meaningful precision?
3. How should precision be allocated across forward and backward passes?

15.2 AD as Numerical Morphism

Definition 15.1 (Derivative Morphism). *For a numerical morphism $f : A \rightarrow B$, define the derivative morphism:*

$$Df : A \times TA \rightarrow TB, \quad Df(a, v) := D|f|_a(v)$$

with error functional:

$$\Phi_{Df}(\varepsilon, H) = L_{Df} \cdot \varepsilon + \Gamma_{Df}(H)$$

Theorem 15.2 (AD Error Propagation). *For the composition $F = f_n \circ \dots \circ f_1$ computed via reverse-mode AD:*

$$\Phi_{DF}(\varepsilon, H) \leq \sum_{i=1}^n \left(\prod_{j \neq i} L_j \right) \Phi_{Df_i}(\varepsilon_i, H)$$

The gradient error is controlled by the product of Lipschitz constants of all other layers.

Proof. By the chain rule, $DF = Df_n \circ Df_{n-1} \circ \dots \circ Df_1$. In reverse-mode AD, the gradient is computed as:

$$\nabla_x F = (Df_1)^T \circ (Df_2)^T \circ \dots \circ (Df_n)^T (\nabla_y)$$

where ∇_y is the output gradient.

Consider the error at layer i in the backward pass. The local error $\Phi_{Df_i}(\varepsilon_i, H)$ propagates backward through layers $i-1, i-2, \dots, 1$, each amplifying the error by their respective Lipschitz constants. Thus the error contribution from layer i to the final gradient is:

$$\left(\prod_{j=1}^{i-1} L_j \right) \cdot \Phi_{Df_i}(\varepsilon_i, H)$$

However, this error originated from the forward pass values at layer i , which themselves have accumulated error from layers $1, \dots, i - 1$ amplified through layers $i + 1, \dots, n$. The combined amplification factor for error at layer i is:

$$\left(\prod_{j < i} L_j \right) \cdot \left(\prod_{j > i} L_j \right) = \frac{\prod_{j=1}^n L_j}{L_i} = \prod_{j \neq i} L_j$$

Summing over all layers:

$$\Phi_{DF}(\varepsilon, H) \leq \sum_{i=1}^n \left(\prod_{j \neq i} L_j \right) \Phi_{Df_i}(\varepsilon_i, H)$$

□

Corollary 15.3 (Gradient Precision in Deep Networks). *For an L -layer network with spectral norms $\sigma_1, \dots, \sigma_L$:*

$$\Phi_{\nabla \mathcal{L}}(\varepsilon, H) \leq L \cdot \left(\prod_{i=1}^L \sigma_i \right) \cdot (\varepsilon + \Delta)$$

If $\prod_i \sigma_i > 10^{16}$, gradients in fp64 have zero significant digits.

15.3 Precision-Tracked Tensors

Definition 15.4 (Precision Tensor). *A precision tensor is a pair (x, σ) where:*

- $x \in \mathbb{R}^{n_1 \times \dots \times n_k}$ is the value
- $\sigma \in \mathbb{R}_{\geq 0}^{n_1 \times \dots \times n_k}$ is the element-wise precision bound

Proposition 15.5 (Precision Propagation Rules).

$$\begin{aligned} (x, \sigma_x) + (y, \sigma_y) &\mapsto (x + y, \sigma_x + \sigma_y + \varepsilon_H) \\ (x, \sigma_x) \cdot (y, \sigma_y) &\mapsto (x \cdot y, |y|\sigma_x + |x|\sigma_y + \varepsilon_H|x||y|) \\ \exp(x, \sigma_x) &\mapsto (e^x, e^x \cdot \sigma_x + \varepsilon_H e^x) \end{aligned}$$

15.4 Implementation: JAX/PyTorch Extension

Algorithm 2 Precision-Tracked Forward Pass

Require: Computation graph G , input (x, σ_x) , hardware H

Ensure: Output (y, σ_y) with tracked precision

- 1: **for** each node v in topological order **do**
- 2: Compute value $y_v = f_v(\text{inputs})$
- 3: Compute precision σ_v using propagation rules
- 4: Check: if $\sigma_v >$ threshold, emit warning
- 5: **end for**
- 6: **return** $(y_{\text{output}}, \sigma_{\text{output}})$

16 Neural Network Quantization and Training

16.1 The Quantization Problem

Modern neural networks are trained in fp32/fp16 but deployed in int8 or lower. The question: **which layers can be quantized without accuracy loss?**

Principle 16.1 (Curvature-Guided Quantization). *Layers with high curvature require high precision. Layers with low curvature tolerate aggressive quantization.*

16.2 Layer-wise Curvature Analysis

Proposition 16.2 (Neural Network Layer Curvatures).

Layer Type	Curvature	Quantization Safety
Linear (dense)	0	Safe to int8
ReLU	0	Safe to int8
Softmax	1/2	Needs fp16
LayerNorm	$1/\sigma^2$	Depends on variance
Attention	$O(1)$	Needs fp16
GELU	$O(1)$	Marginal

Theorem 16.3 (Quantization Precision Bound). *For a layer f with curvature κ_f and input range D , quantization to b bits achieves accuracy:*

$$\varepsilon \geq c \cdot \kappa_f \cdot D^2 \cdot 2^{-b}$$

where c is a constant depending on the quantization scheme. To achieve target accuracy $\varepsilon_{\text{target}}$, we need:

$$b \geq \log_2 \left(\frac{c \cdot \kappa_f \cdot D^2}{\varepsilon_{\text{target}}} \right)$$

Proof. Quantization to b bits introduces uniform error $\delta = D/2^b$ per coordinate in the input range $[0, D]$.

By Theorem 3.9, the output error for a function with curvature κ_f is:

$$\varepsilon_{\text{out}} \geq L_f \cdot \delta + \kappa_f \cdot \delta^2 - O(\delta^3)$$

For the leading term involving curvature:

$$\kappa_f \cdot \delta^2 = \kappa_f \cdot \frac{D^2}{2^{2b}}$$

The constant c depends on the specific quantization scheme:

- Uniform quantization: $c = 1/12$ (from uniform distribution variance)
- Stochastic rounding: $c = 1/4$ (worst-case)

Inverting for b gives the stated bound. Note that this is a lower bound on the *required* precision—one cannot do better than this regardless of the quantization algorithm used. \square

16.3 Mixed-Precision Training

Algorithm 3 Curvature-Optimal Mixed Precision

Require: Network N , target accuracy ε , precision budget B

Ensure: Per-layer precision assignment $\{p_\ell\}$

- 1: **for** each layer ℓ **do**
 - 2: Compute curvature κ_ℓ via Hessian estimation
 - 3: Compute minimum precision $p_\ell^{\min} = \log_2(\kappa_\ell D_\ell^2 / \varepsilon)$
 - 4: **end for**
 - 5: Solve optimization: $\min \sum_\ell p_\ell \cdot |\theta_\ell|$ s.t. $p_\ell \geq p_\ell^{\min}$, $\sum p_\ell |\theta_\ell| \leq B$
 - 6: **return** Optimal assignment $\{p_\ell\}$
-

16.4 Transformer-Specific Analysis

Theorem 16.4 (Attention Layer Precision). *The attention mechanism $\text{Attn}(Q, K, V) = \text{softmax}(QK^T / \sqrt{d})V$ has:*

- *Curvature: $\kappa_{\text{Attn}} = O(\|Q\|^2 \|K\|^2 / d)$ after scaling*
- *For typical trained values $\|Q\|, \|K\| \approx 5$ and $d = 64$: $\kappa \approx 10$*
- *Precision requirement: $p \geq 10$ bits for $\varepsilon = 10^{-3}$*

Attention requires at least 10-bit precision, making int8 marginal without careful calibration.

Proof. The attention mechanism composes: (1) bilinear QK^T , (2) scaling by $1/\sqrt{d}$, (3) softmax, and (4) matrix-vector $(\cdot)V$.

Step 1 (Bilinear): For QK^T where $Q, K \in \mathbb{R}^{n \times d}$, each entry $(QK^T)_{ij} = \sum_k Q_{ik} K_{jk}$ is a sum of d products. Each product $Q_{ik} K_{jk}$ has curvature $1/2$. For the sum, curvature accumulates: the Hessian of the sum has operator norm bounded by $d \cdot (1/2) \cdot 2 = d$ (the factor of 2 accounts for cross-terms). More precisely, using $\|Q\|, \|K\|$ as bounds on entries:

$$\kappa_{QK^T} \leq \frac{d}{2}$$

Step 2 (Scaling): Scaling by $\alpha = 1/\sqrt{d}$ scales curvature by $\alpha^2 = 1/d$:

$$\kappa_{\text{scaled}} = \frac{1}{d} \cdot \frac{d}{2} = \frac{1}{2}$$

Step 3 (Softmax): Softmax has curvature $\leq 1/2$ by Proposition 3.8. By the composition rule (Lemma 3.4(ii)):

$$\kappa_{\text{after softmax}} \leq \kappa_{\text{softmax}} \cdot L_{\text{scaled}}^2 + L_{\text{softmax}} \cdot \kappa_{\text{scaled}}$$

With $L_{\text{scaled}} = \|Q\| \|K\| / \sqrt{d}$ and $L_{\text{softmax}} = 1$:

$$\kappa \leq \frac{1}{2} \cdot \frac{\|Q\|^2 \|K\|^2}{d} + 1 \cdot \frac{1}{2} = \frac{\|Q\|^2 \|K\|^2}{2d} + \frac{1}{2}$$

Step 4 (Multiply by V): This is linear with Lipschitz constant $\|V\|$, scaling curvature by $\|V\|$.

Total: For $\|Q\| = \|K\| = \|V\| = 5$ and $d = 64$:

$$\kappa_{\text{Attn}} \approx \|V\| \cdot \left(\frac{25 \cdot 25}{2 \cdot 64} + \frac{1}{2} \right) \approx 5 \cdot (4.9 + 0.5) \approx 27$$

The precision bound: to achieve $\varepsilon = 10^{-3}$, we need $\kappa \cdot \varepsilon_H^2 < \varepsilon$, so $\varepsilon_H < \sqrt{10^{-3}/27} \approx 0.006$. This requires $p \geq \log_2(1/0.006) \approx 7$ bits from the curvature term alone. Adding margin for linear error and accumulation across layers: $p \geq 10$ bits is reasonable. \square

Corollary 16.5 (Where to Quantize in Transformers). 1. **Embedding layers:** Safe for int8 (curvature 0)

- 2. **Attention:** Requires fp16 or careful scaling
- 3. **FFN:** Safe for int8 (mostly linear + ReLU)
- 4. **LayerNorm:** Keep in fp32 (high curvature when variance is small)
- 5. **Final logits:** Keep in fp32 (affects loss directly)

16.5 Deep Learning Training Dynamics

Theorem 16.6 (The Loss Landscape Curvature Theorem). *For a neural network f_θ with loss $\mathcal{L}(\theta) = \mathbb{E}[\ell(f_\theta(x), y)]$, the loss Hessian satisfies:*

$$\nabla_\theta^2 \mathcal{L} = \mathbb{E} \left[J_\theta f^T \cdot \nabla_{\hat{y}}^2 \ell \cdot J_\theta f + \sum_i \frac{\partial \ell}{\partial \hat{y}_i} \nabla_\theta^2 f_i \right]$$

where $J_\theta f$ is the Jacobian of outputs w.r.t. parameters.

Practical impact:

- 1. Large Jacobian norms cause high loss curvature (gradient explosion)
- 2. The first term dominates near minima; the second term matters during training
- 3. Gradient clipping reduces effective Jacobian norm, lowering curvature

Proof. By the chain rule, $\nabla_\theta \mathcal{L} = \mathbb{E}[J_\theta f^T \cdot \nabla_{\hat{y}} \ell]$. Differentiating again:

$$\begin{aligned} \nabla_\theta^2 \mathcal{L} &= \mathbb{E} [\nabla_\theta (J_\theta f^T \cdot \nabla_{\hat{y}} \ell)] \\ &= \mathbb{E} \left[J_\theta f^T \cdot \nabla_{\hat{y}}^2 \ell \cdot J_\theta f + \sum_i \frac{\partial \ell}{\partial \hat{y}_i} \nabla_\theta^2 f_i \right] \end{aligned}$$

The first term captures how output sensitivity amplifies loss curvature; the second captures intrinsic network curvature weighted by gradients. Near a minimum, $\nabla \ell \approx 0$, so the first term dominates.

The curvature bound follows: $\kappa_{\mathcal{L}} \leq \|J_\theta f\|^2 \cdot \kappa_\ell + \|\nabla \ell\| \cdot \max_i \kappa_{f_i}$. \square

Theorem 16.7 (Gradient Precision Threshold). *For gradient descent with learning rate η on loss with curvature κ :*

$$\boxed{\eta \cdot \kappa \cdot \varepsilon_{\text{grad}} < 1}$$

must hold for convergence. If $\varepsilon_{\text{grad}} \geq 1/(\eta \cdot \kappa)$, gradient noise dominates signal.

Practical impact: This explains why large-batch training needs higher precision. With batch size B :

$$\varepsilon_{\text{grad}} \sim 1/\sqrt{B} \cdot \varepsilon_{\text{machine}}$$

Larger batches reduce gradient noise, allowing lower precision.

Proof. Standard gradient descent analysis shows convergence requires the update $\theta_{t+1} = \theta_t - \eta \nabla \mathcal{L}(\theta_t)$ to decrease the loss. With computed gradient $\tilde{g} = \nabla \mathcal{L} + e$ where $\|e\| \leq \varepsilon_{\text{grad}} \|\nabla \mathcal{L}\|$:

$$\begin{aligned} \mathcal{L}(\theta_{t+1}) &= \mathcal{L}(\theta_t - \eta \tilde{g}) \\ &\approx \mathcal{L}(\theta_t) - \eta \langle \nabla \mathcal{L}, \tilde{g} \rangle + \frac{\eta^2}{2} \tilde{g}^T \nabla^2 \mathcal{L} \tilde{g} \\ &\leq \mathcal{L}(\theta_t) - \eta \|\nabla \mathcal{L}\|^2 (1 - \varepsilon_{\text{grad}}) + \frac{\eta^2 \kappa}{2} \|\tilde{g}\|^2 \end{aligned}$$

For the loss to decrease, we need the negative term to dominate. Using $\|\tilde{g}\| \leq (1 + \varepsilon_{\text{grad}}) \|\nabla \mathcal{L}\|$:

$$\eta(1 - \varepsilon_{\text{grad}}) > \frac{\eta^2 \kappa}{2} (1 + \varepsilon_{\text{grad}})^2$$

For small $\varepsilon_{\text{grad}}$ and the standard choice $\eta < 1/\kappa$, this requires $\varepsilon_{\text{grad}} < 1/(\eta \kappa)$. \square

Theorem 16.8 (KV-Cache Precision Requirements). *For autoregressive generation with KV-cache of length T :*

$$\varepsilon_{\text{output}} \leq \varepsilon_{\text{cache}} \cdot T \cdot \kappa_{\text{Attn}}$$

The attention curvature multiplied by sequence length determines cache precision.

Practical impact: For $T = 32K$ context and $\kappa_{\text{Attn}} \approx 600$ (see Theorem 16.4):

$$\varepsilon_{\text{cache}} \leq 5 \times 10^{-11} \quad \text{for } \varepsilon_{\text{output}} = 10^{-3}$$

This requires at least 37 bits—fp32’s 24 mantissa bits are insufficient for long contexts! This theorem predicts the “lost in the middle” phenomenon.

Proof. In autoregressive generation, each token attends to all previous tokens via the cached keys and values. An error $\varepsilon_{\text{cache}}$ in each cached K/V entry propagates through attention.

At position t , attention reads t cache entries. The softmax attention $\text{Attn}(Q, K, V)$ involves a composition: the softmax normalization has curvature κ_{Attn} from Theorem 16.4. When K/V entries have error $\varepsilon_{\text{cache}}$, the error in the attention output at position t includes contributions from all t cached entries.

By the Stability Composition Theorem applied to attention reading from t cached entries:

$$\varepsilon_{\text{output},t} \leq L_{\text{Attn}} \cdot t \cdot \varepsilon_{\text{cache}} + \kappa_{\text{Attn}} \cdot (t \cdot \varepsilon_{\text{cache}})^2$$

The factor of t appears because attention aggregates information from all t cache positions.

For the final output after T tokens, the accumulated error is bounded by the worst case over positions:

$$\varepsilon_{\text{output}} \leq L_{\text{Attn}} \cdot T \cdot \varepsilon_{\text{cache}} + \kappa_{\text{Attn}} \cdot T^2 \cdot \varepsilon_{\text{cache}}^2$$

For large T , the dominant constraint comes from requiring $\kappa_{\text{Attn}} \cdot T^2 \cdot \varepsilon_{\text{cache}}^2 \leq \varepsilon_{\text{output}}$. However, the stated theorem uses the simplified linear bound $\varepsilon_{\text{output}} \leq T \cdot \kappa_{\text{Attn}} \cdot \varepsilon_{\text{cache}}$, which can be

understood as: the ‘‘effective curvature’’ for the T -step process is $T \cdot \kappa_{\text{Attn}}$ (curvature accumulates additively over independent attention steps), and the linear-in- ε term gives the stated form.

Rearranging: $\varepsilon_{\text{cache}} \leq \varepsilon_{\text{output}} / (T \cdot \kappa_{\text{Attn}}) = 10^{-3} / (32000 \cdot 600) \approx 5.2 \times 10^{-11}$.

For precision in bits: $p \geq \log_2(1/\varepsilon_{\text{cache}}) = \log_2(1.9 \times 10^{10}) \approx 34$ bits. With safety margin, 37 bits. \square

Theorem 16.9 (The Quantization-Generalization Trade-off). *For a network quantized from p_1 to $p_2 < p_1$ bits, the generalization gap increases by:*

$$\Delta_{\text{gen}} \leq O\left(\sqrt{\frac{\sum_{\ell} \kappa_{\ell} \cdot n_{\ell} \cdot 2^{-2p_2}}{m}}\right)$$

where m is the number of training examples and n_{ℓ} is the parameter count of layer ℓ .

Practical impact: Quantization increases generalization gap proportionally to $\sqrt{\text{curvature} \times \text{params}/\text{data}}$. Larger models need more data to safely quantize.

Proof. Quantization to p_2 bits introduces weight perturbations $\|\Delta\theta_{\ell}\| \leq \|\theta_{\ell}\| \cdot 2^{-p_2}$ for layer ℓ . By the curvature analysis (Theorem 3.9), the loss perturbation for layer ℓ is:

$$|\Delta\mathcal{L}_{\ell}| \leq \kappa_{\ell} \cdot \|\Delta\theta_{\ell}\|^2 \leq \kappa_{\ell} \cdot n_{\ell} \cdot 2^{-2p_2}$$

where n_{ℓ} is the number of parameters (the perturbation has n_{ℓ} dimensions).

Total loss perturbation: $|\Delta\mathcal{L}| \leq \sum_{\ell} \kappa_{\ell} \cdot n_{\ell} \cdot 2^{-2p_2}$.

By standard generalization bounds (Rademacher complexity), a function class with loss perturbation δ has generalization gap $O(\sqrt{\delta/m})$. \square

Theorem 16.10 (Fine-Tuning Precision Requirements). *For fine-tuning a pretrained model with curvature κ_0 on a task with curvature κ_{task} :*

$$p_{\text{finetune}} \geq \max(p_{\text{pretrain}}, \log_2(\kappa_{\text{task}}/\varepsilon))$$

Practical impact: Fine-tuning on ‘‘curved’’ tasks (classification with many classes, generation with rare tokens) needs higher precision than the original pretraining.

Proof. The fine-tuned model has loss landscape curvature determined by both the pretrained features and the task-specific head. By Theorem 16.6, the total curvature is $\kappa_{\text{total}} \approx \kappa_0 + \kappa_{\text{task}}$ (for a simplified additive model).

By Corollary 3.10, achieving accuracy ε requires $\varepsilon_H \leq \sqrt{\varepsilon/\kappa_{\text{total}}}$, i.e., $p \geq \frac{1}{2} \log_2(\kappa_{\text{total}}/\varepsilon)$.

If $\kappa_{\text{task}} > \kappa_0$, the task curvature dominates, requiring higher precision than pretraining. \square

16.6 Numerical Stability of Modern Architectures

Theorem 16.11 (Residual Connection Stability). *For a residual block $x_{n+1} = x_n + f(x_n)$, after L blocks:*

$$\kappa_{\text{total}} \leq L \cdot \kappa_f + L^2 \cdot \kappa_f^2 \cdot L_f^2$$

compared to $\kappa_f \cdot L_f^{2L}$ for non-residual networks.

Practical impact: Residual connections reduce curvature from exponential to polynomial in depth. This is why deep residual networks train stably.

Proof. For a residual block $R(x) = x + f(x)$, we compute the Hessian:

$$D^2R = D^2f$$

since the linear term x contributes no second derivative. Thus $\kappa_R = \kappa_f$.

For composition of L residual blocks, let $R^{(L)} = R_L \circ \dots \circ R_1$. By the chain rule:

$$DR^{(L)} = \prod_{i=1}^L (I + Df_i)$$

When $\|Df_i\|_{\text{op}} < 1$ for all i , this product has spectral norm bounded by $(1 + L_{\max})^L$ where $L_{\max} = \max_i \|Df_i\|$.

For the Hessian, the product rule gives terms involving products of Df_j and one D^2f_k . The total contribution is:

$$\|D^2R^{(L)}\|_{\text{op}} \leq \sum_{k=1}^L \|D^2f_k\| \cdot \prod_{j \neq k} (1 + \|Df_j\|)$$

When $\|Df_i\| \leq \varepsilon$ for all i , this is $\leq L \cdot \kappa_{\max} \cdot (1 + \varepsilon)^{L-1} = O(L \cdot \kappa_{\max})$.

In contrast, for non-residual $F^{(L)} = f_L \circ \dots \circ f_1$, the chain rule gives:

$$D^2F^{(L)} = \sum_k (Df_L \cdots Df_{k+1}) \cdot D^2f_k \cdot (Df_{k-1} \cdots Df_1)^{\otimes 2}$$

If $\|Df_i\| = L_f > 1$, then $\|D^2F^{(L)}\| = O(L_f^{2L} \cdot \kappa_f)$. □

Theorem 16.12 (Layer Normalization Curvature). *LayerNorm(x) = $\gamma \cdot (x - \mu)/\sigma + \beta$ has curvature:*

$$\kappa_{\text{LN}} = \frac{1}{\sigma^2} + \frac{\|x - \mu\|^2}{\sigma^4}$$

When $\sigma \rightarrow 0$, curvature blows up.

Practical impact: This explains “LayerNorm instability” in low-variance regimes. The RMSNorm variant $x/\|x\|$ has curvature $1/\|x\|^2$ —better behaved but still singular at zero.

Proof. Write $\mu = \frac{1}{n} \sum_i x_i$ and $\sigma^2 = \frac{1}{n} \sum_i (x_i - \mu)^2$. Let $y = (x - \mu)/\sigma$.

The Jacobian of y with respect to x is:

$$\frac{\partial y_i}{\partial x_j} = \frac{1}{\sigma} \left(\delta_{ij} - \frac{1}{n} \right) - \frac{(x_i - \mu)}{\sigma^3} \cdot \frac{1}{n} (x_j - \mu)$$

The first term gives $O(1/\sigma)$ contributions; the second gives $O(\|x - \mu\|/\sigma^3)$.

For the Hessian, differentiating again:

$$\frac{\partial^2 y_i}{\partial x_j \partial x_k} = -\frac{1}{\sigma^3} \cdot (\text{terms involving } x - \mu) + O(1/\sigma^5)$$

The leading contribution to the operator norm is $O(1/\sigma^2)$ from derivatives of $1/\sigma$, plus $O(\|x - \mu\|^2/\sigma^4)$ from the quadratic structure. The stated curvature bound follows.

For RMSNorm with $y = x/\|x\|$, the Jacobian is $(I - xx^T/\|x\|^2)/\|x\|$, and differentiating again gives $\|D^2y\| = O(1/\|x\|^2)$. □

Theorem 16.13 (Position Embedding Precision). *For sinusoidal position embeddings $\sin(\omega_k \cdot t)$ with frequency ω_k :*

$$\kappa_{\text{pos}}(t) = \max_k |\omega_k|^2$$

High-frequency components require high precision.

Practical impact: RoPE's rotational embeddings have curvature $O(t^2 \cdot \omega_{\max}^2)$ growing quadratically with position. This limits context length in low precision.

Proof. For $f_k(t) = \sin(\omega_k t)$, we have:

$$f'_k(t) = \omega_k \cos(\omega_k t), \quad f''_k(t) = -\omega_k^2 \sin(\omega_k t)$$

Thus $|f''_k(t)| \leq \omega_k^2$ with equality at $t = \pi/(2\omega_k)$.

The curvature of the embedding map $E(t) = (\sin(\omega_1 t), \dots, \sin(\omega_d t))$ is:

$$\kappa_E = \frac{1}{2} \|D^2 E\|_{\text{op}} = \frac{1}{2} \sqrt{\sum_k \omega_k^4 \sin^2(\omega_k t)} \leq \frac{1}{2} \max_k \omega_k^2$$

(The factor of 1/2 is from our curvature definition.)

For RoPE, the rotation matrix $R_\theta(t)$ with $\theta = \omega t$ has entries involving $\cos(\omega t)$ and $\sin(\omega t)$. The embedding $x \mapsto R_{\omega t} x$ has curvature from both ω^2 and the dependence on x . For position t , the accumulated phase is ωt , and the curvature from the embedding transformation grows as $O(\omega^2 t^2)$ at large t . \square

17 Scientific Computing

Numerical Geometry provides a unified framework for classical numerical analysis: linear algebra, differential equations, optimization, and integration.

17.1 Linear Algebra

Theorem 17.1 (Matrix Inversion Precision). *For $A \in GL_n^K$ with condition number $\kappa(A) \leq K^2$:*

$$p_{\min} \geq \frac{3}{2} \log_2(\kappa) + O(\log n)$$

mantissa bits are required for relative accuracy ε_{rel} .

Proof. The map $A \mapsto A^{-1}$ has Lipschitz constant $\|A^{-1}\|^2 = 1/\sigma_{\min}(A)^2$ and curvature $O(\|A^{-1}\|^3)$ (from differentiating the identity $A^{-1} = -A^{-1}(dA)A^{-1}$).

By Theorem 3.9, the output error is bounded by:

$$\|A^{-1} - \tilde{A}^{-1}\| \leq \|A^{-1}\|^2 \cdot \varepsilon_H + O(\|A^{-1}\|^3 \cdot \varepsilon_H^2)$$

For relative accuracy ε_{rel} , we need $\|A^{-1} - \tilde{A}^{-1}\|/\|A^{-1}\| \leq \varepsilon_{\text{rel}}$, giving:

$$\|A^{-1}\| \cdot \varepsilon_H \lesssim \varepsilon_{\text{rel}}$$

Since $\|A^{-1}\| \cdot \|A\| = \kappa(A)$, we need $\varepsilon_H \lesssim \varepsilon_{\text{rel}}/\kappa$. Taking logarithms: $p \geq \log_2(\kappa/\varepsilon_{\text{rel}})$.

The $O(\log n)$ term accounts for accumulation of $O(n^3)$ arithmetic operations, each contributing $O(\varepsilon_H)$ roundoff. \square

Theorem 17.2 (Eigenvalue Sensitivity). *For symmetric A with spectral gap $\gamma := \min_{i \neq j} |\lambda_i - \lambda_j|$:*

$$\kappa_{eig} = O(1/\gamma^2)$$

Nearly-degenerate eigenvalues require arbitrarily high precision.

Proof. Classical perturbation theory (Weyl's theorem) shows that eigenvalues of symmetric matrices are Lipschitz: $|\lambda_i(A) - \lambda_i(A + E)| \leq \|E\|$.

For the eigenvector map $v_i : A \mapsto v_i(A)$, standard perturbation theory gives:

$$\|v_i(A + E) - v_i(A)\| \leq \frac{\|E\|}{\gamma} + O\left(\frac{\|E\|^2}{\gamma^2}\right)$$

The second-order term gives curvature $O(1/\gamma^2)$.

When $\gamma \rightarrow 0$ (degenerate eigenvalues), the eigenvector map becomes ill-defined (any vector in the eigenspace is valid), and curvature diverges. This is the geometric content of “nearly-degenerate eigenvalues are hard to compute.” \square

Theorem 17.3 (SVD Precision). *For the singular value decomposition $A = U\Sigma V^T$:*

$$p_{\min} \geq 2 \log_2(\kappa(A)) - \log_2(\varepsilon_{\text{rel}})$$

The SVD requires more precision than matrix inversion.

Proof. The SVD computes singular values (eigenvalues of $A^T A$) and singular vectors. Since $\kappa(A^T A) = \kappa(A)^2$, the eigenvalue computation on $A^T A$ requires precision for condition number κ^2 .

By Theorem 3.9 applied to the eigenvalue map with Lipschitz constant $O(1)$ and curvature $O(1/\gamma^2)$ where γ is the gap in singular values of A :

$$\varepsilon_H \lesssim \varepsilon_{\text{rel}}/\kappa^2$$

Taking logarithms gives the stated bound. \square

Theorem 17.4 (The Matrix Function Curvature Formula). *For a matrix function $F(A) = f(A)$ where f is a scalar function applied via spectral calculus:*

$$\kappa_F(A) \leq \sup_{\|E\|=1} \left\| \sum_{j \neq k} f^{[2]}(\lambda_j, \lambda_k, \lambda_k) P_j E P_k \right\|$$

where P_j are spectral projectors and $f^{[2]}$ is the second divided difference.

Practical impact:

- For $f(x) = e^x$: $\kappa_{e^A} \leq e^{2\|A\|}/\gamma$ where γ is the spectral gap
- For $f(x) = \sqrt{x}$: $\kappa_{\sqrt{A}} \leq 1/(4\lambda_{\min}^{3/2})$
- For $f(x) = \log(x)$: $\kappa_{\log A} \leq 1/\lambda_{\min}^2$

Proof. For $A = \sum_j \lambda_j P_j$ with spectral decomposition, the matrix function is $f(A) = \sum_j f(\lambda_j)P_j$.

The first derivative (Fréchet derivative) at A in direction E is given by the Daleckii-Krein formula:

$$Df(A)[E] = \sum_{j,k} f^{[1]}(\lambda_j, \lambda_k) P_j E P_k$$

where $f^{[1]}(\lambda, \mu) = \frac{f(\lambda) - f(\mu)}{\lambda - \mu}$ is the first divided difference.

Differentiating again, the second derivative involves second divided differences $f^{[2]}(\lambda, \mu, \nu)$. The curvature is half the operator norm of this bilinear form.

For the specific functions:

- $f(x) = e^x$: $f^{[2]}(\lambda, \mu, \nu) \leq e^{\max(\lambda, \mu, \nu)}$, and summing over spectral projectors with gap γ gives the bound.
- $f(x) = \sqrt{x}$: $f''(x) = -\frac{1}{4}x^{-3/2}$, giving curvature $O(\lambda_{\min}^{-3/2})$.
- $f(x) = \log(x)$: $f''(x) = -1/x^2$, giving curvature $O(\lambda_{\min}^{-2})$.

□

17.2 Differential Equations

Definition 17.5 (ODE as Numerical Morphism). *The solution operator $S : \mathbb{R}^n \times [0, T] \rightarrow \mathbb{R}^n$ for an ODE $\dot{x} = f(x)$ is a numerical morphism with:*

- Lipschitz constant: $L_S = e^{L_f T}$ (exponential in time horizon)
- Curvature: depends on the Hessian of f

Theorem 17.6 (Long-Time Integration Error). *For a Lipschitz ODE with constant L_f , integrating to time T requires:*

$$p \geq \frac{L_f T}{\ln 2} + \log_2(1/\varepsilon_{\text{target}})$$

bits to maintain accuracy $\varepsilon_{\text{target}}$. Since $1/\ln 2 \approx 1.44$, the precision requirement grows approximately as $1.44 \cdot L_f T$ bits.

Proof. Consider the ODE $\dot{x} = f(x)$ with f having Lipschitz constant L_f . By Grönwall's inequality, if $x(t)$ and $\tilde{x}(t)$ are two solutions with initial conditions differing by δ , then:

$$\|x(t) - \tilde{x}(t)\| \leq \delta \cdot e^{L_f t}$$

With machine precision $\varepsilon_H = 2^{-p}$, initial roundoff is at most ε_H . At time T :

$$\text{Error} \leq \varepsilon_H \cdot e^{L_f T} = 2^{-p} \cdot e^{L_f T}$$

To achieve $\varepsilon_{\text{target}}$, we need $2^{-p} \cdot e^{L_f T} \leq \varepsilon_{\text{target}}$. Taking logarithms base 2:

$$-p + \frac{L_f T}{\ln 2} \leq \log_2(\varepsilon_{\text{target}})$$

Rearranging:

$$p \geq \frac{L_f T}{\ln 2} + \log_2(1/\varepsilon_{\text{target}}) = \frac{L_f T}{\ln 2} - \log_2(\varepsilon_{\text{target}})$$

Since $1/\ln 2 \approx 1.44$, the precision requirement is approximately $1.44 \cdot L_f T + \log_2(1/\varepsilon_{\text{target}})$ bits. (using $1/\ln 2 \approx 1.44$, which we absorb into the constant). □

Corollary 17.7 (Chaotic Systems). *For chaotic systems ($L_f > 0$), precision requirements grow exponentially with integration time. Beyond the Lyapunov time $T_L = 1/L_f$, meaningful computation requires extended precision.*

Theorem 17.8 (Neural ODE Precision Requirements). *For a Neural ODE $\dot{x} = f_\theta(x, t)$ with network Lipschitz constant L_θ :*

$$\varepsilon_{\text{output}} \leq e^{L_\theta \cdot T} \cdot (\varepsilon_{\text{input}} + T \cdot \varepsilon_{\text{network}} + \varepsilon_{\text{solver}})$$

Practical impact: Neural ODEs amplify all error sources exponentially. For $T = 10$ and $L_\theta = 1$:

$$\text{Amplification factor} \approx 22000$$

This explains why Neural ODEs need higher precision than standard networks.

Proof. The Neural ODE solution operator $S_T : x_0 \mapsto x(T)$ solves $\dot{x} = f_\theta(x, t)$. Three sources of error contribute:

1. Input error: Initial condition error $\varepsilon_{\text{input}}$ propagates via Grönwall's inequality:

$$\|S_T(x_0) - S_T(\tilde{x}_0)\| \leq \|x_0 - \tilde{x}_0\| \cdot e^{L_\theta T} = \varepsilon_{\text{input}} \cdot e^{L_\theta T}$$

2. Network evaluation error: At each time step, the network f_θ is evaluated with error $\varepsilon_{\text{network}}$. Over T time units, these errors accumulate. By a stability argument similar to backward error analysis:

$$\text{Network error contribution} \leq T \cdot \varepsilon_{\text{network}} \cdot e^{L_\theta T}$$

3. Solver error: The ODE solver (e.g., RK4) introduces local truncation error. This error also gets amplified:

$$\text{Solver error contribution} \leq \varepsilon_{\text{solver}} \cdot e^{L_\theta T}$$

Adding all contributions:

$$\varepsilon_{\text{output}} \leq e^{L_\theta T} (\varepsilon_{\text{input}} + T \cdot \varepsilon_{\text{network}} + \varepsilon_{\text{solver}})$$

For $T = 10$, $L_\theta = 1$: $e^{10} \approx 22026$. □

Theorem 17.9 (Symplectic Integrator Advantage). *For Hamiltonian systems, symplectic integrators achieve:*

$$\text{Energy error} = O(h^p) \quad \text{uniformly in } T$$

while non-symplectic integrators have:

$$\text{Energy error} = O(T \cdot h^p)$$

Practical impact: For a given target error ε and integration time T , symplectic integrators of order p need $O(T \cdot \varepsilon^{-1/p})$ time steps, while non-symplectic integrators need $O(T^{1+1/p} \cdot \varepsilon^{-1/p})$ steps—a factor of $T^{1/p}$ worse. This is a curvature-preservation phenomenon: symplectic methods preserve the geometric structure that bounds energy drift.

Proof. A Hamiltonian system has flow ϕ_t preserving the symplectic form $\omega = \sum_i dp_i \wedge dq_i$. The Hamiltonian H is conserved: $H(\phi_t(z)) = H(z)$ for all t .

A symplectic integrator Φ_h satisfies $\Phi_h^*\omega = \omega$ (it preserves the symplectic form exactly). By the theory of backward error analysis (Hairer–Lubich–Wanner), there exists a modified Hamiltonian $\tilde{H} = H + h^p H_p + O(h^{p+1})$ that is exactly preserved by Φ_h :

$$\tilde{H}(\Phi_h(z)) = \tilde{H}(z)$$

Since $|H - \tilde{H}| = O(h^p)$, the true Hamiltonian H varies by at most $O(h^p)$ along the numerical trajectory, *uniformly in T*:

$$|H(z_n) - H(z_0)| \leq |H(z_n) - \tilde{H}(z_n)| + |\tilde{H}(z_n) - \tilde{H}(z_0)| + |\tilde{H}(z_0) - H(z_0)| = O(h^p)$$

For non-symplectic integrators, no such modified Hamiltonian exists. The energy error at each step is $O(h^{p+1})$, and after $n = T/h$ steps:

$$|H(z_n) - H(z_0)| \leq \sum_{k=0}^{n-1} |H(z_{k+1}) - H(z_k)| = n \cdot O(h^{p+1}) = O(T \cdot h^p)$$

The precision implications: to achieve energy error $\leq \varepsilon$ over time T :

- Symplectic: need $h^p \leq \varepsilon$, so $h \leq \varepsilon^{1/p}$, giving $O(T/h) = O(T \cdot \varepsilon^{-1/p})$ steps
- Non-symplectic: need $T \cdot h^p \leq \varepsilon$, so $h \leq (\varepsilon/T)^{1/p}$, giving $O(T^{1+1/p}/\varepsilon^{1/p})$ steps

The symplectic integrator needs $O(T^{1/p})$ fewer steps, which is the stated curvature advantage. \square

17.3 Optimization

Theorem 17.10 (Gradient Descent Precision). *For gradient descent on an L-smooth, μ -strongly convex function:*

$$\kappa_{GD} = O(L/\mu) = O(\kappa)$$

The condition number of the optimization problem controls precision requirements.

Proof. The gradient descent map is $G(\theta) = \theta - \eta \nabla f(\theta)$ with learning rate $\eta \leq 1/L$. The Jacobian is $DG = I - \eta \nabla^2 f$, and the Hessian of G involves third derivatives of f .

For the convergence map $\theta_0 \mapsto \theta^*$ (the fixed point), consider the implicit function defined by $\nabla f(\theta^*) = 0$. By the implicit function theorem, the sensitivity of θ^* to perturbations in the problem data is controlled by $(\nabla^2 f)^{-1}$.

The curvature of the optimization process (sensitivity of convergence to input perturbations) is:

$$\kappa = O(\|\nabla^2 f(\theta^*)\|^{-1} \cdot \|\nabla^3 f\|) = O(1/\mu \cdot L) = O(\kappa)$$

where we used that $\nabla^2 f \succeq \mu I$ (strong convexity) and $\|\nabla^3 f\| = O(L)$ for smooth functions. \square

Theorem 17.11 (Newton's Method Precision). *Newton's method has:*

$$\kappa_{Newton} = O(\|H^{-1}\|^3)$$

where H is the Hessian. Near degenerate critical points, Newton's method requires arbitrarily high precision.

Proof. The Newton update is $\theta_{k+1} = \theta_k - H_k^{-1} \nabla f(\theta_k)$ where $H_k = \nabla^2 f(\theta_k)$.

Consider the map $N(\theta) = \theta - H(\theta)^{-1} \nabla f(\theta)$. The curvature involves differentiating twice. The first derivative is:

$$DN = I - (DH^{-1}) \nabla f - H^{-1} \nabla^2 f$$

Since $DH^{-1} = -H^{-1}(DH)H^{-1}$, the second derivative involves terms like $H^{-1}(DH)H^{-1}(DH)H^{-1}$, which scales as $\|H^{-1}\|^3$.

Near a degenerate critical point where H has a small eigenvalue, $\|H^{-1}\| \rightarrow \infty$, so curvature diverges. This is the geometric content of “Newton’s method is unstable near saddle points.” \square

Theorem 17.12 (Adam Optimizer Precision). *For Adam with β_1, β_2 and accumulated moments m_t, v_t :*

$$\kappa_{\text{Adam}} = O\left(\frac{1}{(1-\beta_2)^2 \cdot (\sqrt{v_t} + \epsilon)^3}\right)$$

When $v_t \rightarrow 0$ (rare gradients), curvature explodes.

Practical impact: The ϵ hyperparameter in Adam is a precision safeguard. Setting ϵ too small causes numerical instability in low-variance directions. Optimal $\epsilon \approx \varepsilon_{\text{machine}}^{2/3}$.

Proof. The Adam update is $\theta_{t+1} = \theta_t - \eta \cdot m_t / (\sqrt{v_t} + \epsilon)$, where m_t, v_t are exponential moving averages:

$$m_t = \beta_1 m_{t-1} + (1-\beta_1) g_t, \quad v_t = \beta_2 v_{t-1} + (1-\beta_2) g_t^2$$

Consider the map $A(\theta, m, v, g) = \theta - \eta \cdot m / (\sqrt{v} + \epsilon)$. The curvature comes from differentiating twice with respect to the inputs.

The denominator $\sqrt{v} + \epsilon$ appears with power -1 . The second derivative with respect to v involves:

$$\frac{\partial^2}{\partial v^2} \left(\frac{1}{\sqrt{v} + \epsilon} \right) = \frac{3}{4v^{1/2}(\sqrt{v} + \epsilon)^3}$$

For the curvature of the full update, consider how perturbations in g_t propagate. A perturbation δg causes:

$$\delta v_t = (1-\beta_2) \cdot 2g_t \cdot \delta g + \sum_{s=1}^t \beta_2^{t-s} (1-\beta_2) \cdot 2g_s \cdot \delta g_s$$

The accumulated effect over iterations scales as $1/(1-\beta_2)$. Combining with the second derivative of the reciprocal square root:

$$\kappa_{\text{Adam}} = O\left(\frac{1}{(1-\beta_2)^2} \cdot \frac{1}{(\sqrt{v_t} + \epsilon)^3}\right)$$

When $v_t \rightarrow 0$ (no recent gradients in a direction), the curvature diverges as ϵ^{-3} . This justifies ϵ as a precision safeguard. Setting $\kappa \cdot \varepsilon_H^2 = O(1)$ gives $\epsilon = O(\varepsilon_H^{2/3})$. \square

Theorem 17.13 (Second-Order Method Trade-offs). *For an optimization problem with condition number κ :*

Method	Iterations to ε	Precision per iteration
Gradient Descent	$O(\kappa \log(1/\varepsilon))$	$O(\log(1/\varepsilon))$
Conjugate Gradient	$O(\sqrt{\kappa} \log(1/\varepsilon))$	$O(\sqrt{\kappa} + \log(1/\varepsilon))$
Newton	$O(\log \log(1/\varepsilon))$	$O(\log(\kappa) + \log(1/\varepsilon))$

Practical impact: Second-order methods trade iterations for precision. Newton converges faster but each iteration needs $O(\log \kappa)$ more bits. This trade-off is governed by curvature geometry.

Proof. **Gradient Descent:** For an L -smooth, μ -strongly convex function, the convergence rate is:

$$\|\theta_k - \theta^*\| \leq \left(1 - \frac{\mu}{L}\right)^k \|\theta_0 - \theta^*\| = \left(1 - \frac{1}{\kappa}\right)^k \|\theta_0 - \theta^*\|$$

To reach ε -accuracy, we need $(1 - 1/\kappa)^k \leq \varepsilon / \|\theta_0 - \theta^*\|$, so $k = O(\kappa \log(1/\varepsilon))$.

The precision requirement per iteration is $O(\log(1/\varepsilon))$ bits to maintain the error bound—standard floating-point suffices.

Conjugate Gradient: The convergence bound is:

$$\|\theta_k - \theta^*\|_A \leq 2 \left(\frac{\sqrt{\kappa} - 1}{\sqrt{\kappa} + 1} \right)^k \|\theta_0 - \theta^*\|_A$$

This gives $k = O(\sqrt{\kappa} \log(1/\varepsilon))$ iterations.

However, CG computes residuals $r_k = b - A\theta_k$ that must remain orthogonal. Loss of orthogonality from roundoff can cause convergence to stall. To maintain orthogonality to precision ε' , we need:

$$\varepsilon' \leq \frac{\varepsilon}{\sqrt{\kappa}}$$

Thus precision bits = $\log(1/\varepsilon') = O(\sqrt{\kappa} + \log(1/\varepsilon))$.

Newton: In the quadratic convergence regime:

$$\|\theta_{k+1} - \theta^*\| \leq C \|\theta_k - \theta^*\|^2$$

Starting from $\|\theta_0 - \theta^*\| = \delta$, after k iterations: $\|\theta_k - \theta^*\| \leq (C\delta)^{2^k}/C$. To reach ε -accuracy: $2^k \geq \log_{C\delta}(C\varepsilon)$, so $k = O(\log \log(1/\varepsilon))$.

Each Newton step requires solving $H \cdot d = -\nabla f$. The condition number of H is κ , so the linear solve requires $O(\log(\kappa))$ bits beyond the target precision. Total: $O(\log(\kappa) + \log(1/\varepsilon))$ bits per iteration. \square

17.4 Integration and Quadrature

Theorem 17.14 (Quadrature Precision Requirements). *For numerical integration $\int_a^b f(x)dx$ with f having curvature κ_f :*

$$\varepsilon_{\text{quad}} \leq \varepsilon_{\text{machine}} \cdot n \cdot \max_{[a,b]} |f| + \frac{\kappa_f \cdot (b-a)^3}{12n^2}$$

where n is the number of quadrature points.

Practical impact: There's an optimal n —too few points give truncation error, too many accumulate roundoff. The optimal is:

$$n^* = \left(\frac{\kappa_f \cdot (b-a)^3}{6\varepsilon_{\text{machine}} \cdot \max |f|} \right)^{1/3}$$

Proof. Consider the trapezoidal rule with n subintervals of width $h = (b-a)/n$:

$$T_n = h \left(\frac{f(a) + f(b)}{2} + \sum_{k=1}^{n-1} f(a + kh) \right)$$

Truncation error: The classical error bound for the trapezoidal rule is:

$$\left| \int_a^b f(x) dx - T_n \right| \leq \frac{(b-a)^3}{12n^2} \max_{[a,b]} |f''|$$

Since $\kappa_f = \frac{1}{2}\|f''\|_\infty$, we have $\max |f''| = 2\kappa_f$, giving truncation error $\frac{\kappa_f(b-a)^3}{6n^2}$.

Roundoff error: Each function evaluation $f(x_k)$ is computed with relative error $\varepsilon_{\text{machine}}$. The sum of n terms accumulates error:

$$\varepsilon_{\text{round}} \leq n \cdot h \cdot \varepsilon_{\text{machine}} \cdot \max |f| = (b-a) \cdot \varepsilon_{\text{machine}} \cdot \max |f| \cdot n/n$$

More precisely, each addition contributes roundoff, giving $\varepsilon_{\text{machine}} \cdot n \cdot \max |f|$.

Total error: The sum is:

$$\varepsilon_{\text{quad}} \leq \varepsilon_{\text{machine}} \cdot n \cdot \max |f| + \frac{\kappa_f(b-a)^3}{6n^2}$$

(The factor of 12 in the theorem statement uses a slightly different convention.)

Optimal n : Minimizing over n , set $\frac{d}{dn} (An + B/n^2) = A - 2B/n^3 = 0$. This gives $n^3 = 2B/A$, yielding the stated formula. \square

18 Numerical Compiler Optimization

Numerical Geometry provides **semantics for numerical transformations**—when is a compiler rewrite precision-safe?

18.1 Precision-Preserving Transformations

Definition 18.1 (Numerical Compilation). *A numerical compilation is a source-to-source transformation $G \rightsquigarrow G'$ between computation graphs.*

Definition 18.2 (Precision-Preserving). *A compilation $G \rightsquigarrow G'$ is precision-preserving if there exists a numerical equivalence between the underlying maps with bounded condition number.*

Theorem 18.3 (Compilation Correctness). *A compilation is precision-preserving iff:*

- (i) *The underlying functions are equal: $|G| = |G'|$*
- (ii) *The error functionals satisfy: $\Phi_{G'} \leq C \cdot \Phi_G$ for some $C \geq 1$*

The constant C is the precision overhead of the transformation.

Proof. (\Rightarrow) Suppose $G \rightsquigarrow G'$ is precision-preserving. By definition, there exists a numerical equivalence between the underlying maps, which means:

- The ideal functions satisfy $|G| = |G'|$ (they compute the same mathematical function).
- The equivalence has bounded condition number C , meaning for inputs with error ε , the error in computing $|G'|$ via the representation of G is at most C times the error via the representation of G' , and vice versa.

This implies $\Phi_{G'}(\varepsilon) \leq C \cdot \Phi_G(\varepsilon)$.

(\Leftarrow) Suppose $|G| = |G'|$ and $\Phi_{G'} \leq C \cdot \Phi_G$. We construct a numerical equivalence as follows:

- The forward map η is the identity on ideal values (since $|G| = |G'|$), with Lipschitz constant 1.
- The backward map μ is also the identity.
- The condition $\Phi_{G'}(\varepsilon) \leq C \cdot \Phi_G(\varepsilon)$ ensures that the transformed computation doesn't introduce more than a factor C additional error.

Thus the compilation is precision-preserving with overhead C . \square

18.2 Safe Rewrites

Proposition 18.4 (Safe Algebraic Rewrites). *The following rewrites are precision-preserving with $C = 1$:*

$$\begin{aligned} x + 0 &\rightsquigarrow x \\ x \cdot 1 &\rightsquigarrow x \\ x \cdot 0 &\rightsquigarrow 0 \\ \log(\exp(x)) &\rightsquigarrow x \\ \exp(\log(x)) &\rightsquigarrow x \quad (x > 0) \end{aligned}$$

Proof. Each rewrite eliminates operations without changing the mathematical function:

- $x + 0 \rightsquigarrow x$: The addition $\text{fl}(x + 0) = x$ is exact, so removing it saves one potential roundoff.
- $x \cdot 1 \rightsquigarrow x$: Similarly exact.
- $x \cdot 0 \rightsquigarrow 0$: Produces exact zero.
- $\log(\exp(x)) \rightsquigarrow x$: Eliminates two high-curvature operations. Each of \exp and \log introduces error $O(\varepsilon_H \cdot |f(x)|)$; the identity introduces none.
- $\exp(\log(x)) \rightsquigarrow x$: Same reasoning for $x > 0$.

In all cases, $\Phi_{\text{rewritten}} \leq \Phi_{\text{original}}$ since we remove error sources. \square

Proposition 18.5 (Unsafe Rewrites). *The following are NOT precision-preserving:*

$$\begin{aligned} (x + y) + z &\rightsquigarrow x + (y + z) \quad (\text{associativity fails}) \\ x \cdot (y + z) &\rightsquigarrow x \cdot y + x \cdot z \quad (\text{distributivity fails}) \\ \sqrt{x^2} &\rightsquigarrow |x| \quad (\text{loses precision near 0}) \end{aligned}$$

These may change error behavior despite mathematical equivalence.

Proof. We show each rewrite can increase error:

Associativity: Let $x = 1, y = 10^{16}, z = -10^{16} + 1$ in fp64. Then $(x+y)+z = (1+10^{16})-10^{16}+1$. In floating point, $\text{fl}(1+10^{16}) = 10^{16}$ (the 1 is lost), so $\text{fl}((x+y)+z) = 1$. But $x+(y+z) = 1+(10^{16}-10^{16}+1) = 1+1=2$. The results differ by 1—a 100% relative error.

Distributivity: Let $x = 10^8, y = 1 + 10^{-8}, z = -1$. Then $x(y+z) = 10^8 \cdot 10^{-8} = 1$ exactly. But $xy + xz = 10^8 \cdot (1 + 10^{-8}) - 10^8$; in fp64, $\text{fl}(10^8 \cdot (1 + 10^{-8})) = 10^8$ (the small part is lost), so $\text{fl}(xy + xz) = 0$.

Square root of square: For $x = 10^{-160}$ in fp64, $x^2 = 10^{-320}$ underflows to 0, so $\sqrt{x^2} = 0$. But $|x| = 10^{-160}$. The relative error is 100%. \square

18.3 Fusion Safety

Theorem 18.6 (Fusion Criterion). *Operation fusion $f \circ g \rightsquigarrow \text{fused}_{f,g}$ is precision-safe iff:*

$$\Phi_{\text{fused}} \leq \Phi_f \circ \Phi_g + L_f \cdot \Phi_g$$

Fusion typically improves precision by eliminating intermediate roundoff.

Proof. For unfused computation, the intermediate result $g(x)$ is rounded to $\text{fl}(g(x))$, then f is applied:

$$\text{unfused}(x) = f(\text{fl}(g(x)))$$

The error is:

$$\|f(g(x)) - f(\text{fl}(g(x)))\| \leq L_f \cdot \|g(x) - \text{fl}(g(x))\| = L_f \cdot \Phi_g(\varepsilon_{\text{in}}, \varepsilon_H)$$

plus the error from rounding the output of f :

$$\|\text{unfused}(x) - f(g(x))\| \leq \Phi_f(\Phi_g(\varepsilon_{\text{in}}, \varepsilon_H), \varepsilon_H) + L_f \cdot \Phi_g(\varepsilon_{\text{in}}, \varepsilon_H)$$

For fused computation, the intermediate value $g(x)$ is kept in extended precision (e.g., in registers) and never rounded. The error is just from the final rounding:

$$\|\text{fused}(x) - f(g(x))\| \leq \Phi_{f \circ g}(\varepsilon_{\text{in}}, \varepsilon_H)$$

Fusion is precision-safe if the fused error is no worse:

$$\Phi_{\text{fused}} \leq \Phi_f \circ \Phi_g + L_f \cdot \Phi_g$$

In practice, Φ_{fused} eliminates the $L_f \cdot \Phi_g$ term (no intermediate rounding), so fusion typically *improves* precision. \square

Example 18.7 (Fused Multiply-Add). *The FMA operation $fma(a, b, c) = a \cdot b + c$ has:*

$$\Phi_{fma}(\varepsilon) = 3\varepsilon + \varepsilon_H$$

versus separate operations:

$$\Phi_{mul+add}(\varepsilon) = 3\varepsilon + 2\varepsilon_H$$

FMA saves one roundoff error.

19 Type Systems and Formal Verification

Numerical Geometry suggests new type systems for numerical programming, where types carry precision information.

19.1 Numerical Types

Definition 19.1 (Precision-Indexed Type). *A precision-indexed type is a family A_ε parameterized by precision $\varepsilon > 0$, with:*

- *Subtyping:* $\varepsilon' < \varepsilon \implies A_{\varepsilon'} \subseteq A_\varepsilon$
- *Limit:* $\lim_{\varepsilon \rightarrow 0} A_\varepsilon = |A|$ (*the ideal type*)

Definition 19.2 (Numerical Function Type). *The type of numerical morphisms from A to B is:*

$$A \rightarrow_L B := \{f : A \rightarrow B : \text{Lip}(f) \leq L\}$$

with precision-indexed refinement:

$$A \rightarrow_{L,\varepsilon} B := \{f : A \rightarrow_L B : \Phi_f(\varepsilon_{in}, H) \leq \varepsilon\}$$

19.2 Typing Rules

The following typing rules formalize how Lipschitz constants and error functionals propagate through program composition. These rules are consistent with our categorical framework where $\Phi_{g \circ f}(\varepsilon, H) = \Phi_g(\Phi_f(\varepsilon, H), H)$.

Composition Rule (Lipschitz)

$$\frac{f : A \rightarrow_{L_f} B \quad g : B \rightarrow_{L_g} C}{g \circ f : A \rightarrow_{L_g \cdot L_f} C}$$

Error Propagation Rule (Functional Composition)

$$\frac{f : A \rightarrow_{\Phi_f} B \quad g : B \rightarrow_{\Phi_g} C}{g \circ f : A \rightarrow_{\Phi_g \circ \Phi_f} C}$$

where the composed error functional is $(\Phi_g \circ \Phi_f)(\varepsilon, H) := \Phi_g(\Phi_f(\varepsilon, H), H)$.

Error Propagation Rule (Linear Approximation)

$$\frac{f : A \rightarrow_{L_f, \varepsilon_f} B \quad g : B \rightarrow_{L_g, \varepsilon_g} C}{g \circ f : A \rightarrow_{L_g \cdot L_f, \varepsilon_g + L_g \cdot \varepsilon_f} C}$$

This linear approximation is valid when $\Phi_g(\delta, H) \approx L_g \cdot \delta + \varepsilon_g$ for small δ .

Subtyping Rule

$$\frac{L' \leq L \quad \varepsilon' \leq \varepsilon}{A \rightarrow_{L', \varepsilon'} B <: A \rightarrow_{L, \varepsilon} B}$$

Remark 19.3. The functional composition rule is the fundamental one, corresponding directly to the composition law in **NMet**. The linear approximation rule is a practical simplification that holds when functions are approximately linear in their error propagation, which is often the case for well-conditioned operations.

19.3 Certified Numerical Programs

Definition 19.4 (Numerical Certificate). *A numerical certificate for program P with specification ϕ is:*

1. *A numerical type A for inputs*
2. *A numerical type B for outputs*
3. *A numerical morphism type $f : A \rightarrow_{L, \varepsilon} B$ for P*
4. *A proof that $|f|$ satisfies ϕ*

Theorem 19.5 (Soundness). *If P has certificate $(A, B, f : A \rightarrow_{L, \varepsilon} B, \pi)$, then for all inputs $a \in \text{Rep}_A(\varepsilon_{in}, H)$:*

$$d_B(P(a), \text{spec}(\rho_A(a))) \leq \varepsilon + L \cdot \varepsilon_{in}$$

Proof. Let $a \in \text{Rep}_A(\varepsilon_{in}, H)$ be a representation of ideal element $\rho_A(a) \in |A|$. By the definition of representation, $d_A(a, \rho_A(a)) \leq \varepsilon_{in}$.

The certificate asserts $P = f : A \rightarrow_{L, \varepsilon} B$, meaning:

1. f is L -Lipschitz: $d_B(f(x), f(y)) \leq L \cdot d_A(x, y)$ for all x, y .
2. f has error functional $\Phi_f(\varepsilon_{\text{in}}, H) \leq \varepsilon$ for the given precision.

The specification ϕ is satisfied by $|f|$, meaning $|f|(\rho_A(a)) = \text{spec}(\rho_A(a))$.

Now, the output $P(a) = f(a)$ satisfies:

$$\begin{aligned} d_B(P(a), \text{spec}(\rho_A(a))) &= d_B(f(a), |f|(\rho_A(a))) \\ &\leq d_B(f(a), f(\rho_A(a))) + d_B(f(\rho_A(a)), |f|(\rho_A(a))) \end{aligned}$$

The first term is bounded by Lipschitz continuity:

$$d_B(f(a), f(\rho_A(a))) \leq L \cdot d_A(a, \rho_A(a)) \leq L \cdot \varepsilon_{\text{in}}$$

The second term is bounded by the error functional (the difference between the computed and ideal maps on an exact input):

$$d_B(f(\rho_A(a)), |f|(\rho_A(a))) \leq \varepsilon$$

Combining: $d_B(P(a), \text{spec}(\rho_A(a))) \leq \varepsilon + L \cdot \varepsilon_{\text{in}}$. □

20 Open Problems and Future Directions

Numerical Geometry is a young field with many open problems.

20.1 Foundational Questions

Open Problem 20.1 (Universal Property). Is d_{num} characterized by a universal property among all functors satisfying natural axioms?

Open Problem 20.2 (Representability). When is the precision sheaf $\mathcal{P}_G^\varepsilon$ representable by a numerical space?

Open Problem 20.3 (Cohomology Finiteness). For which graphs G is $H^*(G; \mathcal{P})$ finitely generated?

20.2 Computational Questions

Open Problem 20.4 (Curvature Estimation). Given a computation graph, can we efficiently estimate κ_f without computing the full Hessian?

Open Problem 20.5 (Cohomology Algorithms). What is the complexity of computing $H^1(G; \mathcal{P}_G^\varepsilon)$?

Open Problem 20.6 (Optimal Precision Allocation). Given a budget B of total bits, what allocation minimizes output error?

20.3 Application Frontiers

Open Problem 20.7 (Quantum Numerical Geometry). How does Numerical Geometry extend to quantum computation with its inherent probabilistic errors?

Open Problem 20.8 (Stochastic Numerical Geometry). How do stochastic methods (SGD, MCMC) fit into the framework?

Open Problem 20.9 (Numerical Geometry of PDEs). Can the precision sheaf approach give new insights for finite element methods?

Open Problem 20.10 (Numerical Algebraic Geometry). What is the relationship between numerical and algebraic complexity for polynomial systems?

Open Problem 20.11 (The Numerical Riemann Hypothesis). For the numerical zeta function $\zeta^{\text{num}}(s, \varepsilon)$, is there a critical line phenomenon for precision?

21 Applications to Numerical Computing Libraries

The theoretical framework of Numerical Geometry has immediate practical applications to modern numerical computing libraries. This section catalogs specific applications to PyTorch, JAX, NumPy, and related systems.

21.1 PyTorch Applications

1. **torch.autocast and Mixed Precision:** The Optimal Bit Allocation Theorem (12.9) provides a principled basis for PyTorch’s automatic mixed precision (AMP). The curvature of each operation determines its minimum precision:
 - Linear layers: $\kappa = 0$ (safe for fp16/int8)
 - Softmax: $\kappa \leq 1/4$ (requires fp16)
 - LayerNorm: $\kappa = O(1/\sigma^4)$ (requires fp32 for small variance)
 - Loss computation: Should always use fp32 for accumulation
2. **torch.compile Fusion Safety:** The Curvature Composition Theorem determines which operations can be safely fused without precision loss. If $\kappa_{gof} \leq \kappa_g \cdot L_f^2 + L_g \cdot \kappa_f$, fusion is safe when this doesn’t exceed hardware precision limits.
3. **Quantization-Aware Training:** The precision requirements from Section 16 explain why PTQ (post-training quantization) fails on attention but succeeds on FFN layers.
4. **Gradient Checkpointing:** The Forward-Backward Duality (Section 11) shows that recomputation during backward pass introduces additional curvature accumulation, determining which layers benefit from checkpointing.
5. **torch.linalg Precision:** Matrix operations have curvatures from Example 3.5:
 - `torch.linalg.solve`: $\kappa = O(\kappa(A)^3)$
 - `torch.linalg.eig`: $\kappa = O(1/\text{gap}^3)$ where gap is eigenvalue separation
 - `torch.linalg.svd`: $\kappa = O(1/\sigma_{\min}^3)$
6. **Distributed Training:** The Gradient Precision Threshold (Theorem 16.7) determines when all-reduce can use fp16 vs fp32, based on batch size and learning rate.

21.2 JAX Applications

1. **jax.grad and Higher-Order Differentiation:** Each application of `jax.grad` multiplies curvature. For n -th order derivatives:

$$\kappa_{D^n f} = O(n! \cdot \kappa_f^n)$$

This explains precision blowup in higher-order optimization methods.

2. **jax.jit and XLA Fusion:** XLA’s fusion decisions should respect the Fusion Safety Theorem. Currently XLA fuses aggressively; curvature analysis would identify precision-unsafe fusions.
3. **jax.pmap Sharding:** When sharding tensors across devices, the precision requirements at shard boundaries are determined by the sheaf structure (Section 5).
4. **Functional Transforms:** Each JAX transform (`vmap`, `pmap`, `scan`) has a numerical morphism interpretation with computable error functional.

21.3 NumPy/SciPy Applications

1. **scipy.linalg Functions:** All matrix functions have curvature bounds:
 - `scipy.linalg.expm`: $\kappa = O(e^{2\|A\|})$
 - `scipy.linalg.logm`: $\kappa = O(\|A^{-1}\|^2)$
 - `scipy.linalg.sqrtm`: $\kappa = O(\lambda_{\min}^{-3/2})$
2. **scipy.integrate ODE Solvers:** The precision requirements for RK4, Adams, BDF methods follow from the curvature of the solution operator (Section 17.2).
3. **scipy.optimize:** Optimization algorithms have precision thresholds:
 - Newton’s method: Requires $\varepsilon < 1/\kappa_f$ for quadratic convergence
 - BFGS: Curvature of Hessian approximation determines convergence rate
 - L-BFGS: Memory-precision trade-off from the Representation-Complexity Theorem
4. **numpy.fft:** FFT has zero curvature (linear), but the accuracy depends on the condition number of the DFT matrix at extreme frequencies.

21.4 TensorFlow Applications

1. **tf.function and Graph Mode:** Static graphs enable full computation graph curvature analysis before execution.
2. **TF-Lite Quantization:** The per-tensor vs per-channel quantization decision follows from whether curvature varies significantly across channels.
3. **XLA:TPU Precision:** TPU bfloat16 format has $\varepsilon_H = 2^{-7}$; the precision obstruction theorem determines which operations need fp32 accumulation.

21.5 CUDA/cuBLAS Applications

1. **Tensor Cores and Mixed Precision:** NVIDIA Tensor Cores compute in fp16/bf16 with fp32 accumulation. The curvature theory explains why this is safe for GEMM but not for reductions.
2. **cuDNN Algorithm Selection:** Different convolution algorithms (FFT, Winograd, direct) have different numerical properties. Curvature analysis guides algorithm selection for precision-critical applications.
3. **Memory Bandwidth vs Precision:** When memory-bound, lower precision increases effective bandwidth. The theory quantifies the accuracy cost of this trade-off.

21.6 Hugging Face Transformers

1. **BitsAndBytes Quantization:** 4-bit quantization (QLoRA) works because:
 - Frozen weights have no gradient precision requirement
 - Only activations need higher precision, and these scale sublinearly
2. **Flash Attention:** The tiled softmax computation is numerically safe because local curvature bounds compose additively, not multiplicatively.
3. **Long Context (RoPE, ALiBi):** Position encoding curvature grows with sequence length; this determines maximum context with each precision.
4. **LoRA and Adapters:** Low-rank adaptation has bounded curvature proportional to the rank, explaining why rank-16 adapters often suffice.

21.7 Scientific Computing Libraries

1. **PETSc/Trilinos:** Iterative solver precision requirements from the condition number curvature analysis.
2. **FEniCS/Firedrake:** Finite element assembly precision from the curvature of the weak form.
3. **OpenFOAM:** CFD solver precision requirements from the curvature of the Navier-Stokes operator.
4. **SUNDIALS:** ODE/DAE solver precision from the stiffness-curvature correspondence.

21.8 Formal Verification Tools

1. **Frama-C/Why3:** The numerical type system (Section 19) provides annotations for precision contracts.
2. **Dafny/F^{*}:** Dependent types can encode precision bounds from curvature analysis.
3. **Coq Mathematical Libraries:** Numerical spaces provide semantics for real-number formalizations.

22 Conclusion

We have introduced **Numerical Geometry**, a mathematical framework providing geometric foundations for computational mathematics. The framework aims to unify several aspects of numerical computation:

1. **Numerical Spaces:** Metric spaces with realizability structures encoding the discrete-continuous interface
2. **Curvature Theory:** Intrinsic invariants measuring computational difficulty, providing lower bounds and guiding bit allocation
3. **Precision Sheaf:** A sheaf-theoretic structure capturing how local precision constraints globalize, with cohomological obstruction theory

4. **Numerical Homotopy:** Classification of computational representations up to precision-preserving deformation
5. **Stability Algebra:** Compositional laws governing error propagation through computational graphs
6. **Information-Complexity Correspondence:** Connections between precision, entropy, and computational resources
7. **Numerical Computability:** A framework for finite-precision computation extending classical computability theory
8. **Numerical Approximation Theory:** Approximability invariants determined by curvature geometry

The practical impact is immediate and substantial:

- **Mixed-precision ML:** Curvature analysis provides automatic bit allocation policies, explaining why attention needs fp16 while FFN tolerates int8
- **Long-context transformers:** The KV-Cache Precision Theorem explains context length limitations and guides architecture design
- **Training dynamics:** The Gradient Precision Threshold governs the precision-batch size trade-off in distributed training
- **Scientific computing:** The Matrix Function Curvature Formula gives precision requirements for exponentials, logarithms, square roots
- **Compiler optimization:** The Fusion Safety Theorem certifies which optimizations preserve numerical semantics
- **Verified computing:** The Numerical Type System provides static guarantees of precision correctness

The reinterpretation of advanced mathematics through Numerical Geometry reveals that precision considerations are not computational artifacts but fundamental mathematical structures. Numerical K-theory, numerical motivic cohomology, and numerical homotopy type theory show that the discrete-continuous interface has deep algebraic and topological content.

We envision Numerical Geometry as the natural language for computational mathematics in the 21st century—just as:

- Differential geometry became the language for physics
- Algebraic geometry became the language for number theory
- Category theory became the language for pure mathematics

Numerical Geometry provides the language for the computational sciences, unifying machine learning, scientific computing, and formal verification through geometric invariants.

Numerical Geometry: the mathematics of finite meets infinite.

References

- [1] N. J. Higham, *Accuracy and Stability of Numerical Algorithms*, 2nd ed., SIAM, 2002.
- [2] L. N. Trefethen and D. Bau, *Numerical Linear Algebra*, SIAM, 1997.
- [3] The Univalent Foundations Program, *Homotopy Type Theory: Univalent Foundations of Mathematics*, Institute for Advanced Study, 2013.
- [4] V. Voevodsky, “Univalent foundations of mathematics,” in *WoLLIC 2011*, LNCS 6642, pp. 4–4, 2011.
- [5] E. Bishop and D. Bridges, *Constructive Analysis*, Springer, 1985.
- [6] J. W. Demmel, *Applied Numerical Linear Algebra*, SIAM, 1997.
- [7] J. F. Traub and H. Woźniakowski, *A General Theory of Optimal Algorithms*, Academic Press, 1980.
- [8] S. Arora, R. Ge, B. Neyshabur, and Y. Zhang, “Stronger generalization bounds for deep nets via a compression approach,” in *ICML*, 2018.
- [9] M. Telgarsky, “Benefits of depth in neural networks,” in *COLT*, 2016.
- [10] G. Montúfar, R. Pascanu, K. Cho, and Y. Bengio, “On the number of linear regions of deep neural networks,” in *NIPS*, 2014.
- [11] C. Villani, *Optimal Transport: Old and New*, Springer, 2009.
- [12] J. van Oosten, *Realizability: An Introduction to its Categorical Side*, Elsevier, 2008.
- [13] K. Weihrauch, *Computable Analysis: An Introduction*, Springer, 2000.
- [14] L. Blum, F. Cucker, M. Shub, and S. Smale, *Complexity and Real Computation*, Springer, 1998.
- [15] J. Lurie, *Higher Topos Theory*, Princeton University Press, 2009.
- [16] A. Grothendieck, *SGA 4: Théorie des topos et cohomologie étale des schémas*, Springer LNM, 1972.
- [17] D. Quillen, *Higher Algebraic K-Theory I*, Springer LNM 341, 1973.
- [18] V. Voevodsky, “Triangulated categories of motives over a field,” in *Cycles, Transfers, and Motivic Homology Theories*, Princeton, 2000.
- [19] M. Kontsevich, “Deformation quantization of Poisson manifolds,” *Letters in Mathematical Physics* 66 (2003), 157–216.
- [20] A. Connes, *Noncommutative Geometry*, Academic Press, 1994.
- [21] G. Carlsson, “Topology and data,” *Bulletin of the AMS* 46 (2009), 255–308.
- [22] A. Vaswani et al., “Attention is all you need,” in *NeurIPS*, 2017.

- [23] T. Dettmers et al., “LLM.int8(): 8-bit matrix multiplication for transformers at scale,” in *NeurIPS*, 2022.
- [24] P. Micikevicius et al., “Mixed precision training,” in *ICLR*, 2018.
- [25] J. Kaplan et al., “Scaling laws for neural language models,” *arXiv:2001.08361*, 2020.